

# $GHC_{(STG, Cmm, asm)}$ illustrated

for hardware persons

*exploring some mental models and implementations*

Takenobu T.

"Any sufficiently advanced technology is  
indistinguishable from **magic**."

Arthur C. Clarke

## NOTE

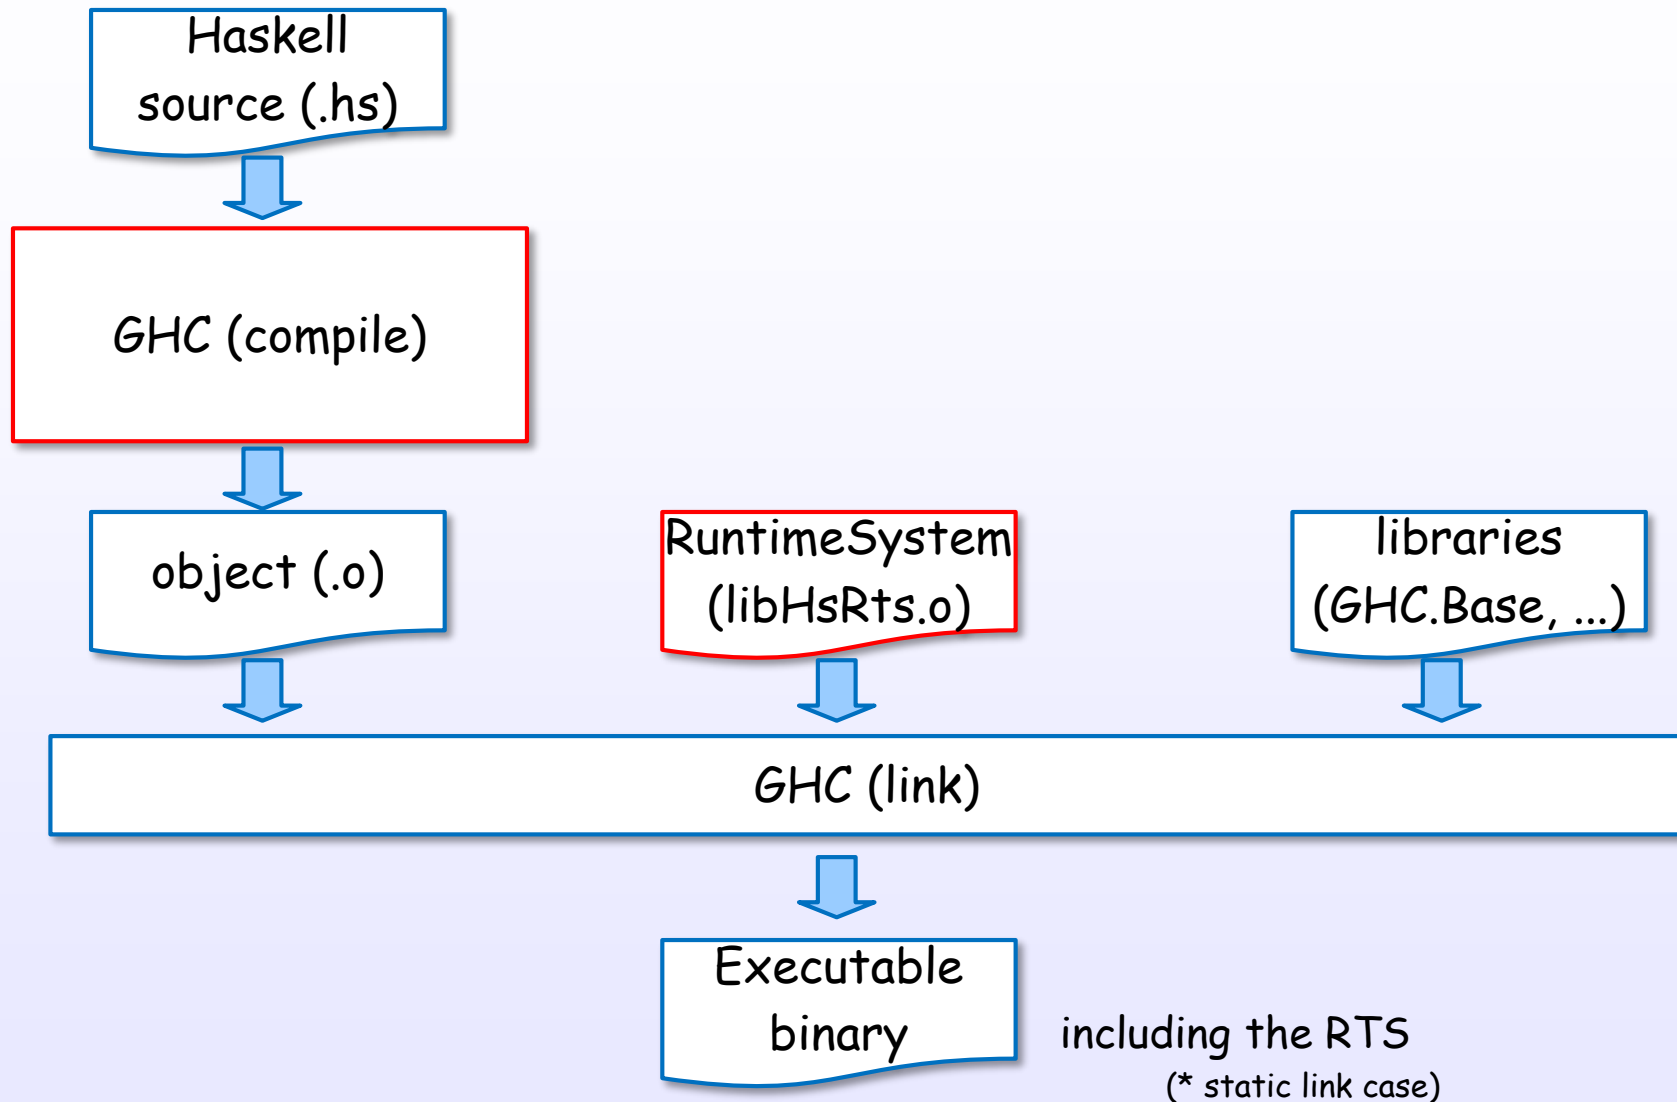
- This is not an official document by the ghc development team.
- Please don't forget "semantics". It's very important.
- This is written for ghc 8.0.

# Contents

- Executable binary
- Compile steps
- Runtime System
- Development languages
- Machine layer/models
- STG-machine
- Heap objects in STG-machine
- STG-machine evaluation
- Pointer tagging
- Thunk and update
- Allocate and free heap objects
- STG - C land interface
- Thread
- Thread context switch
- Creating main and sub threads
- Thread migration
- Heap and Threads
- Threads and GC
- Bound thread
- Spark
- Mvar
- Software transactional memory
- FFI
- IO and FFI
- IO manager
- Bootstrap
- References

Executable binary

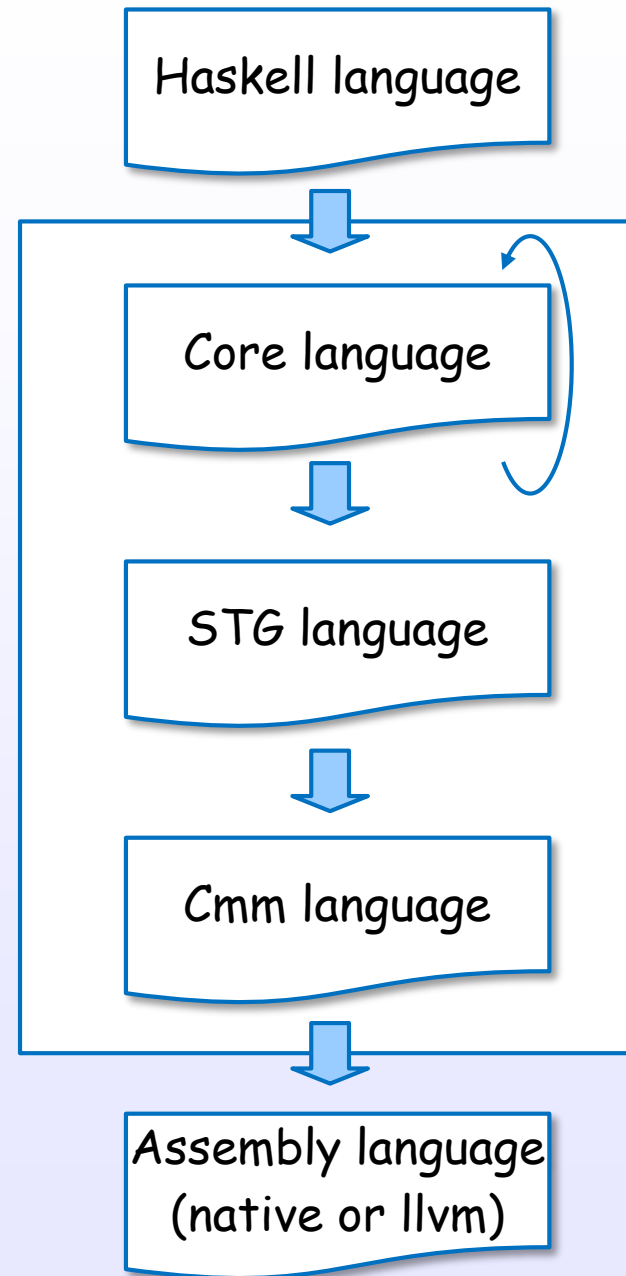
# The GHC = Compiler + Runtime System (RTS)



Compile steps

# GHC transitions between five representations

GHC  
compile  
steps



*each intermediate code can  
be dumped by :*

```
$ ghc -ddump-parsed  
$ ghc -ddump-rn
```

```
$ ghc -ddump-ds  
$ ghc -ddump-simpl  
$ ghc -ddump-prep
```

```
$ ghc -ddump-stg
```

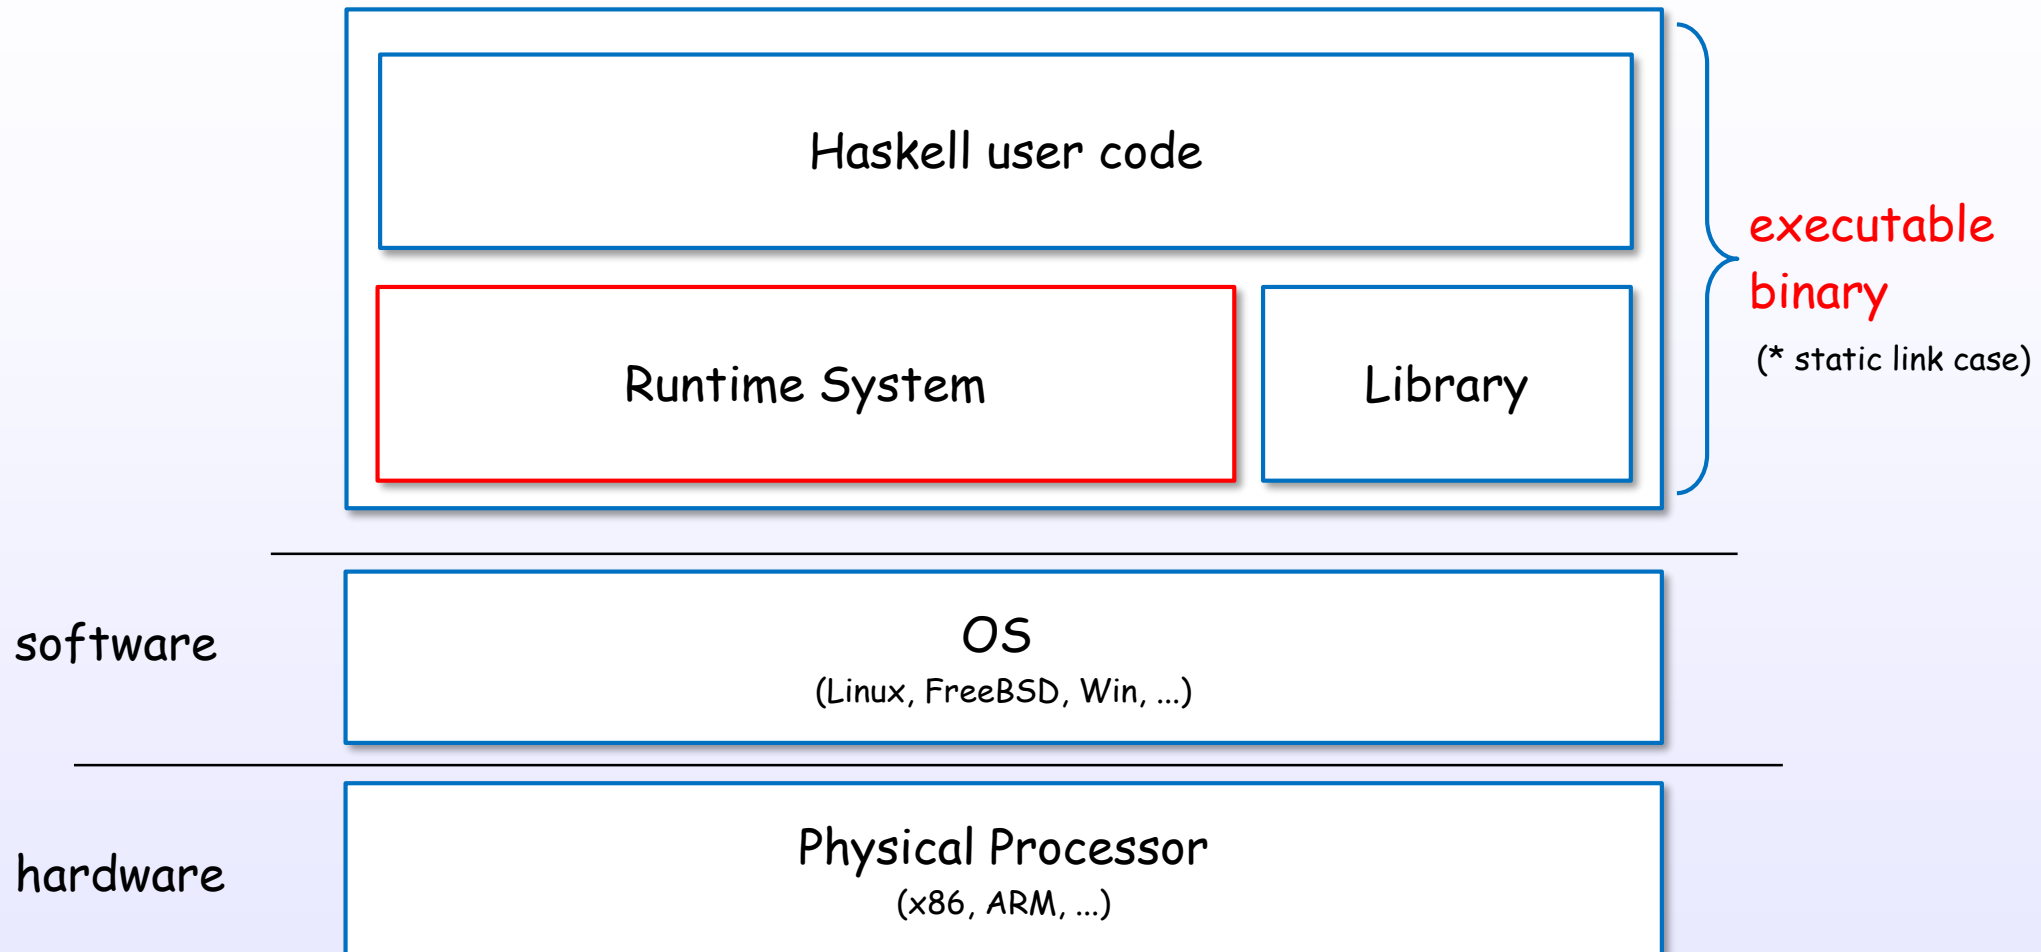
```
$ ghc -ddump-cmm  
$ ghc -ddump-opt-cmm
```

```
$ ghc -ddump-llvm  
$ ghc -ddump-asm
```



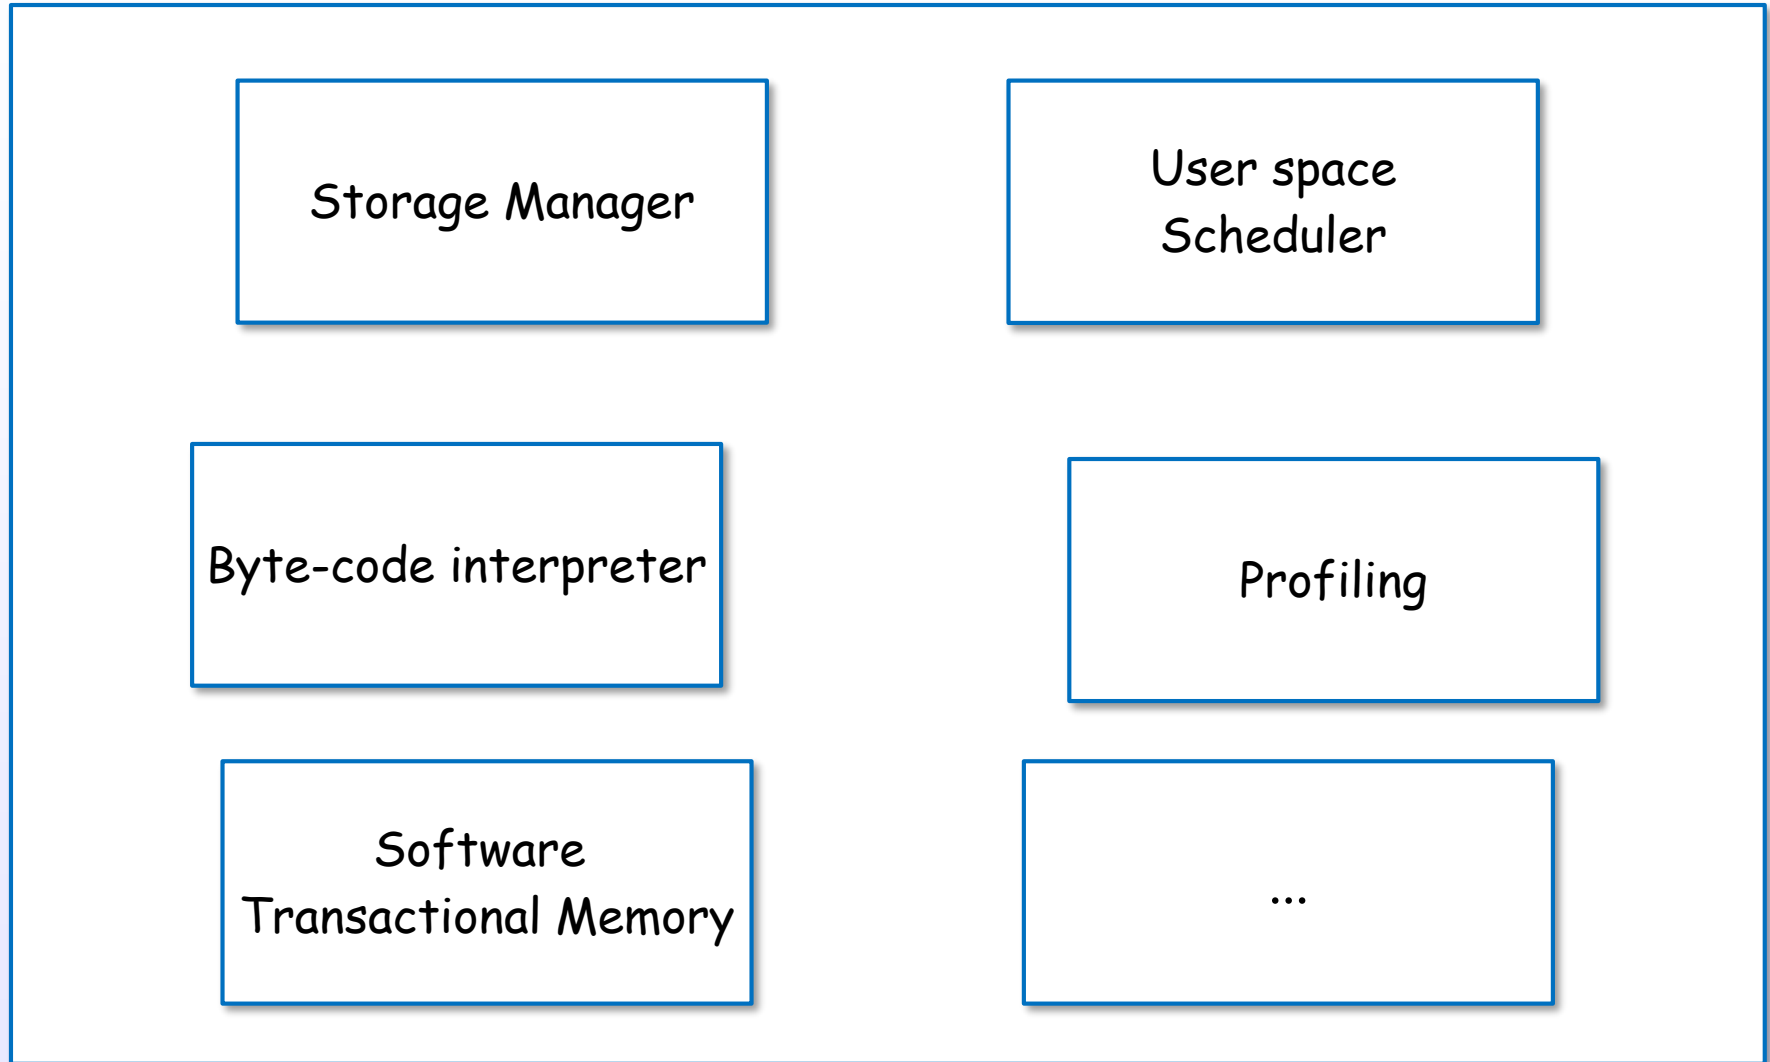
# Runtime System

# Generated binary includes the RTS



# Runtime System includes ...

## Runtime System



Development languages

# The GHC is developed by some languages

compiler

( \$(TOP)/**compiler**/\*)

Haskell

+

Alex (lex)

Happy (yacc)

Cmm (C--)

Assembly

runtime system

( \$(TOP)/**rts**/\*)

C

+

Cmm

Assembly

library

( \$(TOP)/**libraries**/\*)

Haskell

+

C

Machine layer/models

# Machine layer

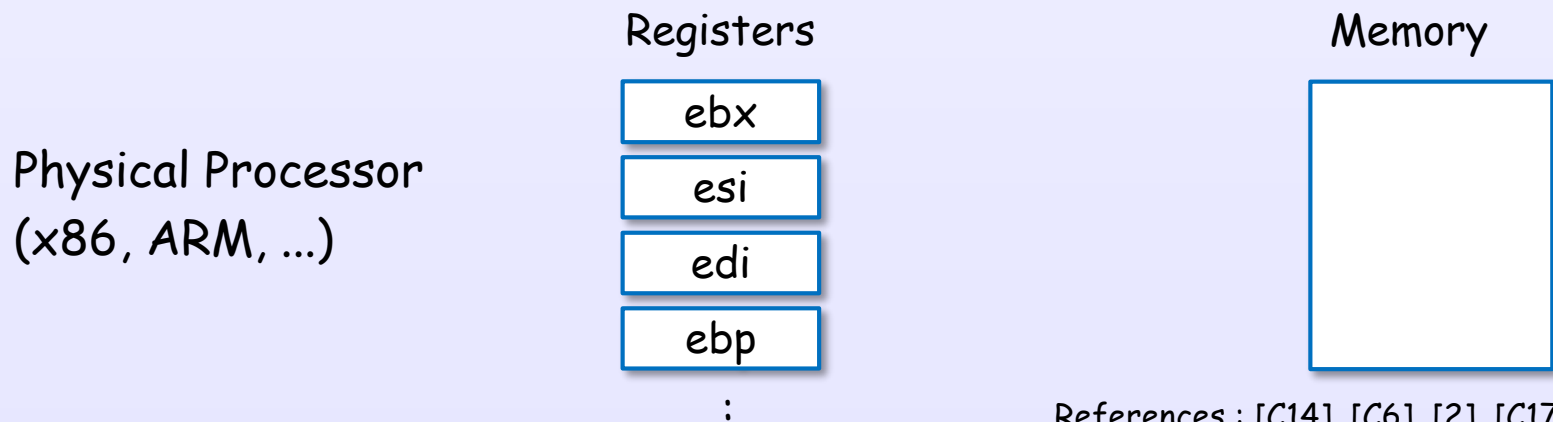
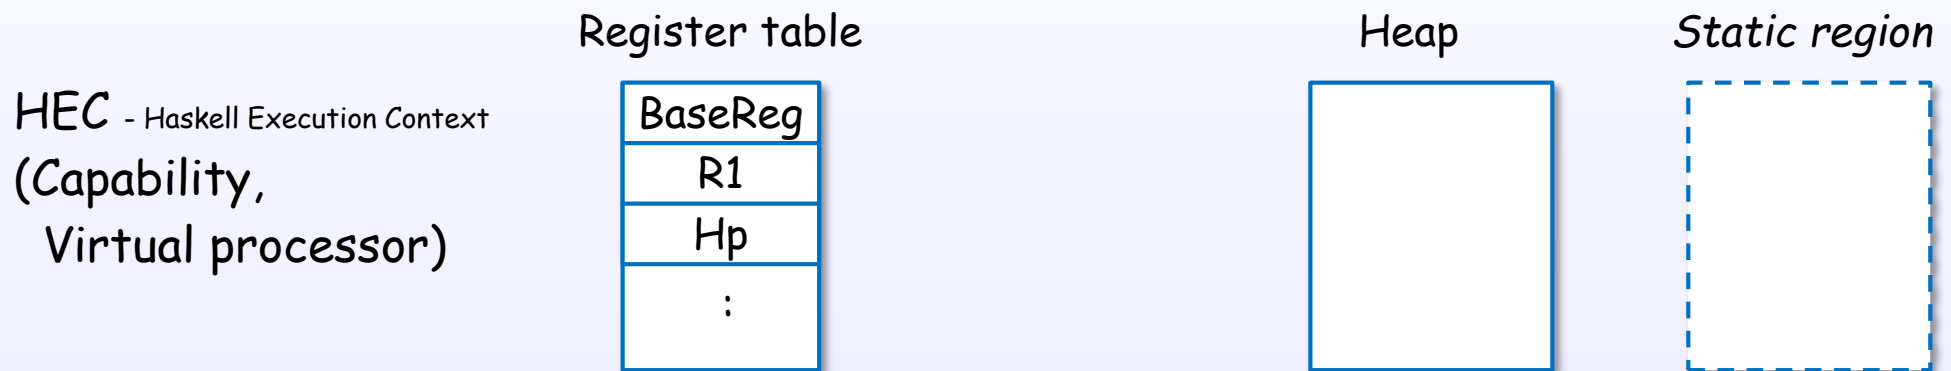
STG-machine  
(Abstract machine)

HEC - Haskell Execution Context  
(Capability, Virtual processor)

Physical Processor  
(x86, ARM, ...)

Each Haskell code is executed in STG semantics.

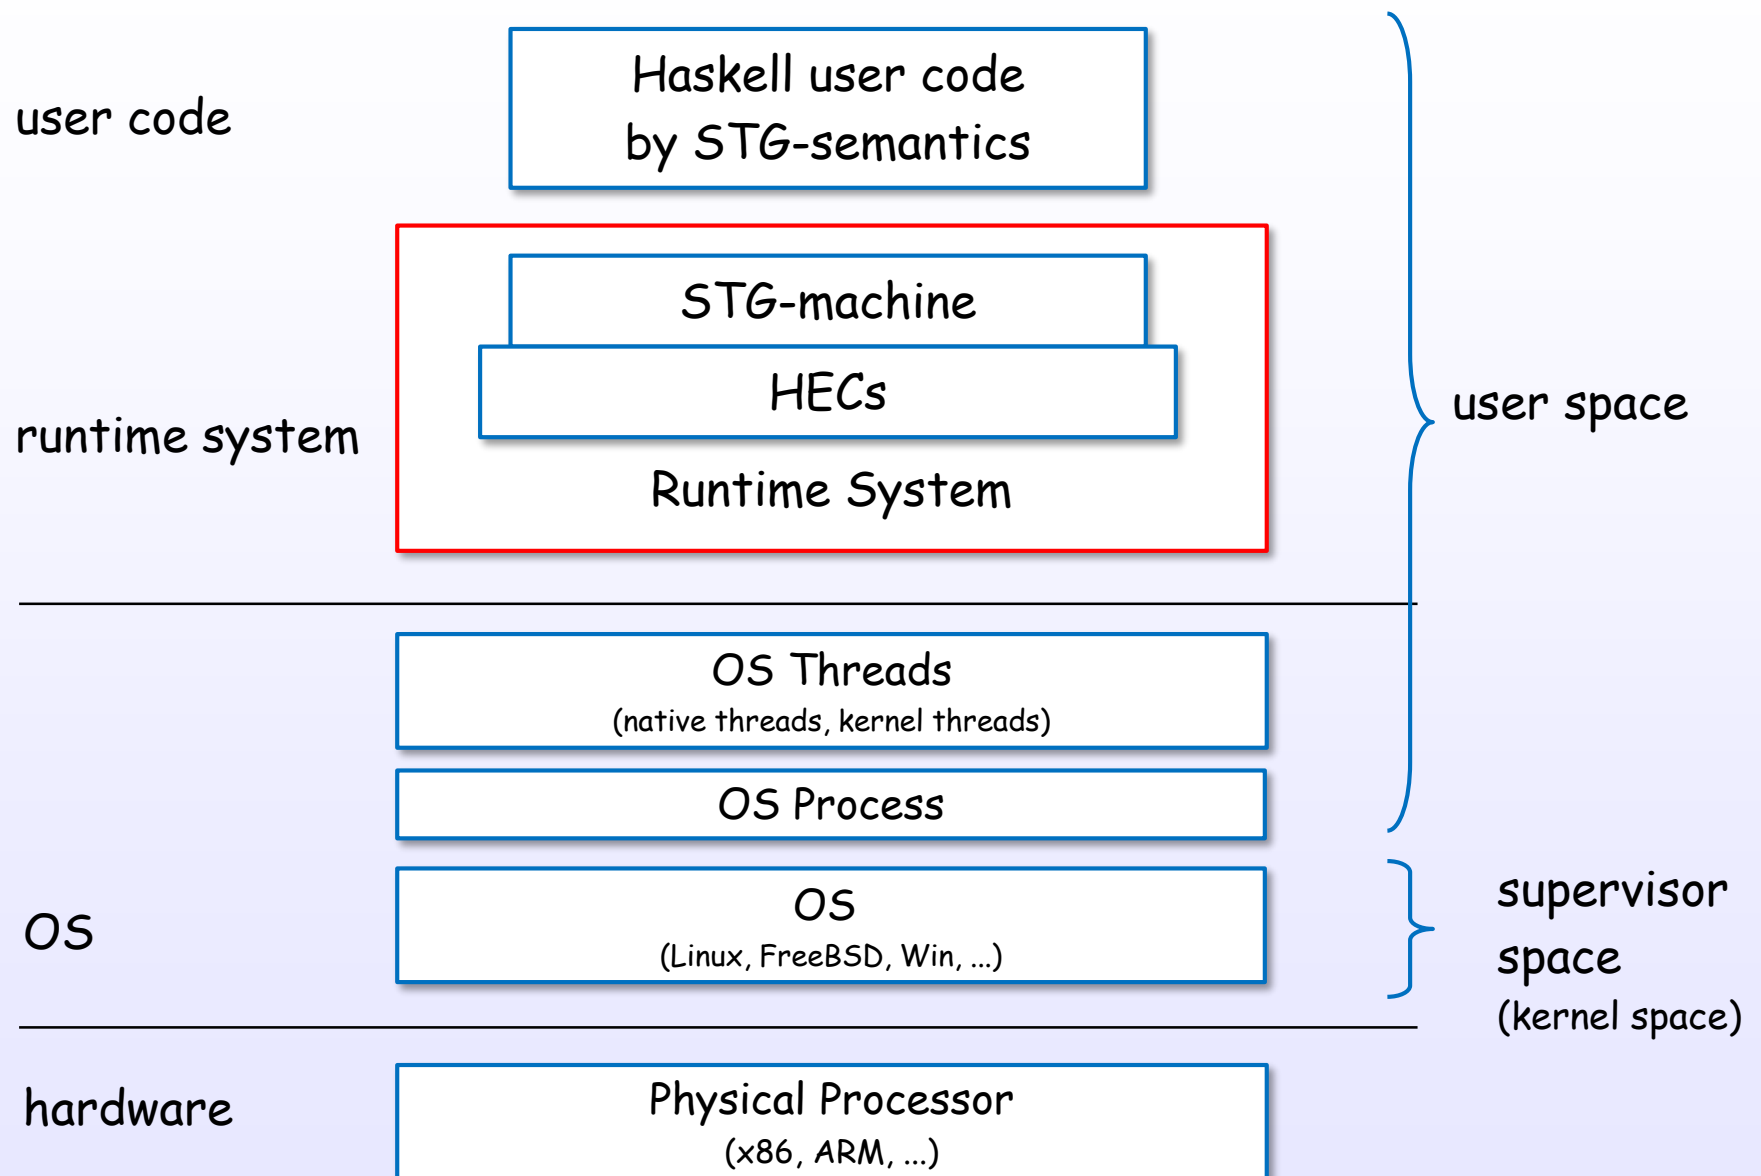
# Machine layer



References : [C14], [C6], [2], [C17], [8], [S15], [S16], [S11]



# Runtime system and HEC

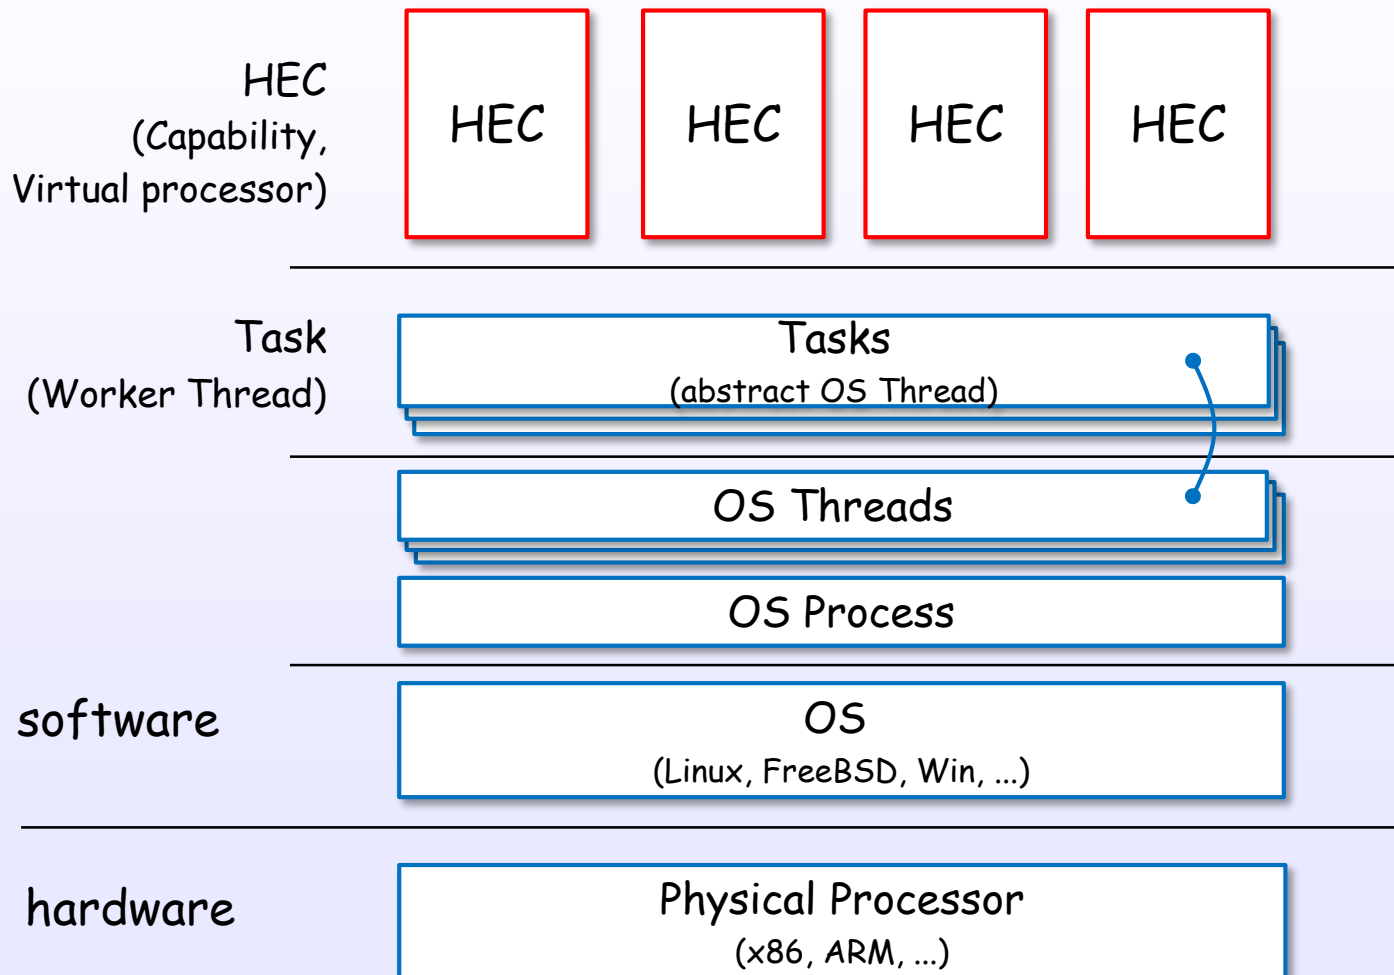


# many HECs

Multi HECs can be generated by compile and runtime options :

```
$ ghc -rtsops -threaded
```

```
$ ./xxx +RTS -N4
```



# HEC (Capability) data structure

[rts/Capability.h] (ghc 8.0)

```
struct Capability_ {
    StgFunTable f;
    StgRegTable r;
    nat no;
    Task *running_task;
    rtsBool in_haskell;
    nat idle;
    rtsBool disabled;
    StgTSO *run_queue_hd;
    StgTSO *run_queue_tl;
    InCall *suspended_ccalls;
    bdescr **mut_lists;
    bdescr **saved_mut_lists;
    bdescr *pinned_object_block;
    bdescr *pinned_object_blocks;
    StgWeak *weak_ptr_list_hd;
    StgWeak *weak_ptr_list_tl;
    int context_switch;
    int interrupt;
    W_ total_allocated;
    #if defined(THREADED_RTS)
        Task *spare_workers;
        nat n_spare_workers;
        Mutex lock;
        Task *returning_tasks_hd;
        Task *returning_tasks_tl;
        Message *inbox;
        SparkPool *sparks;
    #endif
    StgTVarWatchQueue *free_tvar_watch_queues;
    StgInvariantCheckQueue *free_invariant_check_queues;
    StgTRecChunk *free_trec_chunks;
    StgTRecHeader *free_trec_headers;
    nat transaction_tokens;
}
```

*register table*

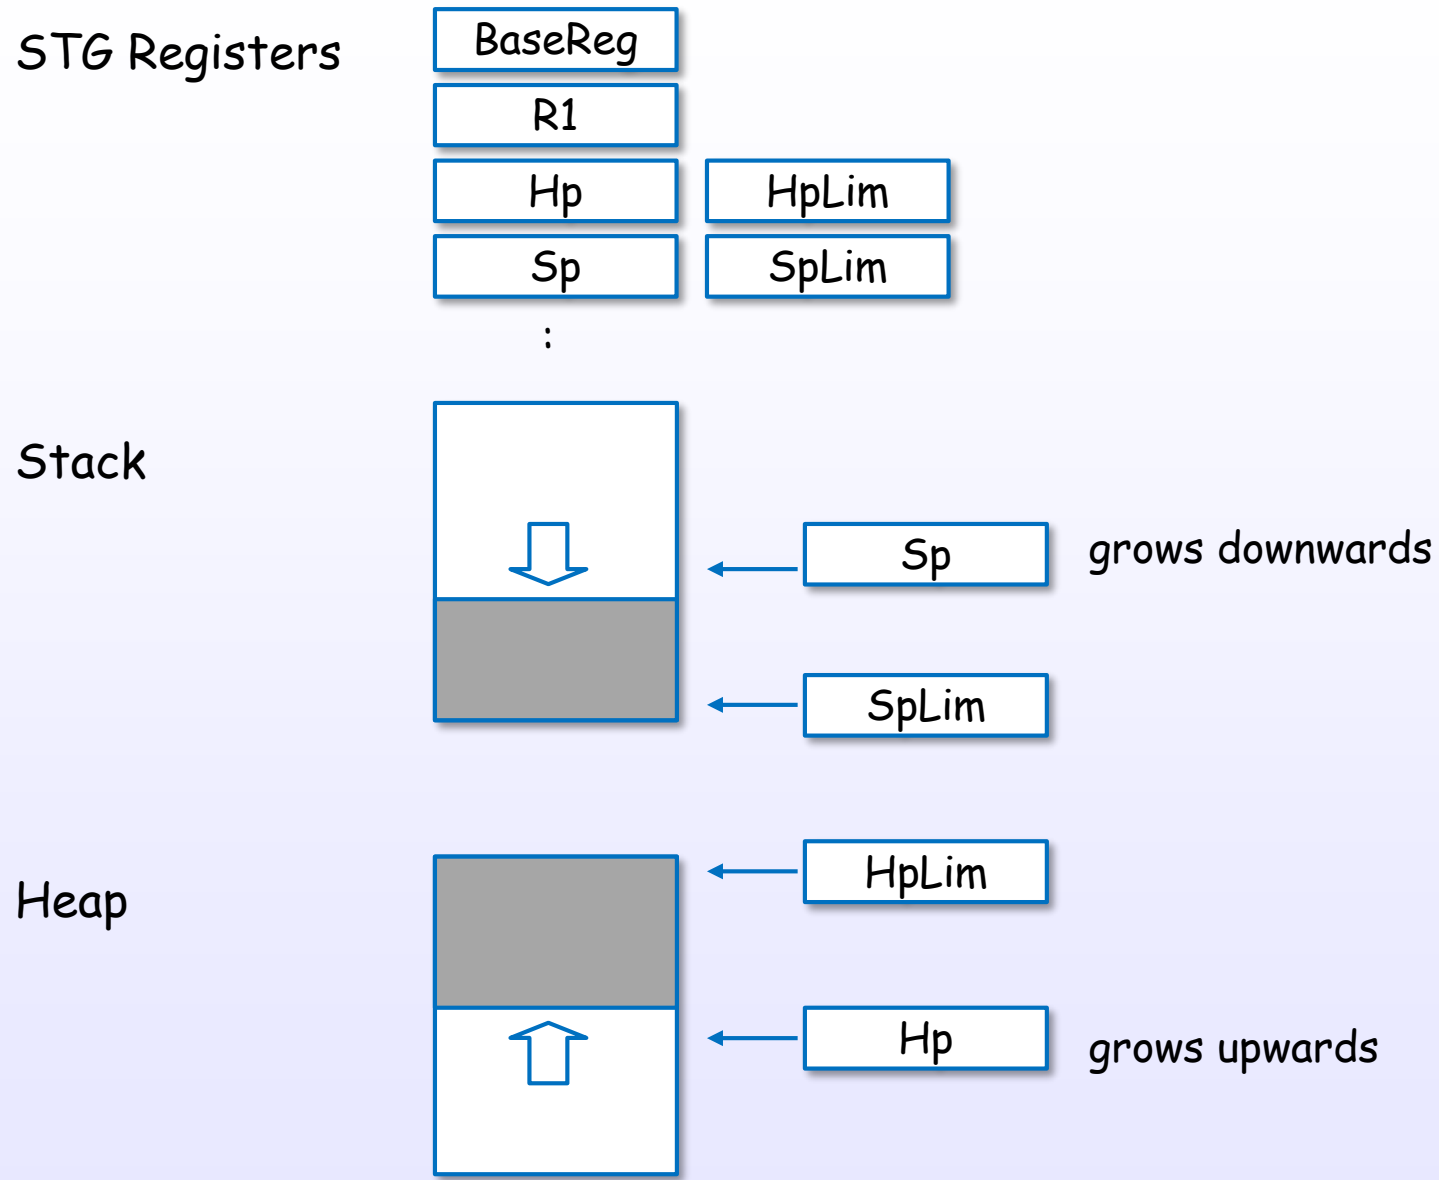
*run queue*

Each HEC (Capability) has a register table and a run queue and ...

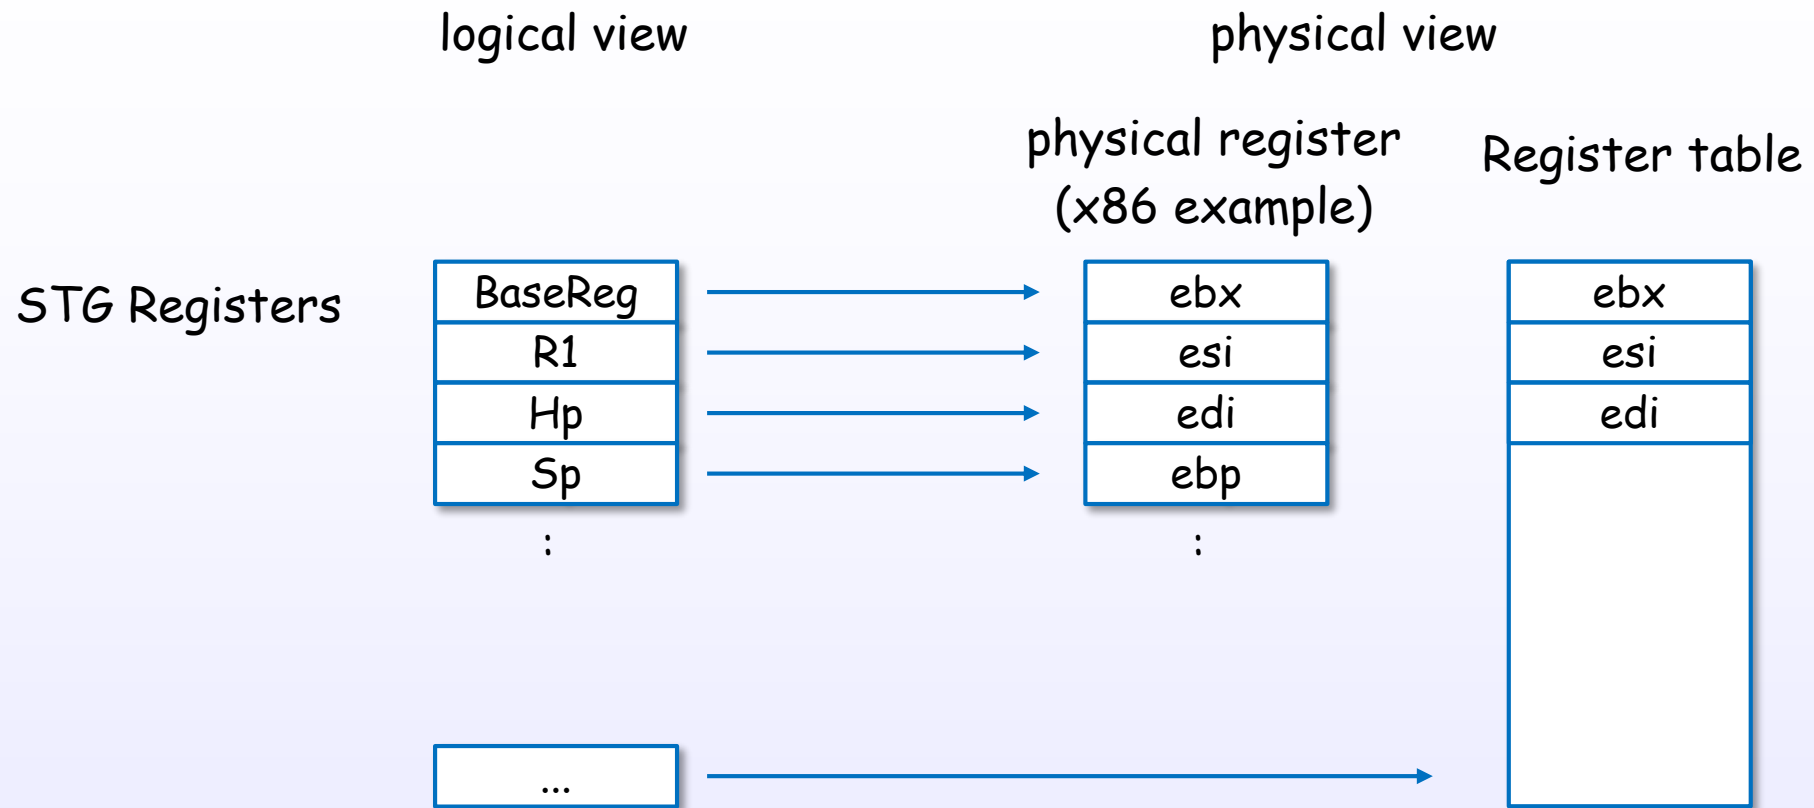
Each HEC (Capability) is initialized at initCapabilities [rts/Capability.c]

STG-machine

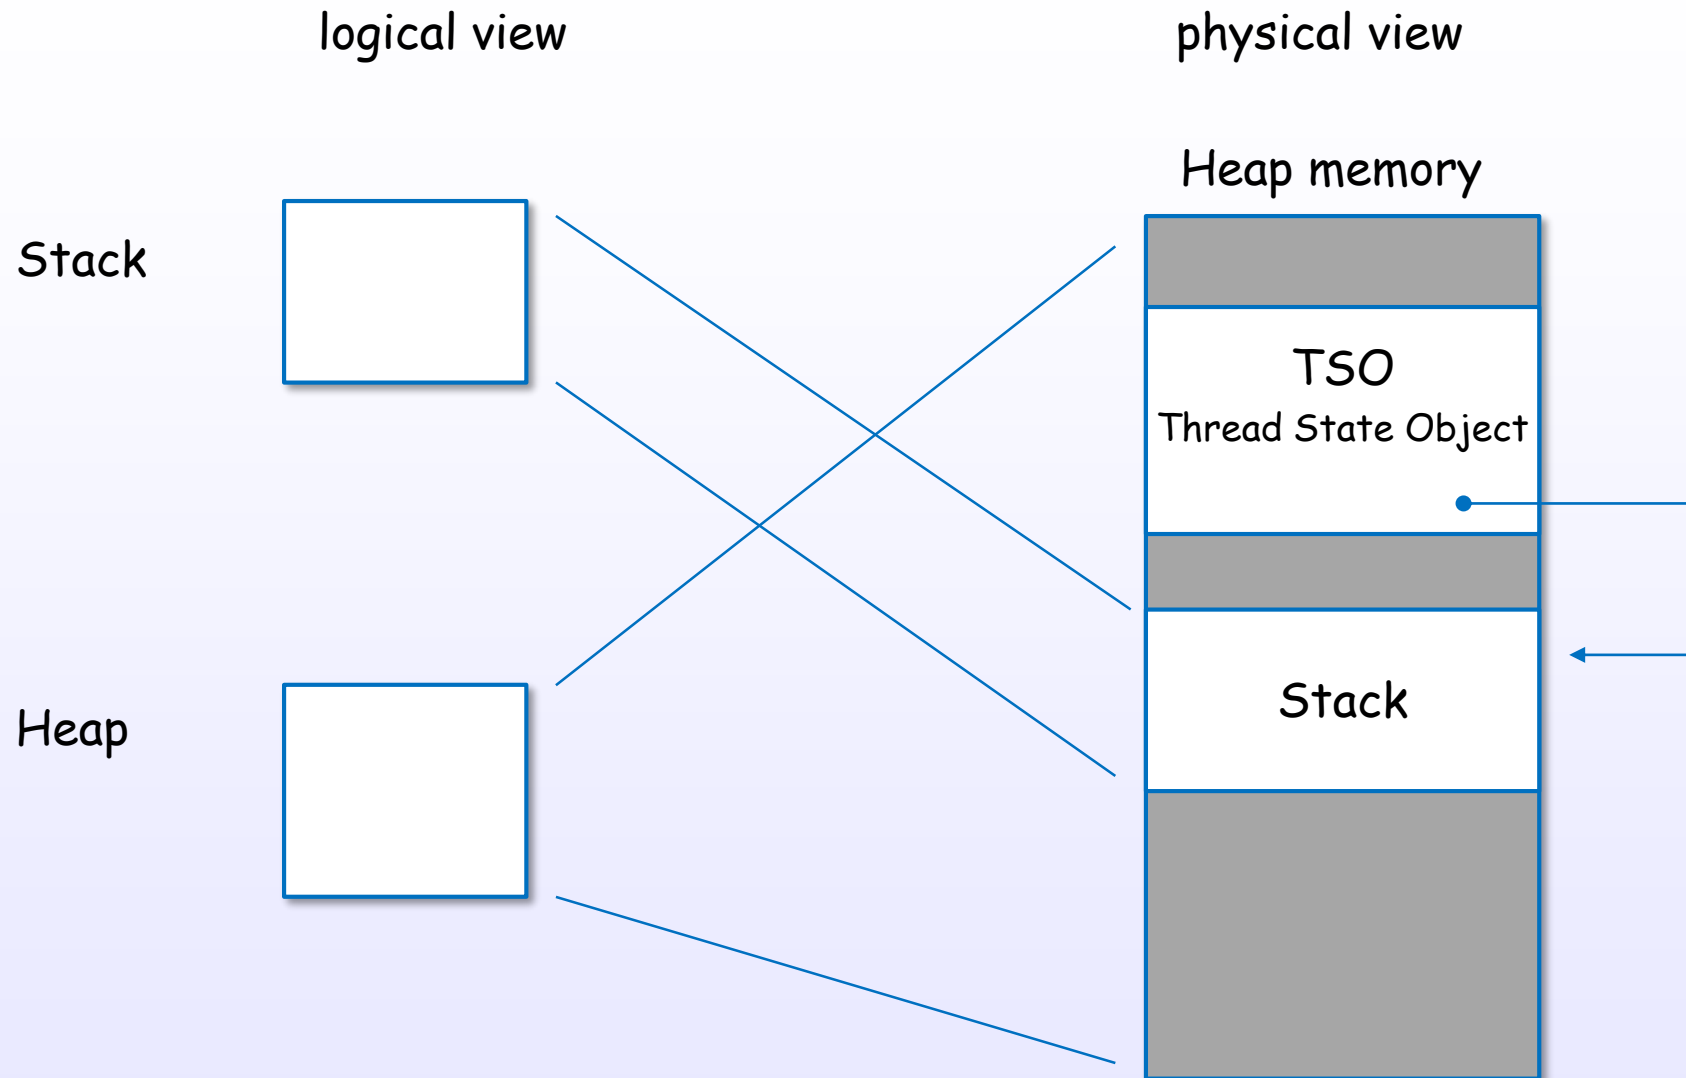
# The STG-machine consists of three parts



## STG-machine is mapped to physical processor



# STG-machine is mapped to physical processor



A stack and a TSO object are in the heap.  
The stack is stored separately from the TSO for size extension and GC.

# TSO data structure

[includes/rts/storage/TSO.h] (ghc 8.0)

```
typedef struct StgTSO_  
  StgHeader          header;  
  struct StgTSO_*     _link;  
  struct StgTSO_*     global_link;  
  struct StgStack_*   *stackobj;  
  StgWord16          what_next;  
  StgWord16          why_blocked;  
  StgWord32          flags;  
  StgTSOBlockInfo    block_info;  
  StgThreadID        id;  
  StgWord32          saved_errno;  
  StgWord32          dirty;  
  struct InCall_*     bound;  
  struct Capability_* cap;  
  struct StgTRecHeader_* trec;  
  struct MessageThrowTo_* blocked_exceptions;  
  struct StgBlockingQueue_* bq;  
  StgInt64          alloc_limit;  
  StgWord32          tot_stack_size;  
} *StgTSOPtr;
```


link to stack object

A TSO object is **only ~18words + stack**. Lightweight!



Heap objects in STG-machine


# Every heap object is represented uniformly




Closure (header + payload) + Info Table + Entry Code

# Heap object (closure)

logical view

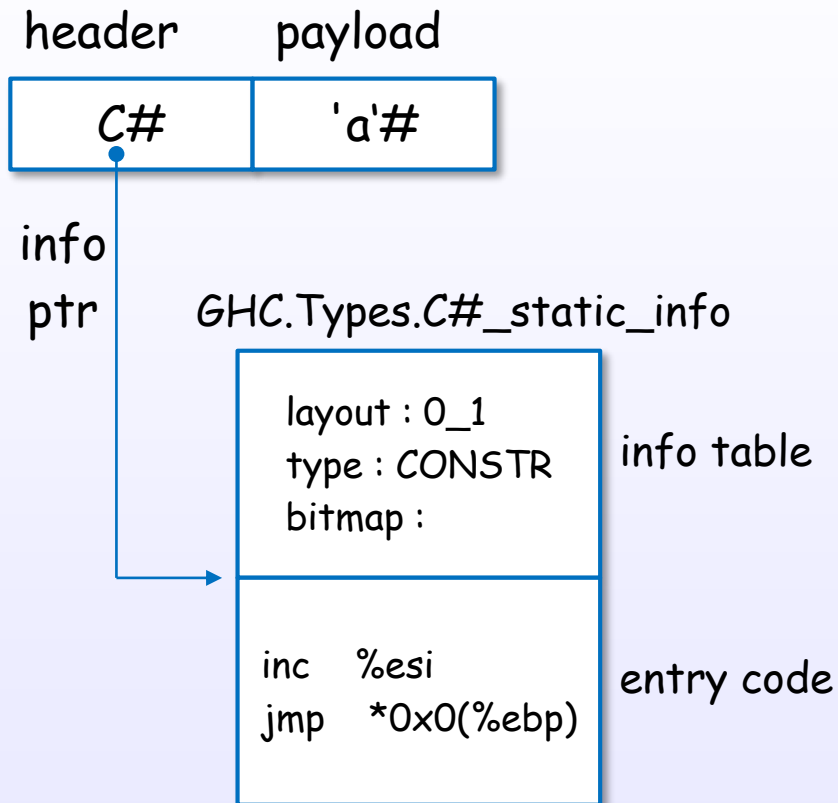


physical view

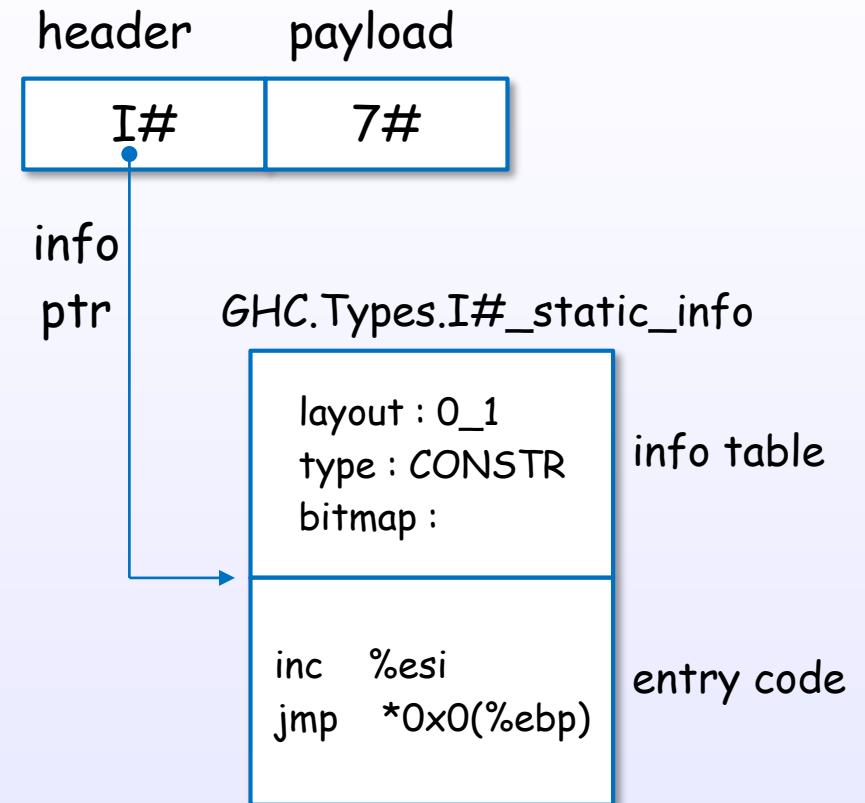


# Closure examples : Char, Int

'a' :: Char



7 :: Int



# Closure example (code)

[Example.hs]

```
module Example where
value1 :: Int
value1 = 7
```

STG

Cmm

[ghc -O -ddump-stg Example.hs]

```
Example.value1 :: GHC.Types.Int
[GblId, Caf=NoCafRefs, Str=DmdType m, Unf=OtherCon []] =
  NO_CCS GHC.Types.I#! [7#];
```

[ghc -O -ddump-opt-cmm Example.hs]

```
section ""data" .
__stginit_main@main:Example" {
  __stginit_main@main:Example:
}

section ""data" . Example.value1_closure" {
  Example.value1_closure:
  const GHC.Types.I#_static_info;
  const 7;
}

section ""readonly" . cHc_str" {
  cHc_str:
  I8[] [109,97,105,110]
}
```

asm

[ghc -O -ddump-asm Example.hs]

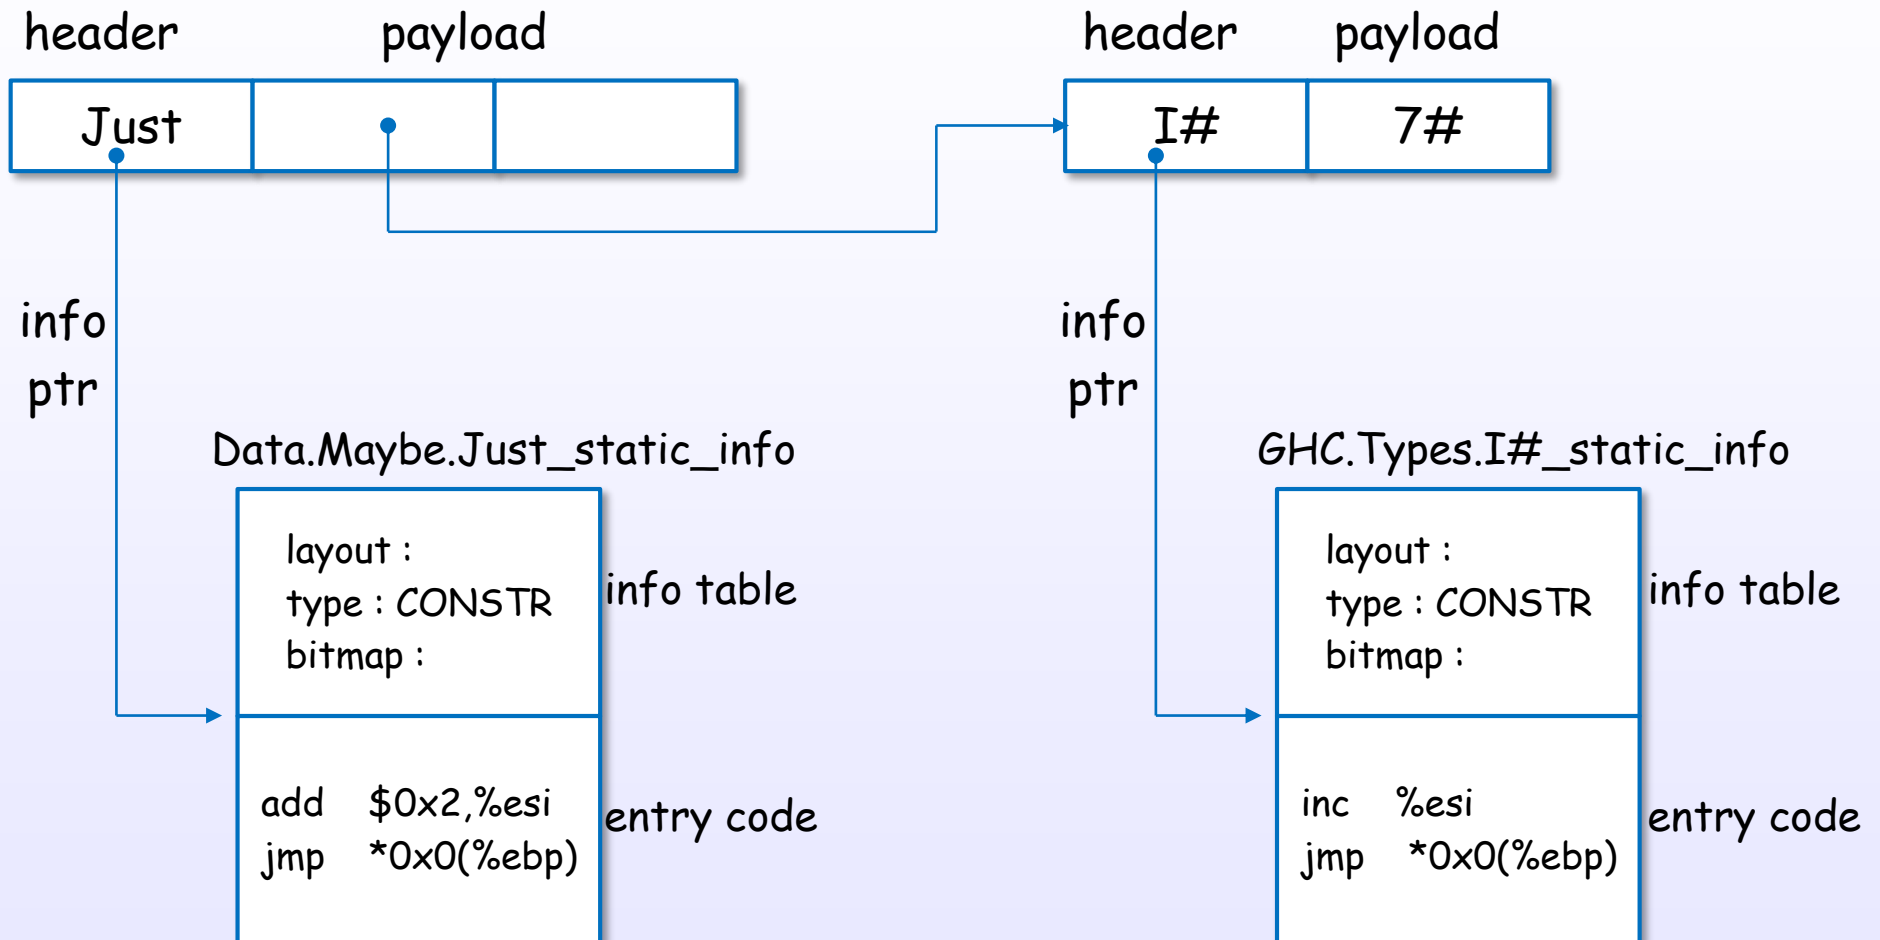
```
.section .data
.align 8
.align 1
.globl __stginit_main@main:Example
__stginit_main@main:Example:

.section .data
.align 8
.align 1
.globl Example.value1_closure
Example.value1_closure:
  .quad GHC.Types.I#_static_info
  .quad 7
```

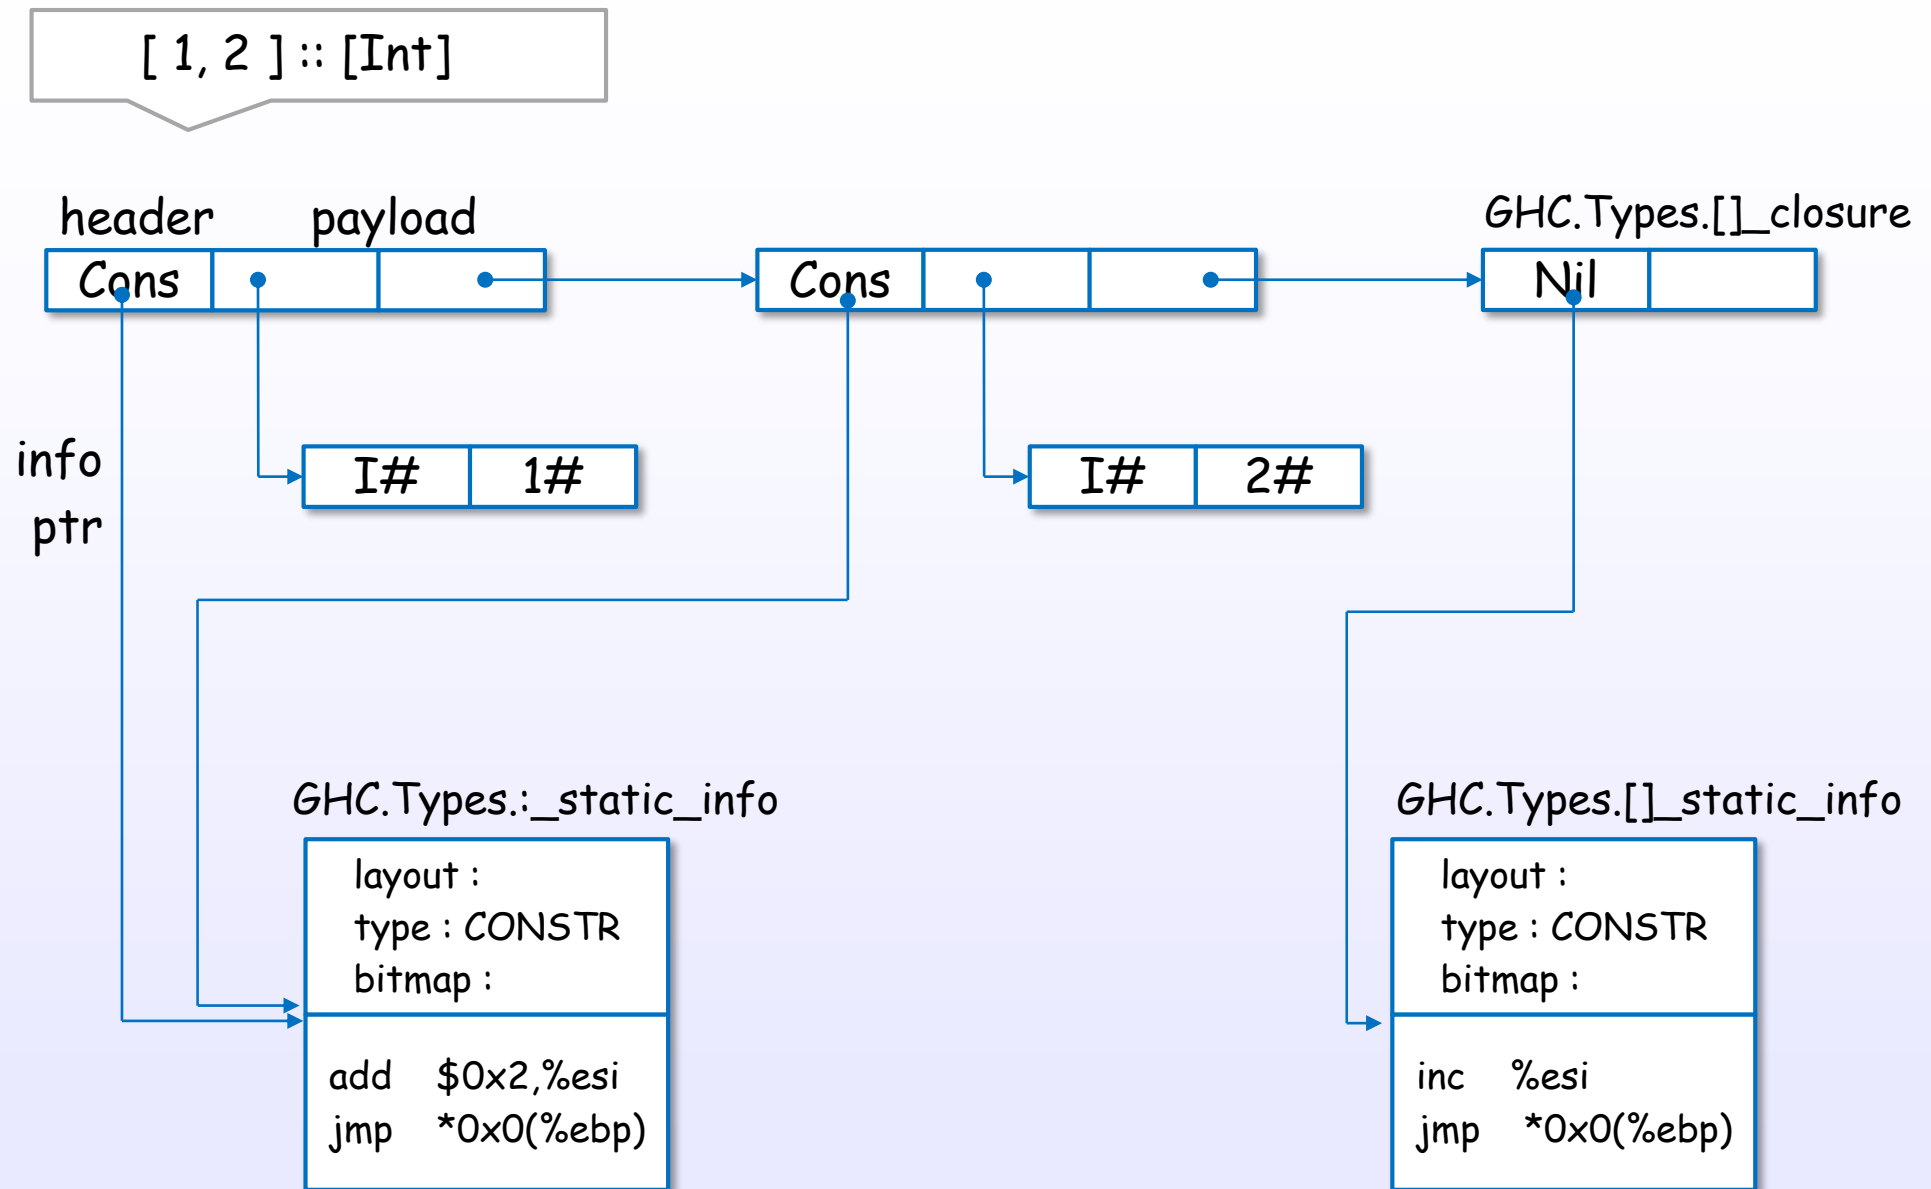
header		payload	
I#		7#	

# Closure examples : Maybe

Just 7 :: Maybe Int



# Closure examples : List




# Closure examples : Thunk


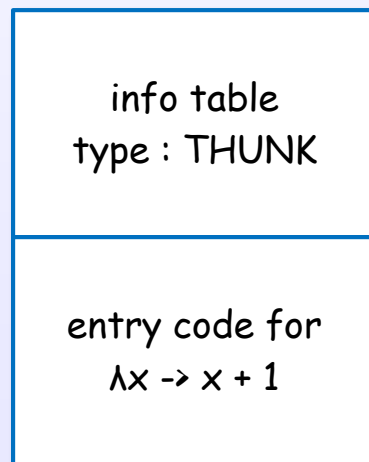
"thunk"

$x + 1 :: \text{Int}$


(free variable :  $x = 7$ )



info  
ptr




info  
ptr





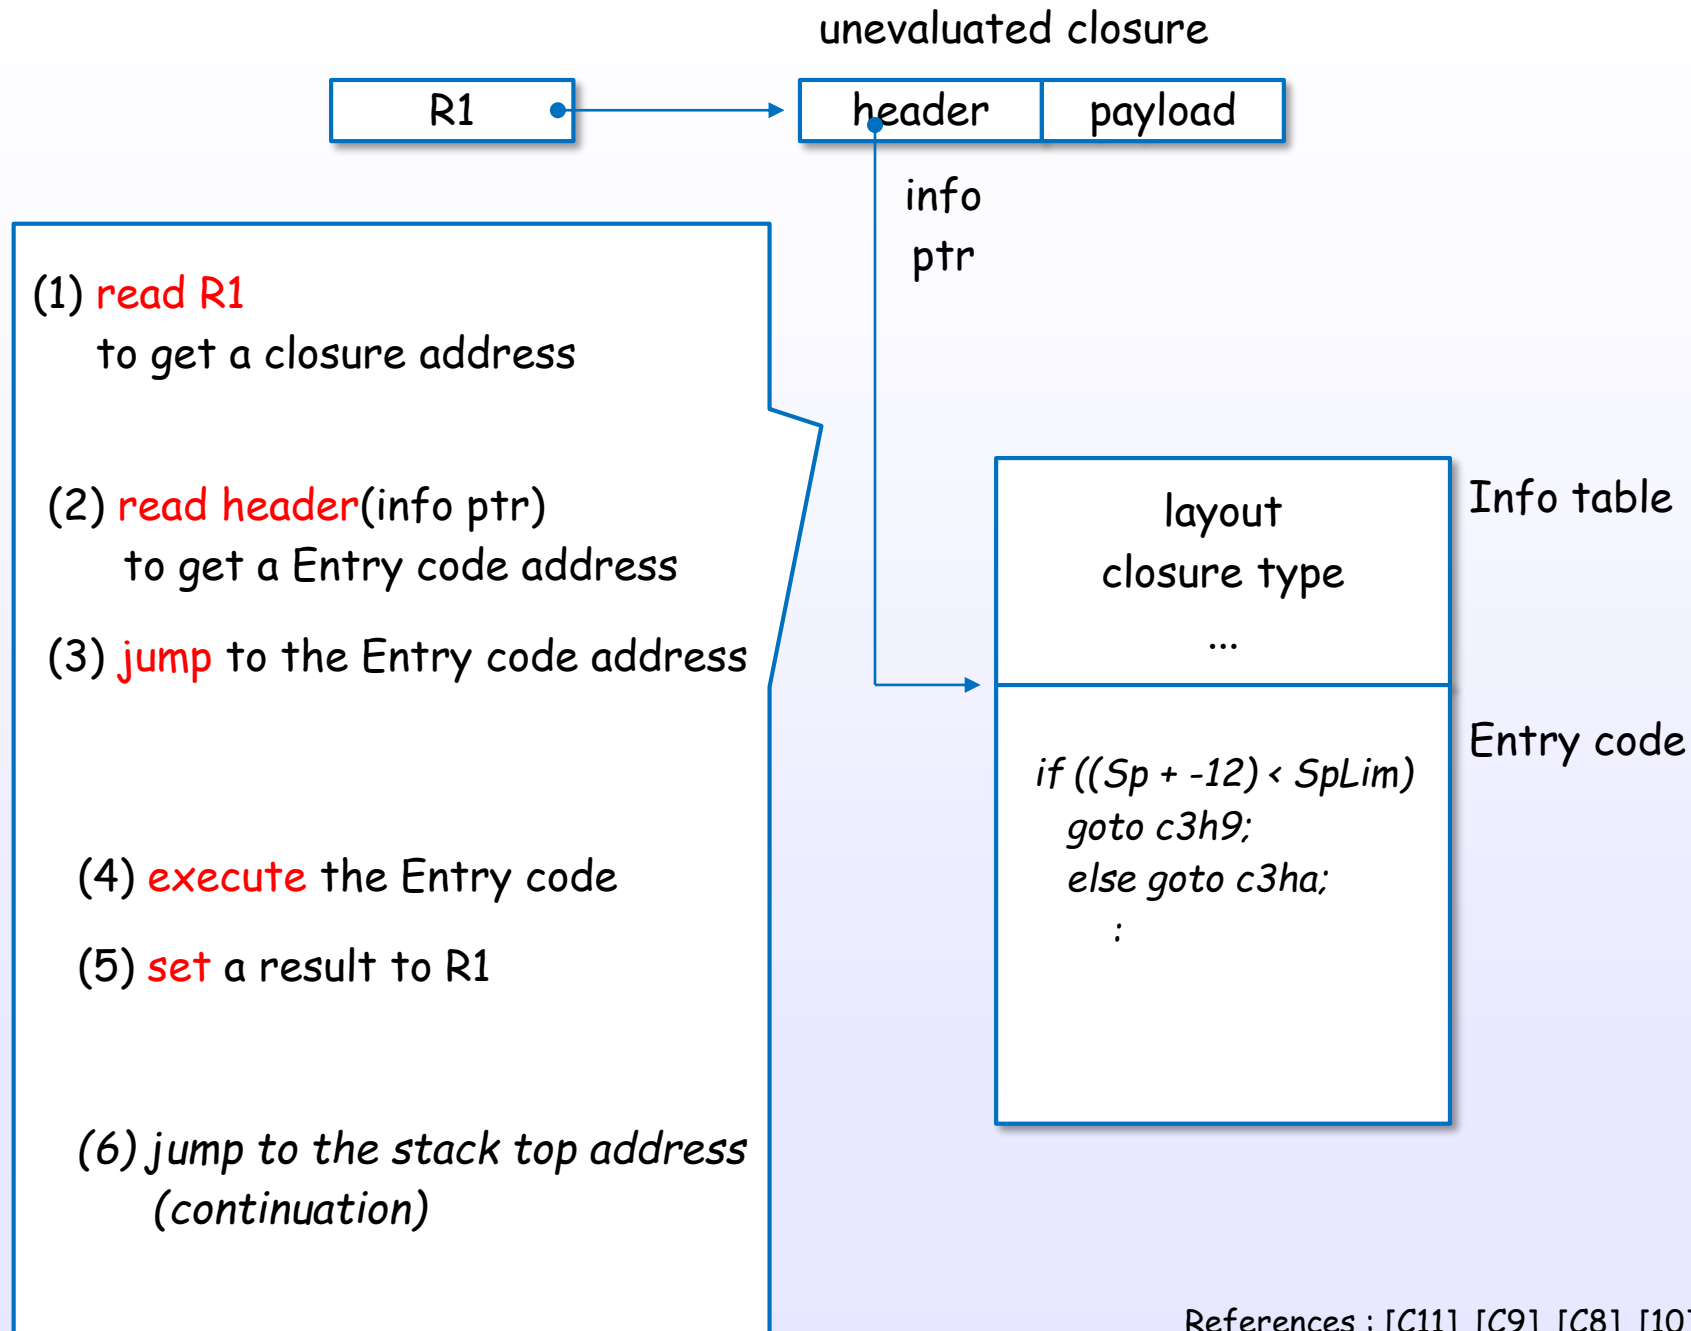
## STG-machine evaluation

# STG evaluation flow




- (1) push a continuation code (next code) to the stack top
- (2) enter to R1 closure
- (3) set a result to R1
- (4) jump (return) to the stack top code
- (5) repeat from (1)

# Enter to a closure



Pointer tagging

# Pointer tagging



pointer



... an unevaluated closure



... an evaluated closure;  
1st constructor value or evaluated.  
(for instance: "Nothing" )



... an evaluated closure; 2nd constructor value.  
(for instance: "Just xx")



... an evaluated closure; 3rd constructor value.

\* 32bit machine case



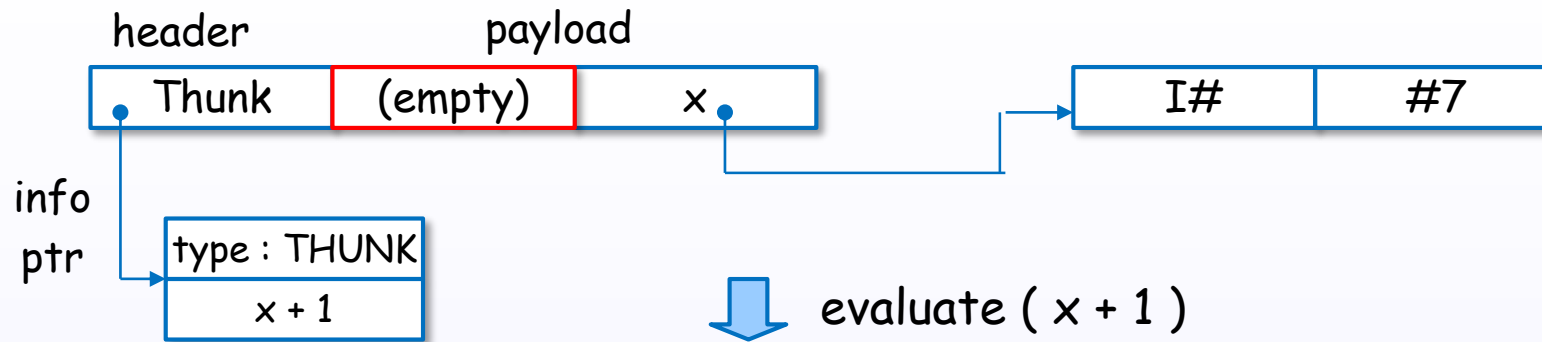
fast judgment!

check only pointer's lower bits without evaluating the closure.

Think and update

# Thunk and update

"thunk"  $x + 1 :: \text{Int}$  (free variable :  $x = 7$ )




GC (eliminate Indirect)


Allocate and free heap objects



# Allocate heap objects

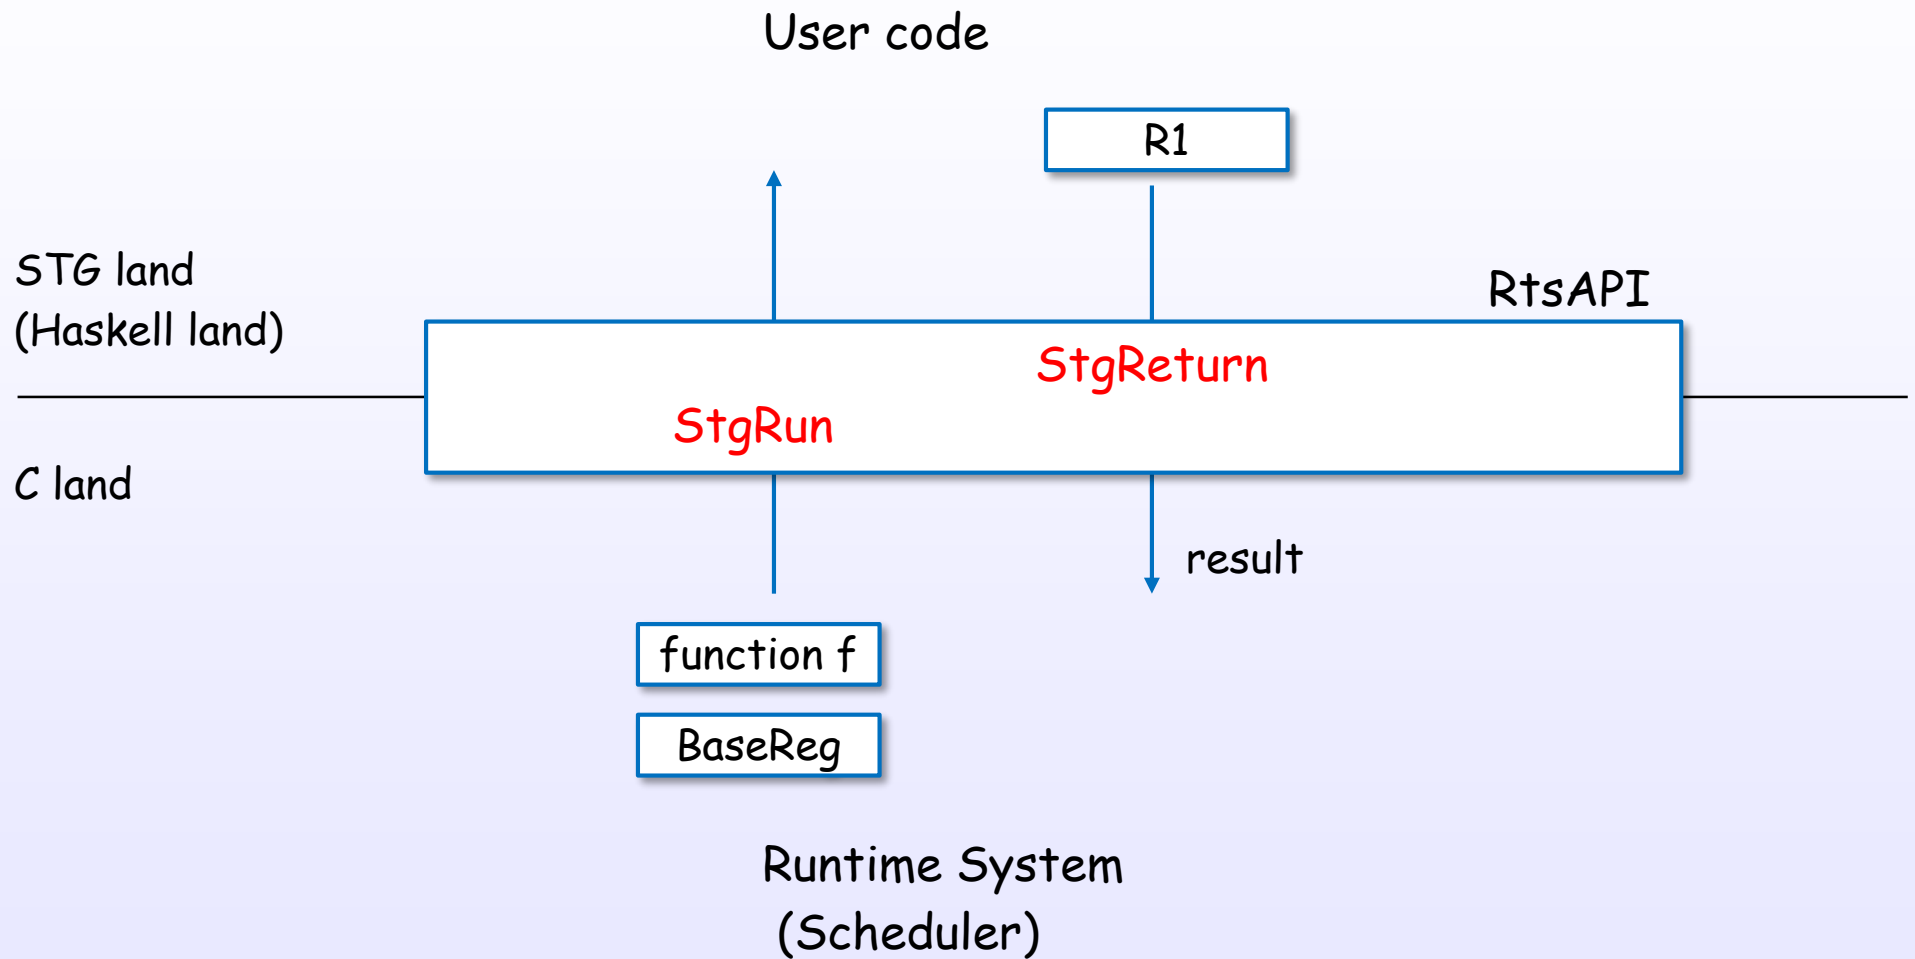


# free and collect heap objects



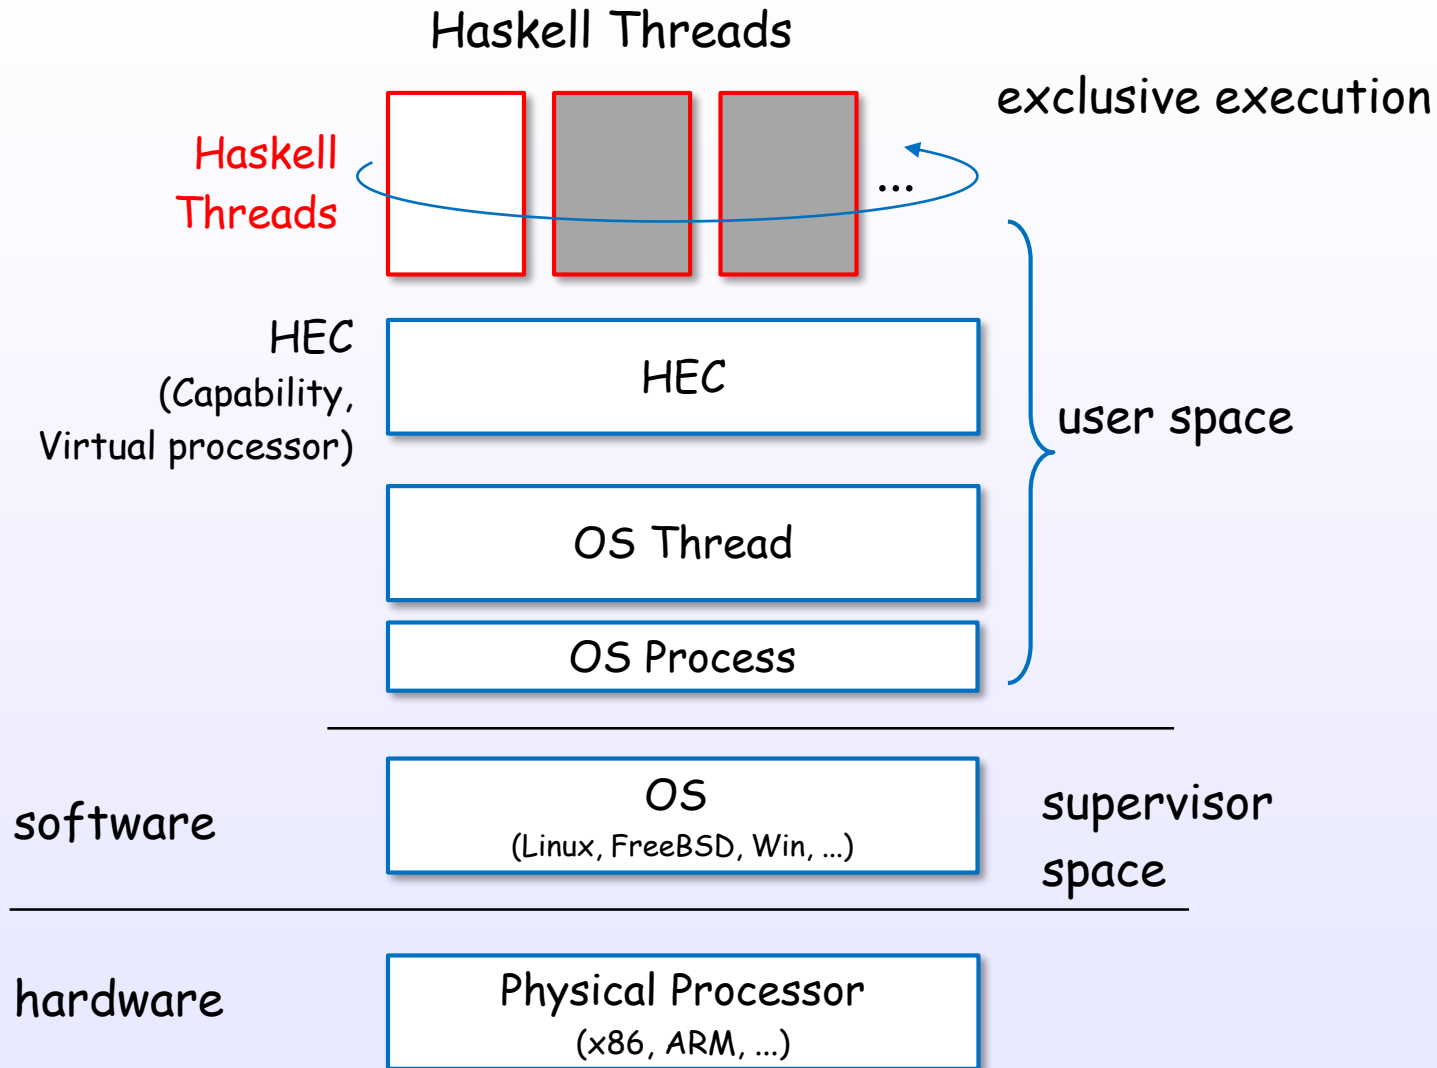
STG - C land interface

# STG (Haskell) land - C land interface

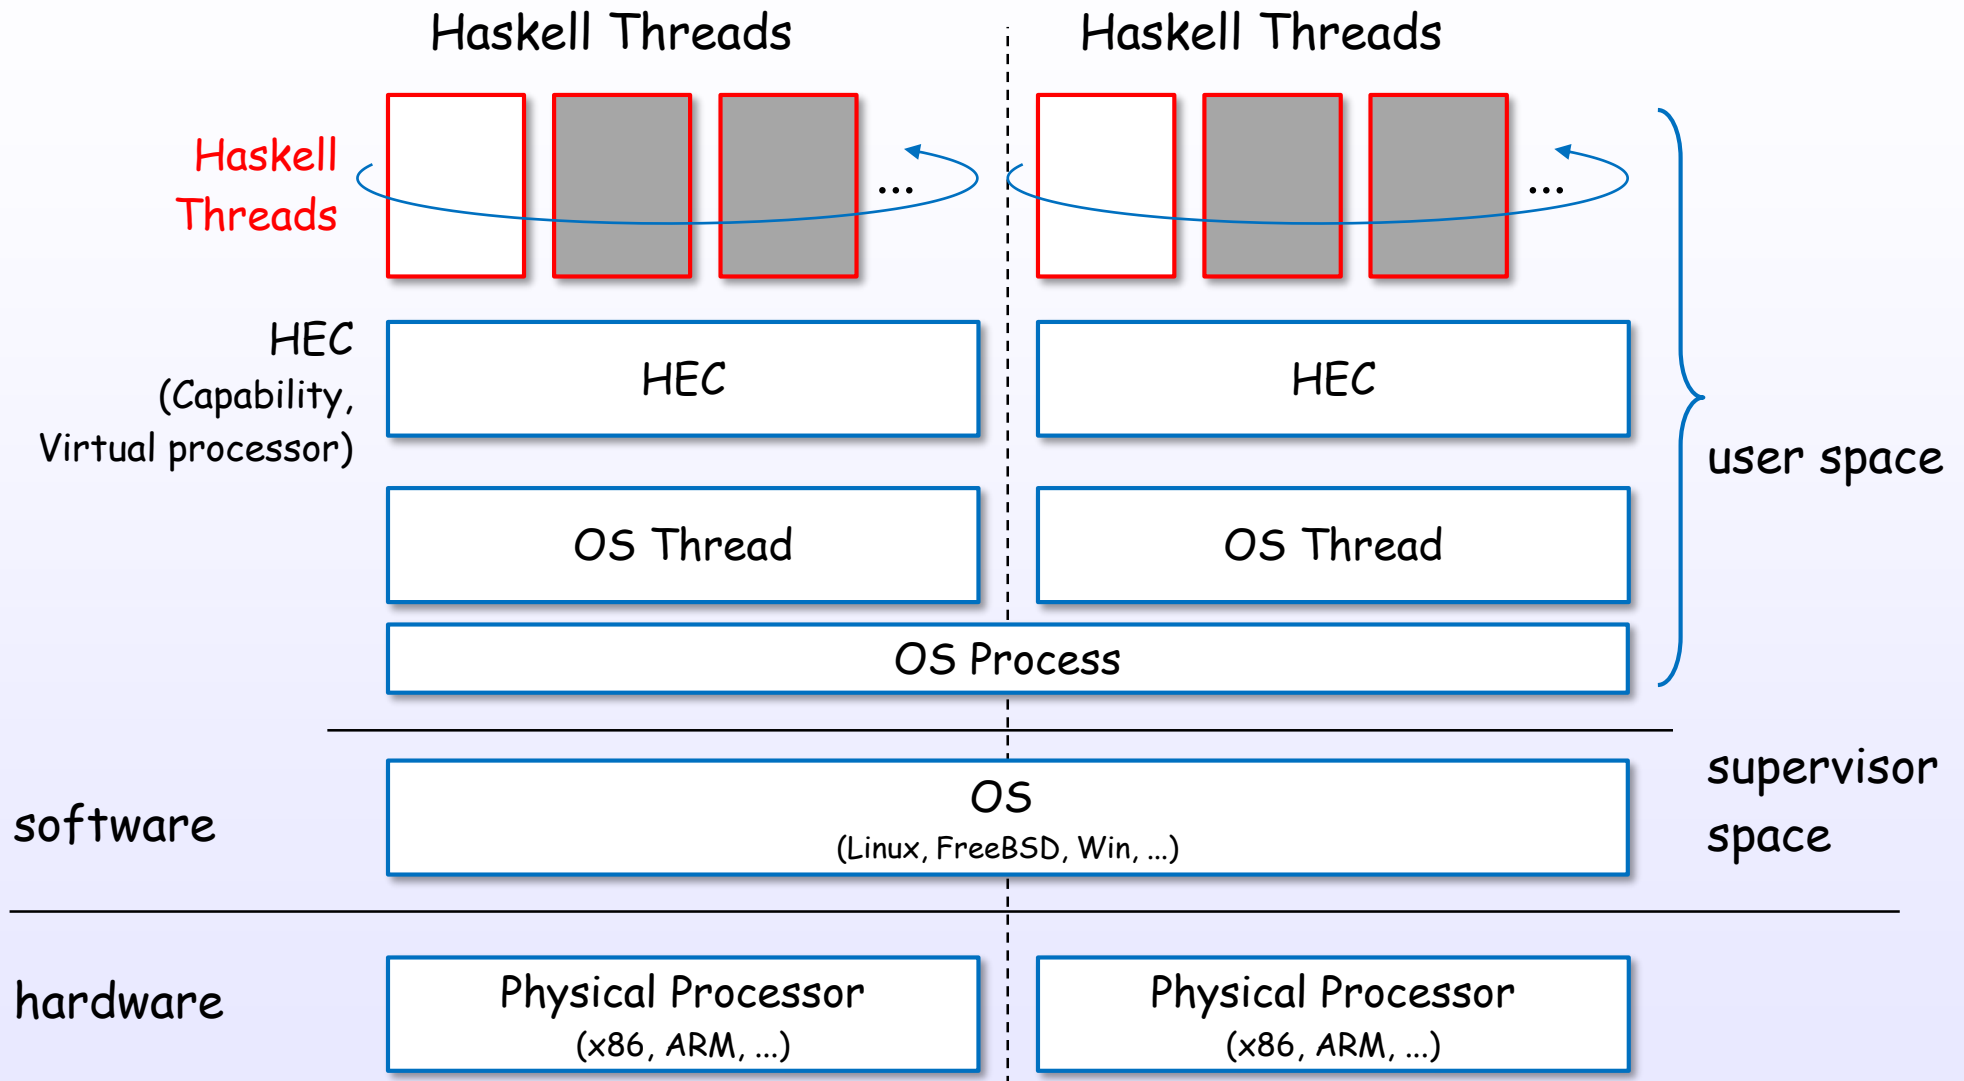


Thread

# Thread layer (single core)



# Thread layer (multi core)



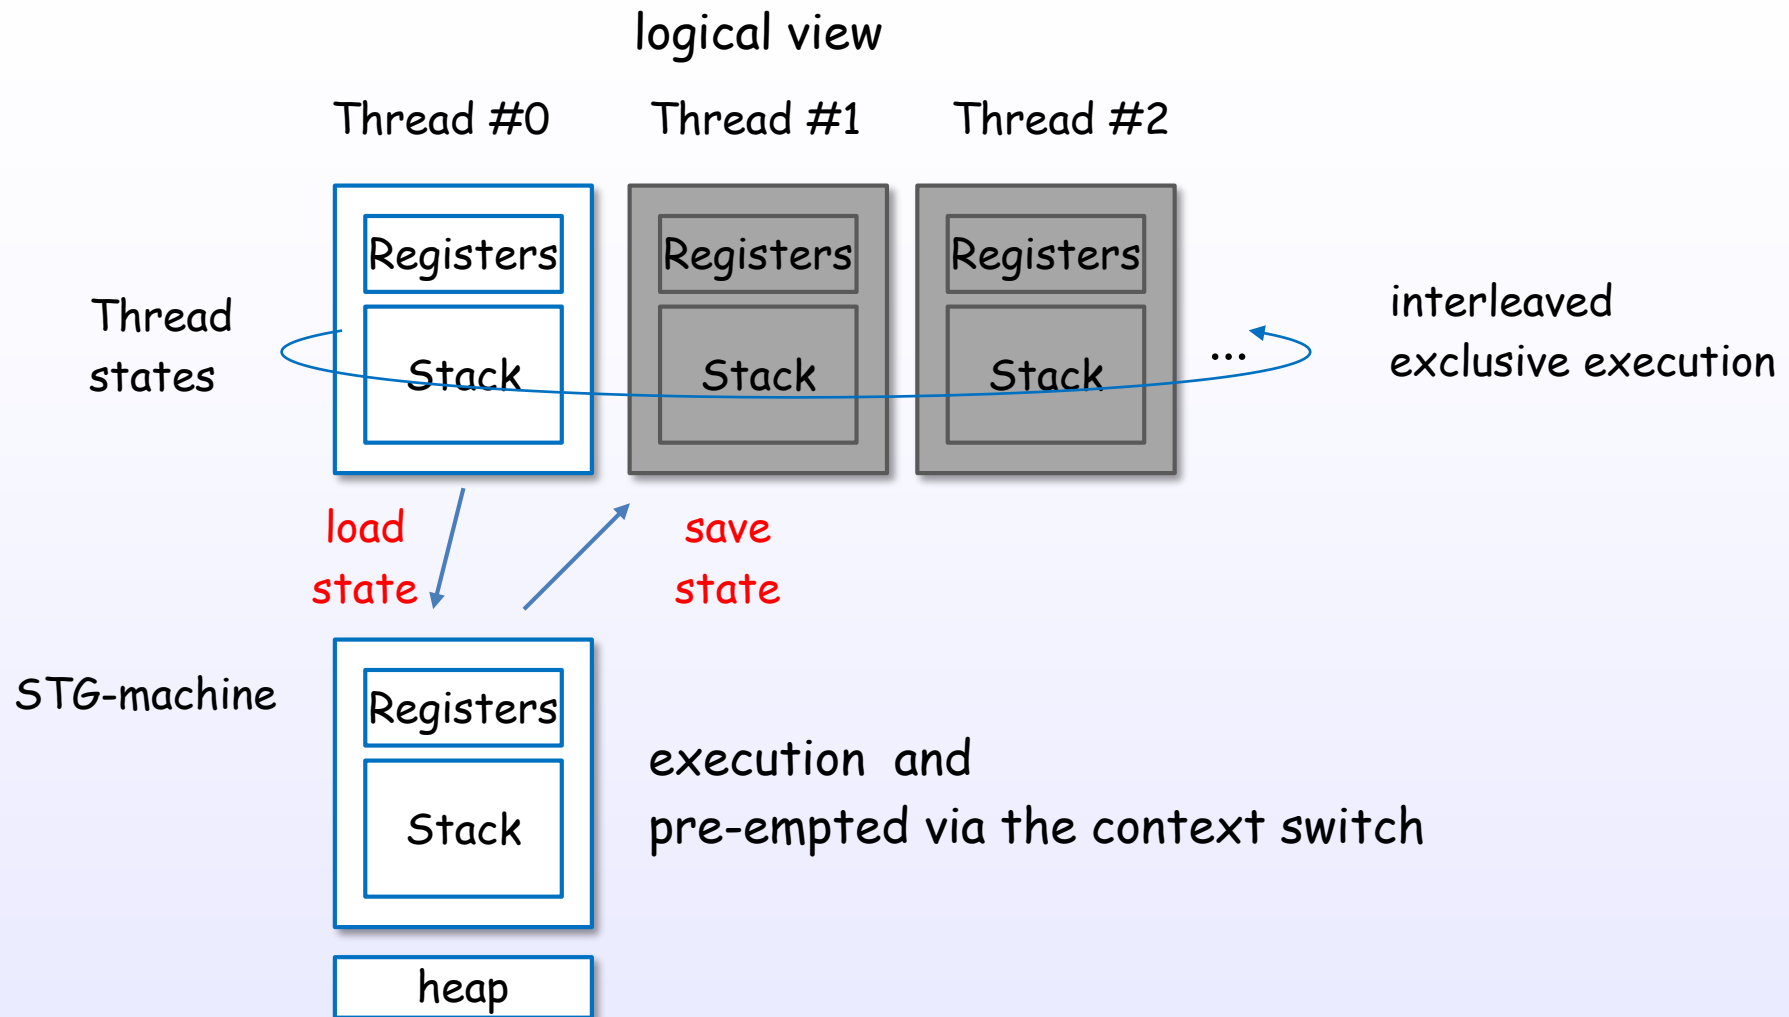
\*Threaded option case (ghc -threaded)

References : [5], [8], [9], [14], [C17], [C11], [19], [S17], [S16], [S23], [S22], [S14]

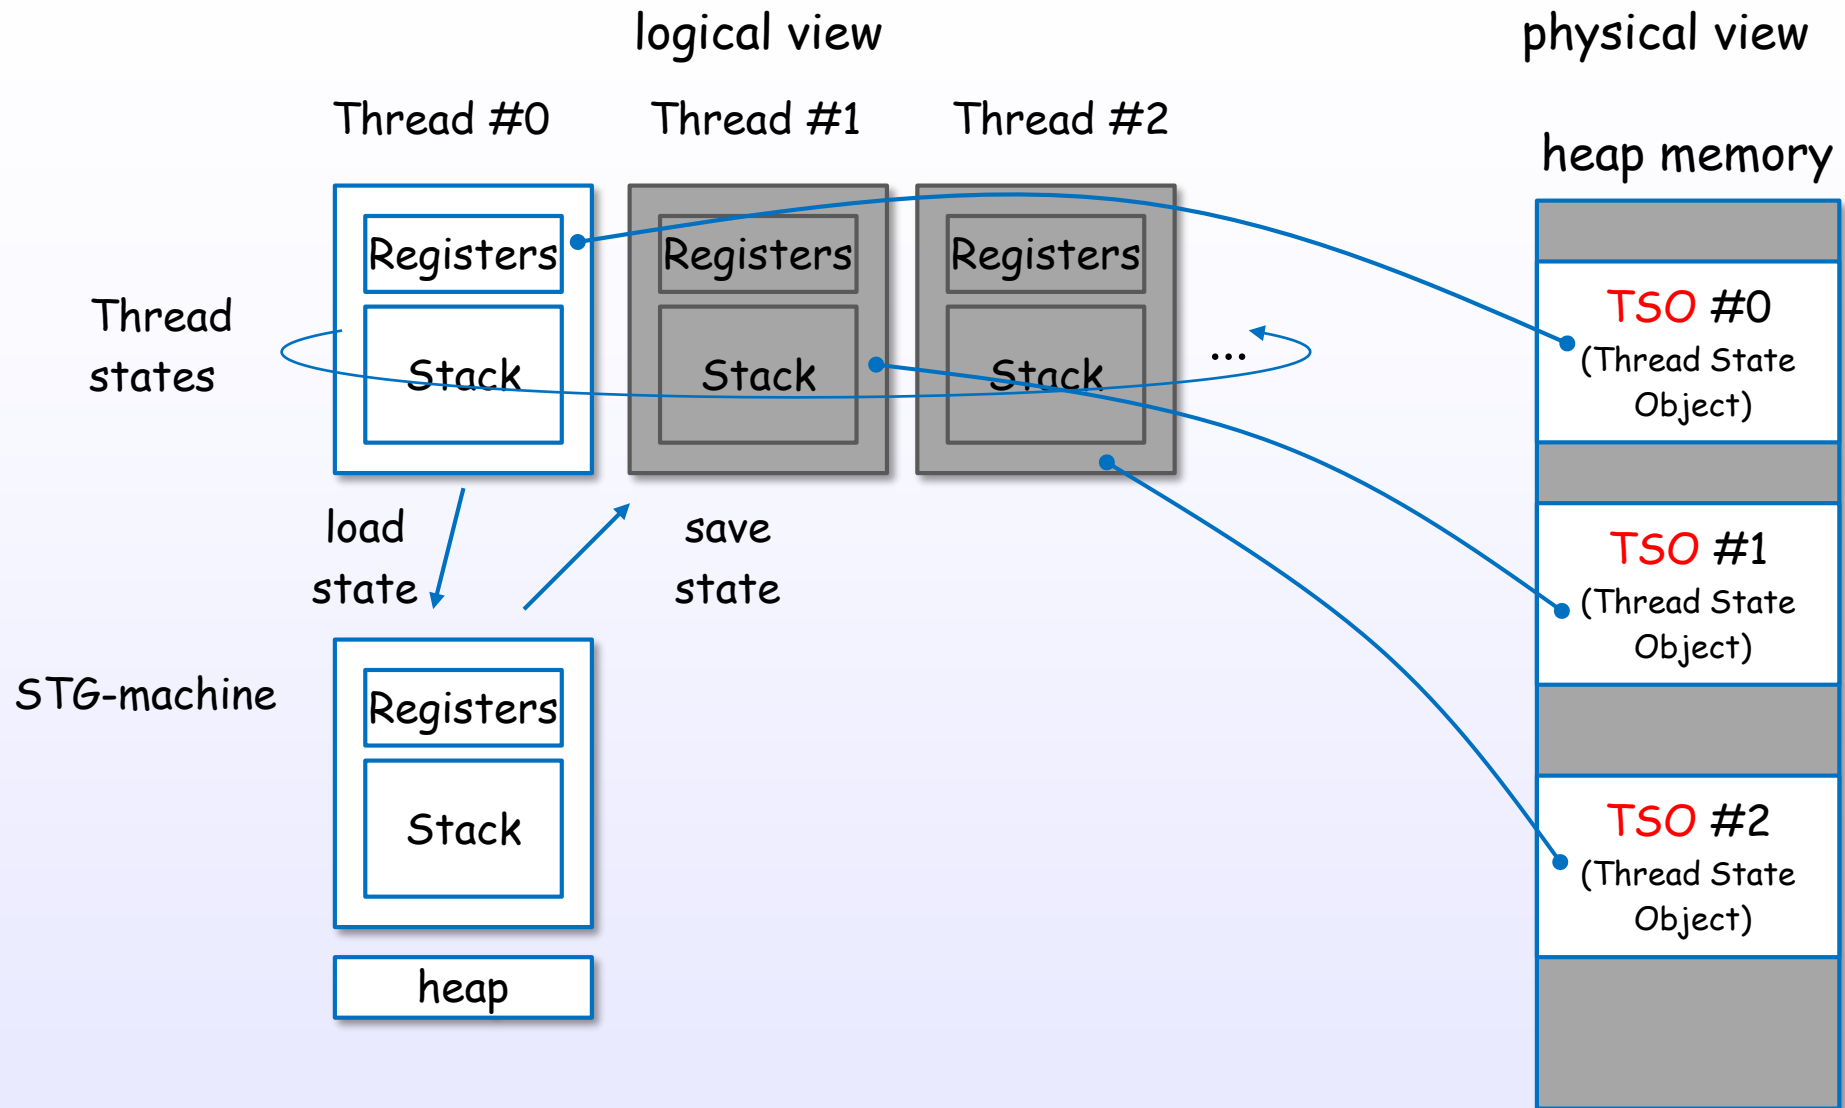
Thread context switch



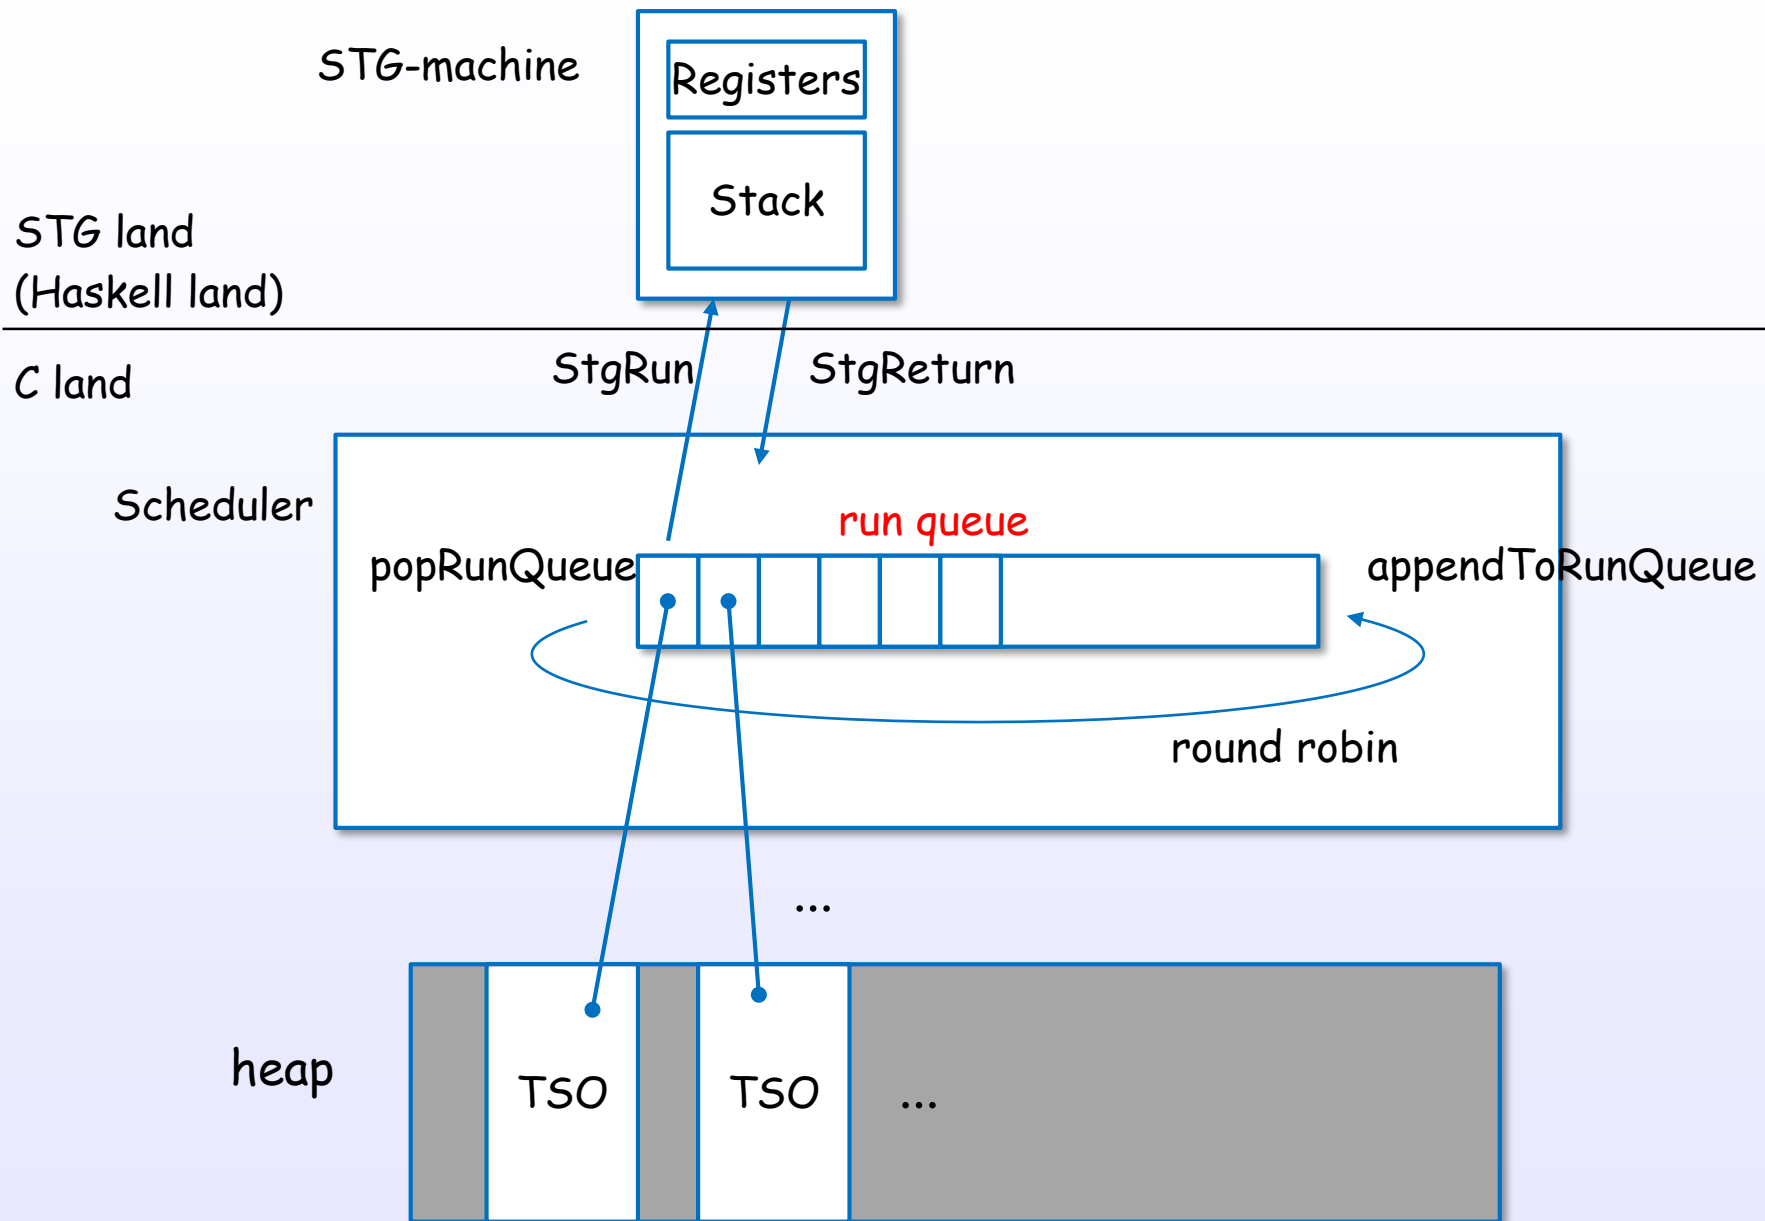
# Threads and context switch



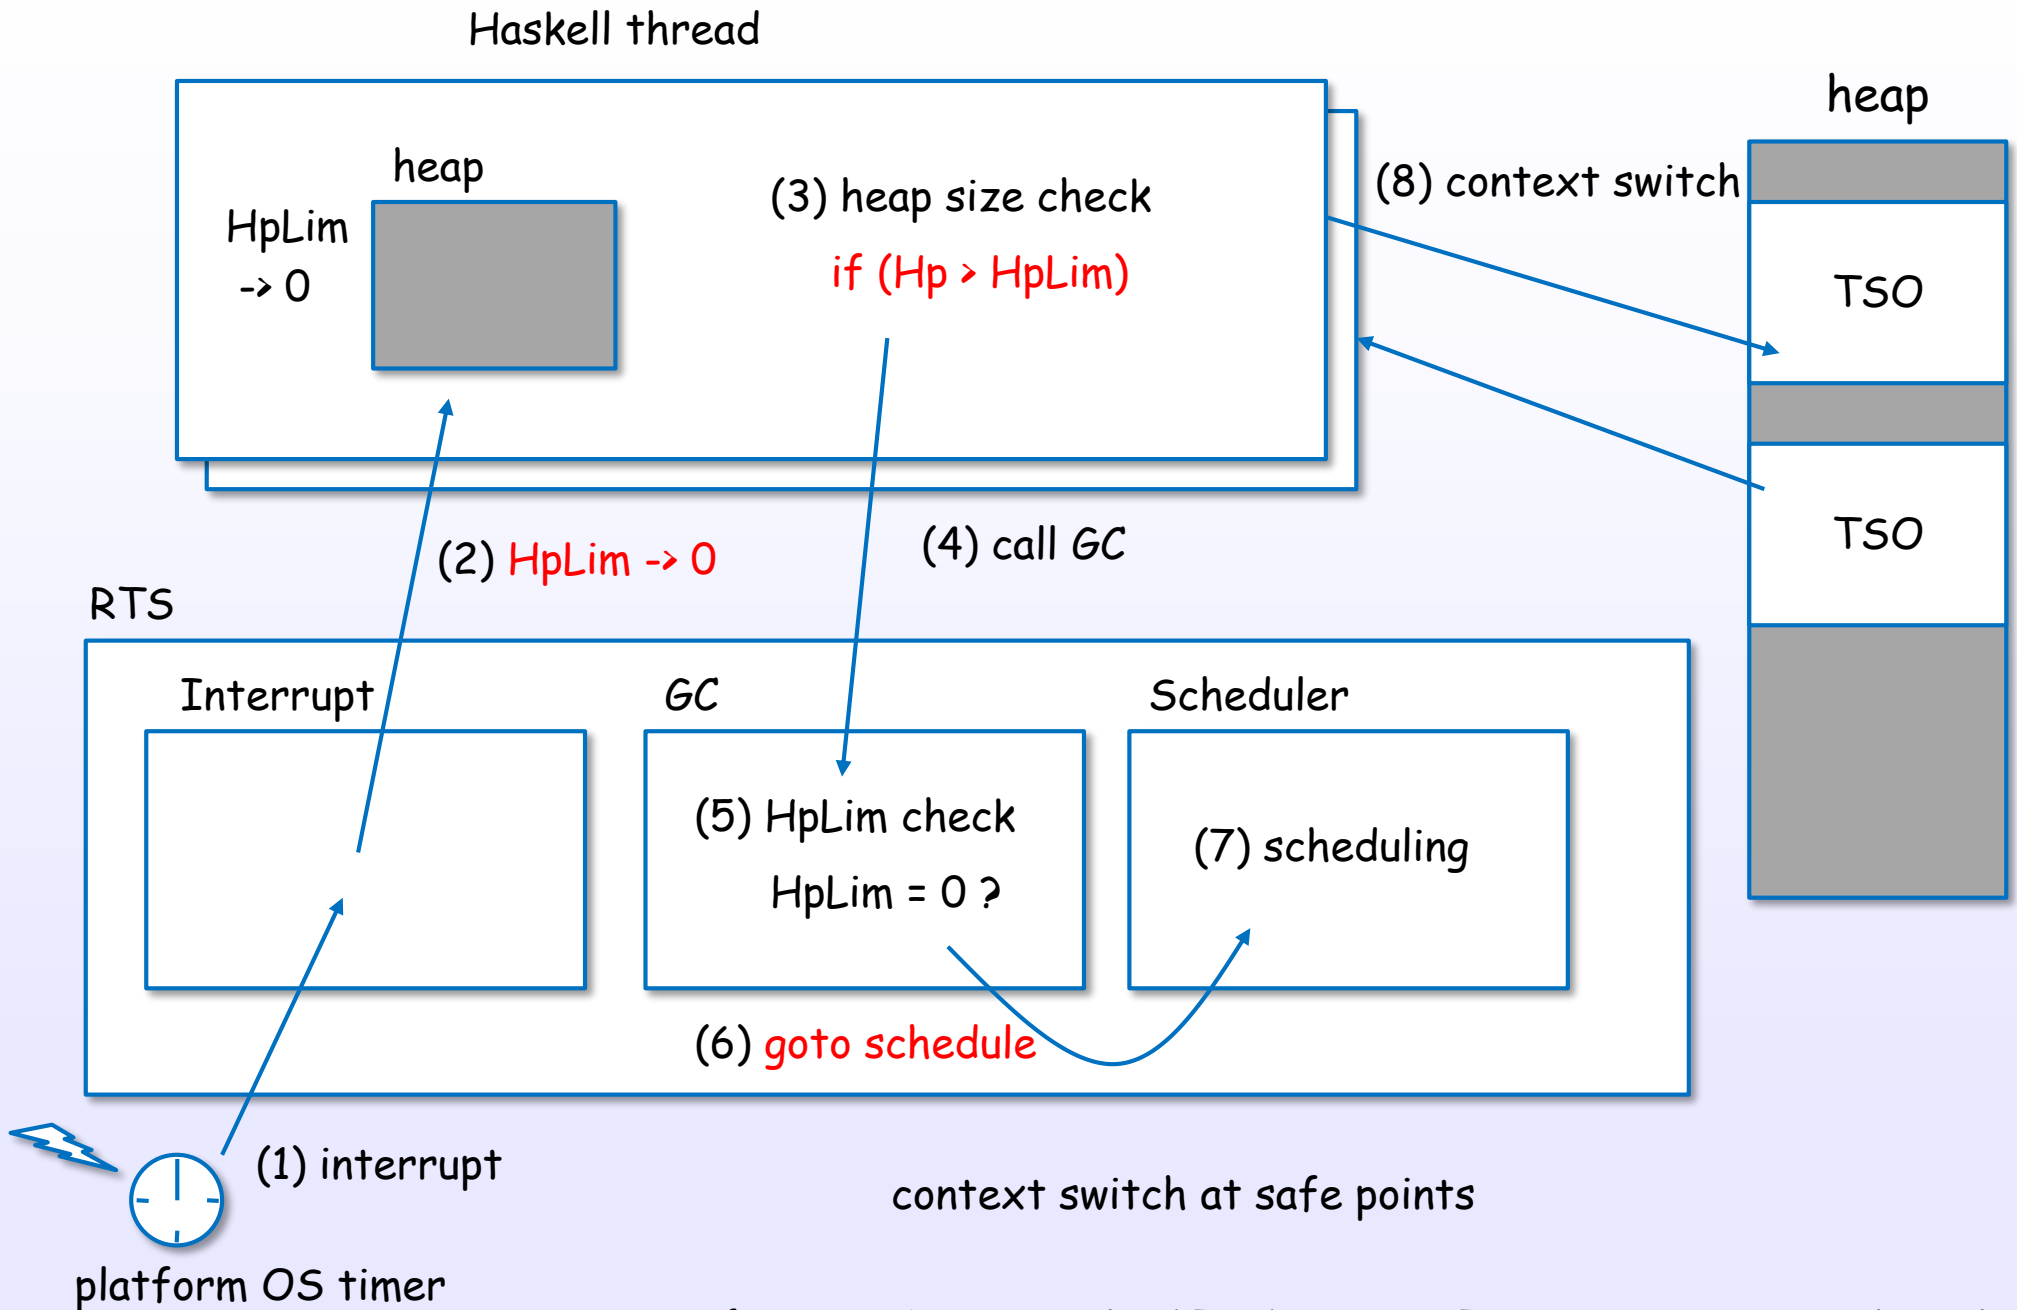
# Threads and TSOs



# Scheduling by run queue



# Context switch flow



# Context switch flow (code)

stg\_gc\_noregs

```
if (HpLim == 0) {  
    jump stg_returnToSched [R1];  
}
```

stg\_returnToSched

```
W_ r1;  
r1 = R1; // foreign calls may clobber R1  
SAVE_THREAD_STATE();  
foreign "C" threadPaused(MyCapability()  
    "ptr", CurrentTSO);  
R1 = r1;  
jump StgReturn [R1];
```

STG land  
(Haskell land)

C land

```
cap->r.rHpLim = NULL;
```

schedule

stopCapability  
contextSwitchCapability  
contextSwitchAllCapabilities  
handle\_tick

CreateTimerQueue

initTicker

initTimer

startTimer

hs\_init\_ghc

hs\_main

next  
handle\_tick ..

OS

\*Windows case

References : [5], [8], [9], [14], [C17], [C11], [19], [S17], [S16], [S21], [S23], [S22], [S14], [S24]

Creating main and sub threads

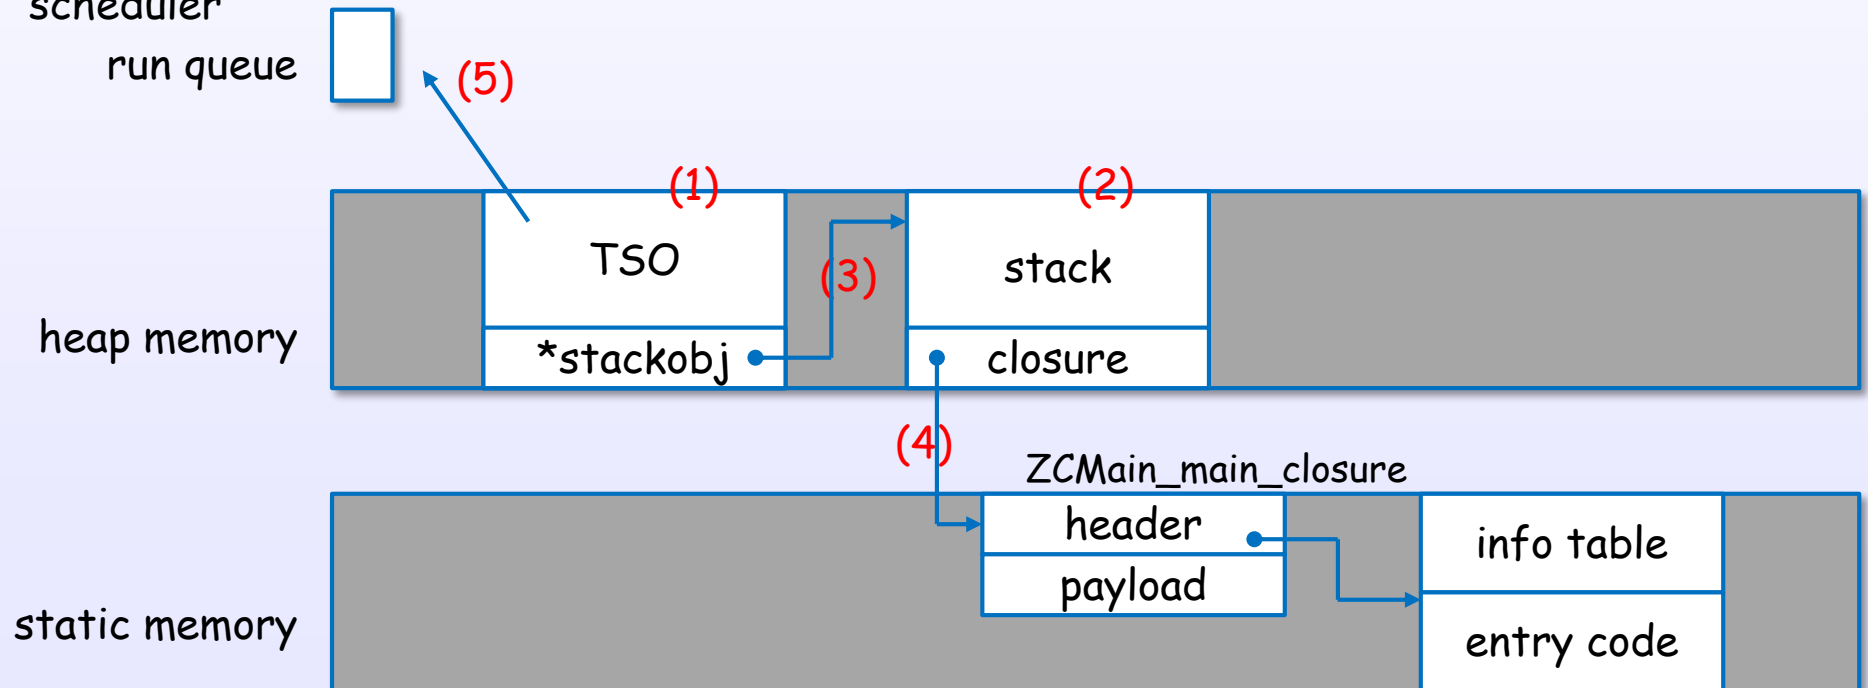
# Create a main thread

Runtime  
System

Runtime system bootstrap code [rts/RtsAPI.c]

```
rts_evalLazyIO
  createIOThread
    createThread ... (1), (2), (3)
    pushClosure ... (4)
  scheduleWaitThread
    appendToRunQueue ... (5)
```

scheduler  
run queue



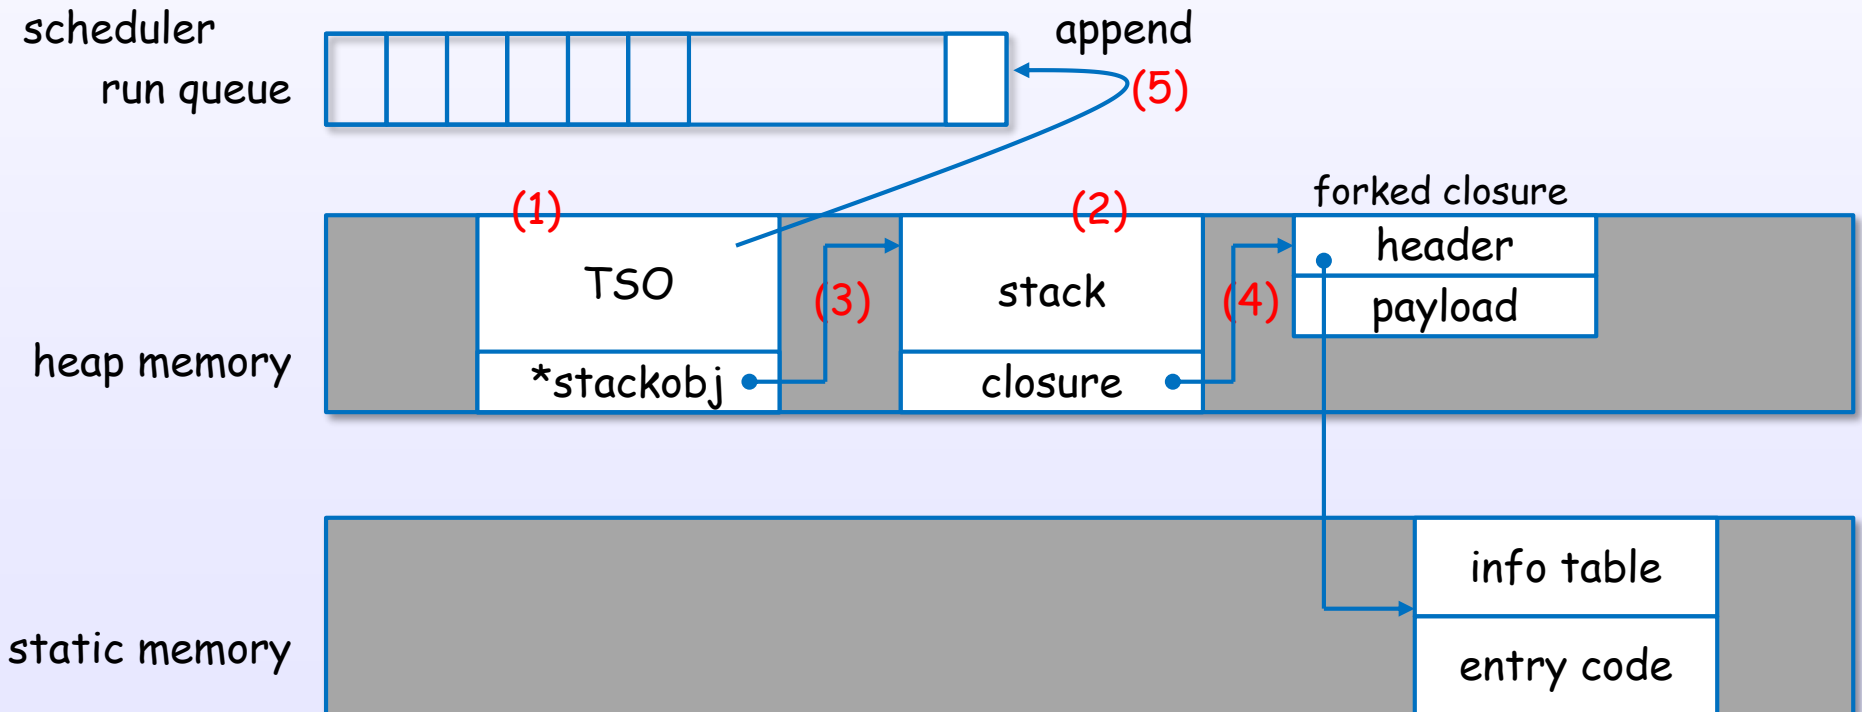
# Create a sub thread using forkIO

## Haskell Threads

```
forkIO
  stg_forkzh
    ccall createIOThread ... (1), (2), (3), (4)
    ccall scheduleThread ... (5)
```

User code

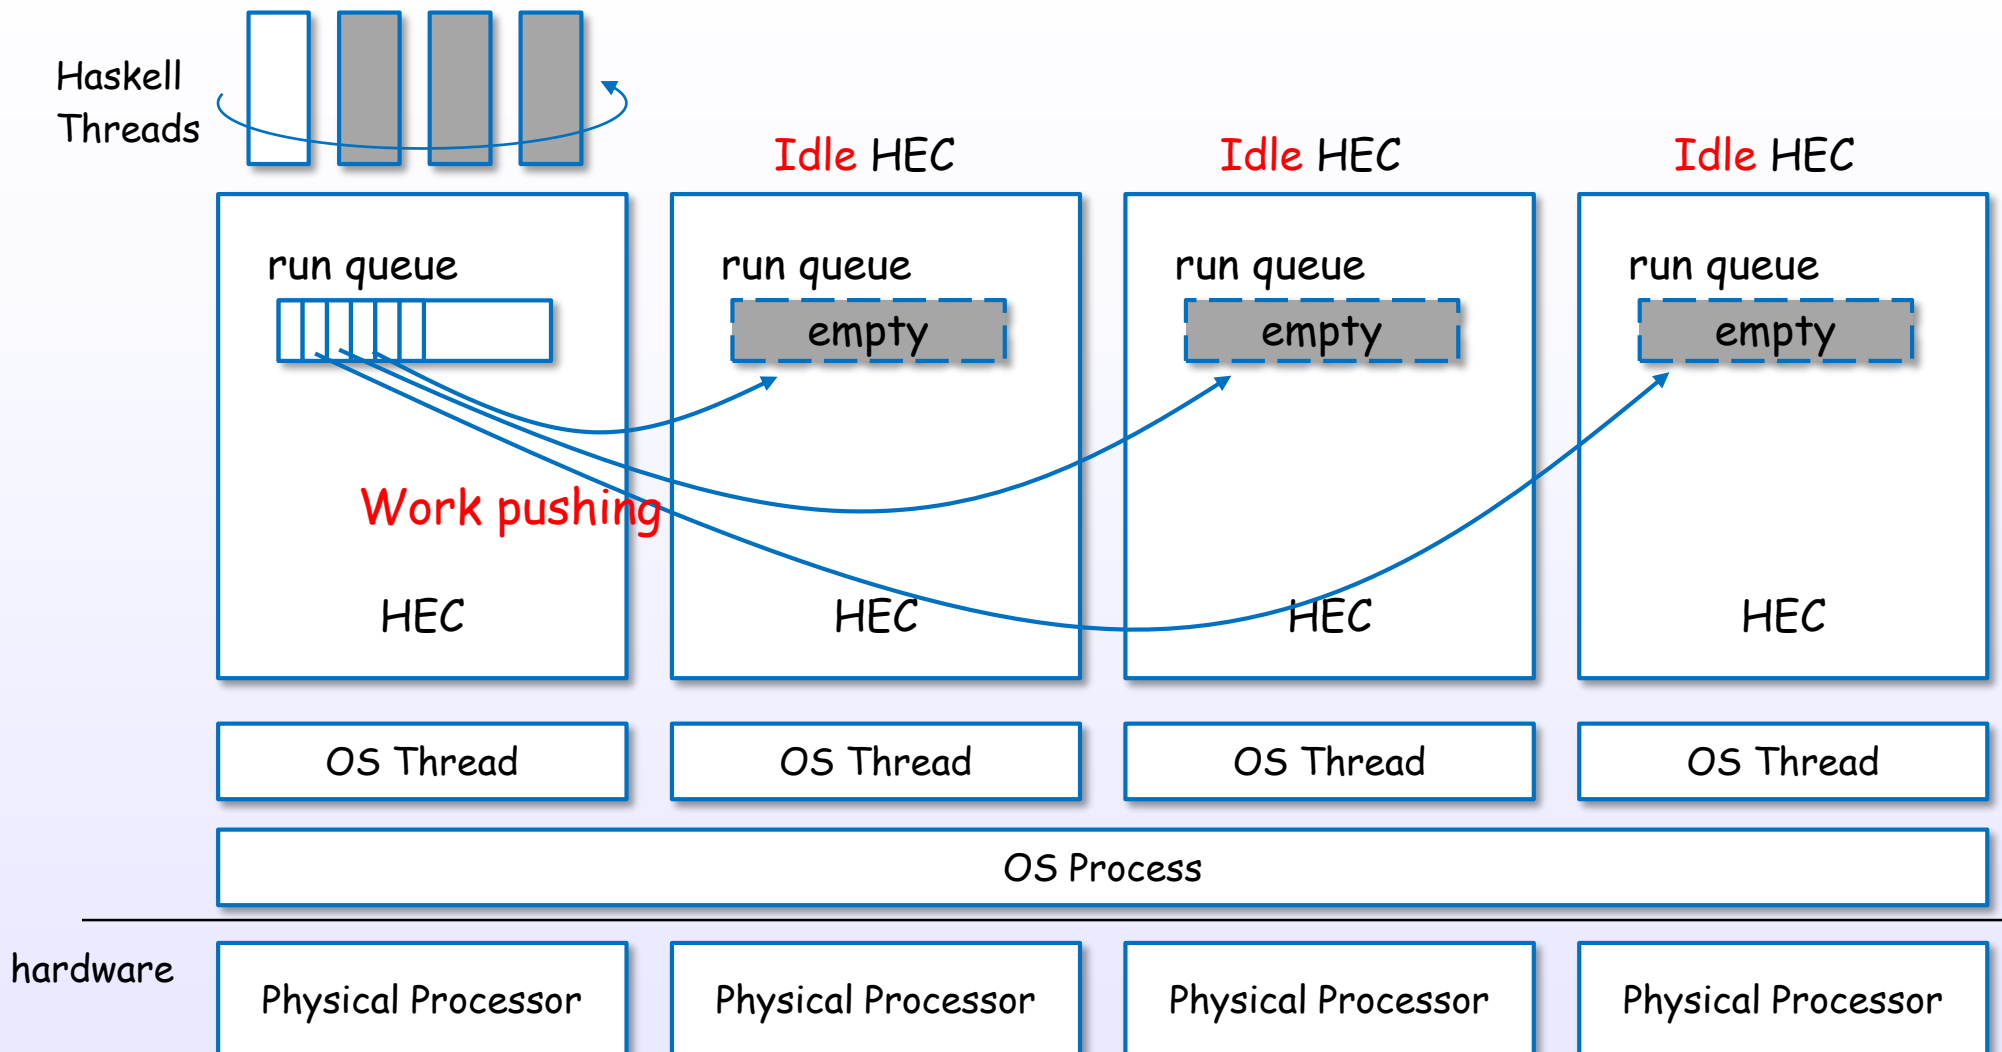
Runtime System





# Thread migration

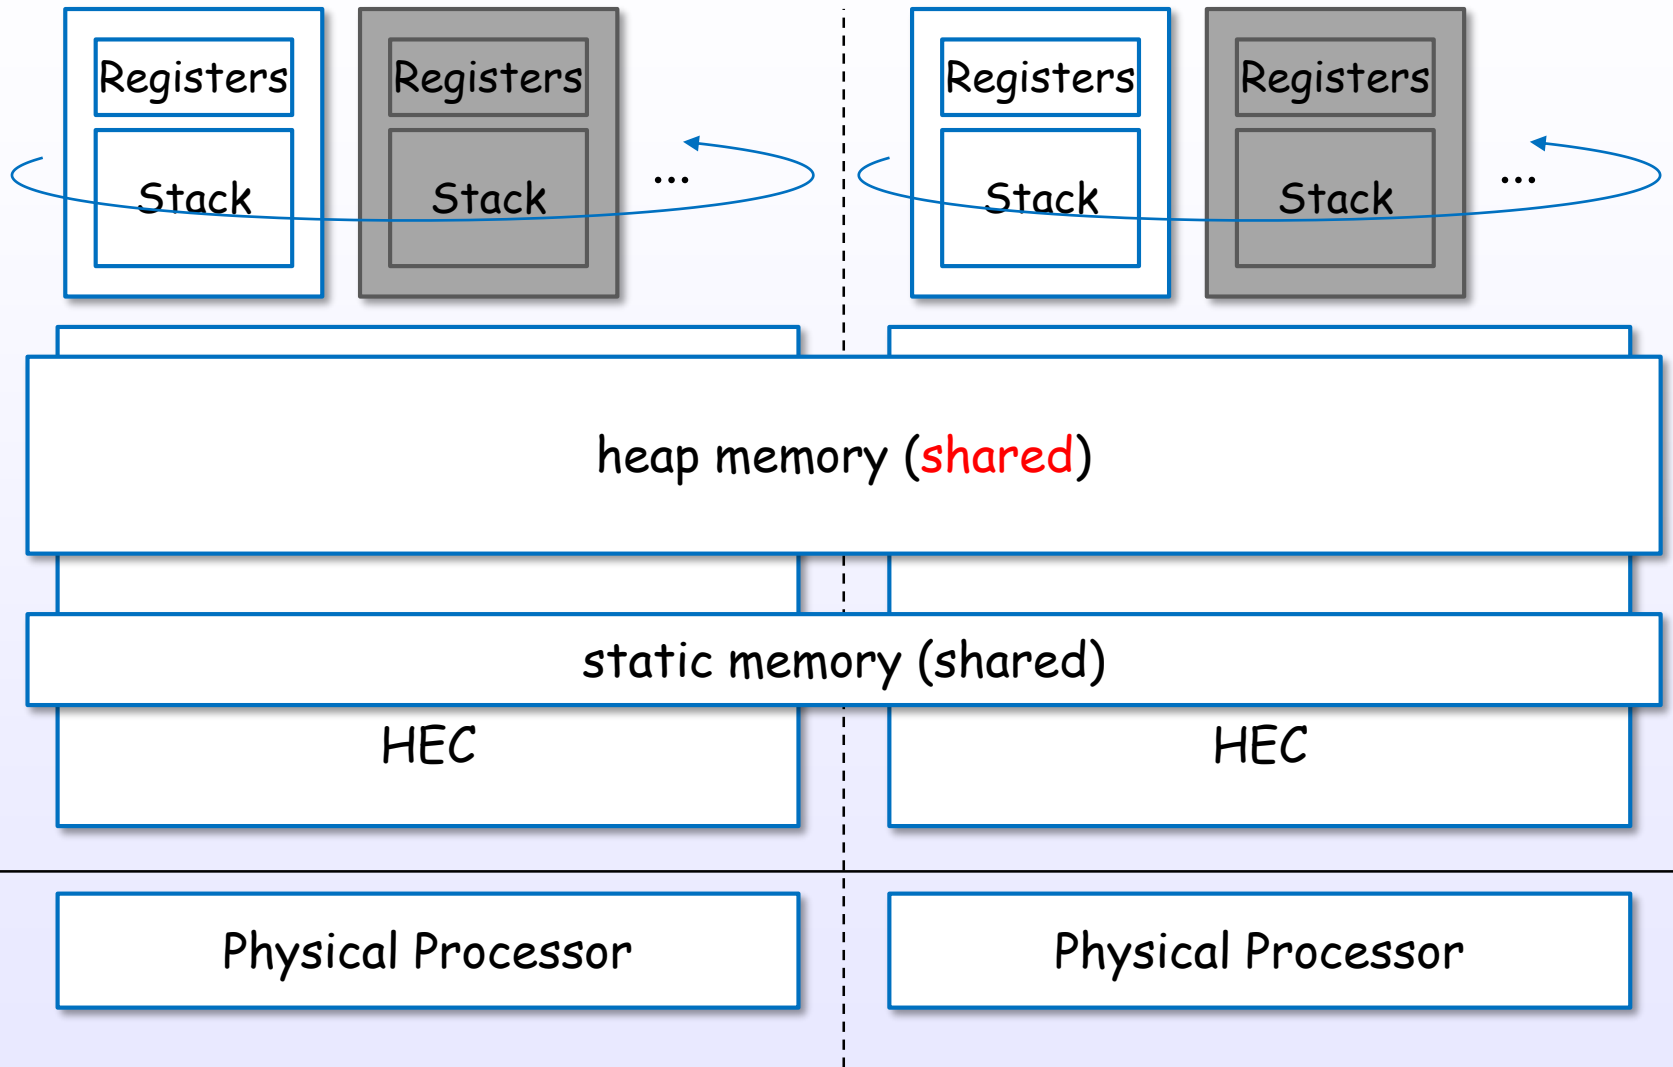
# Threads are migrated to idle HECs



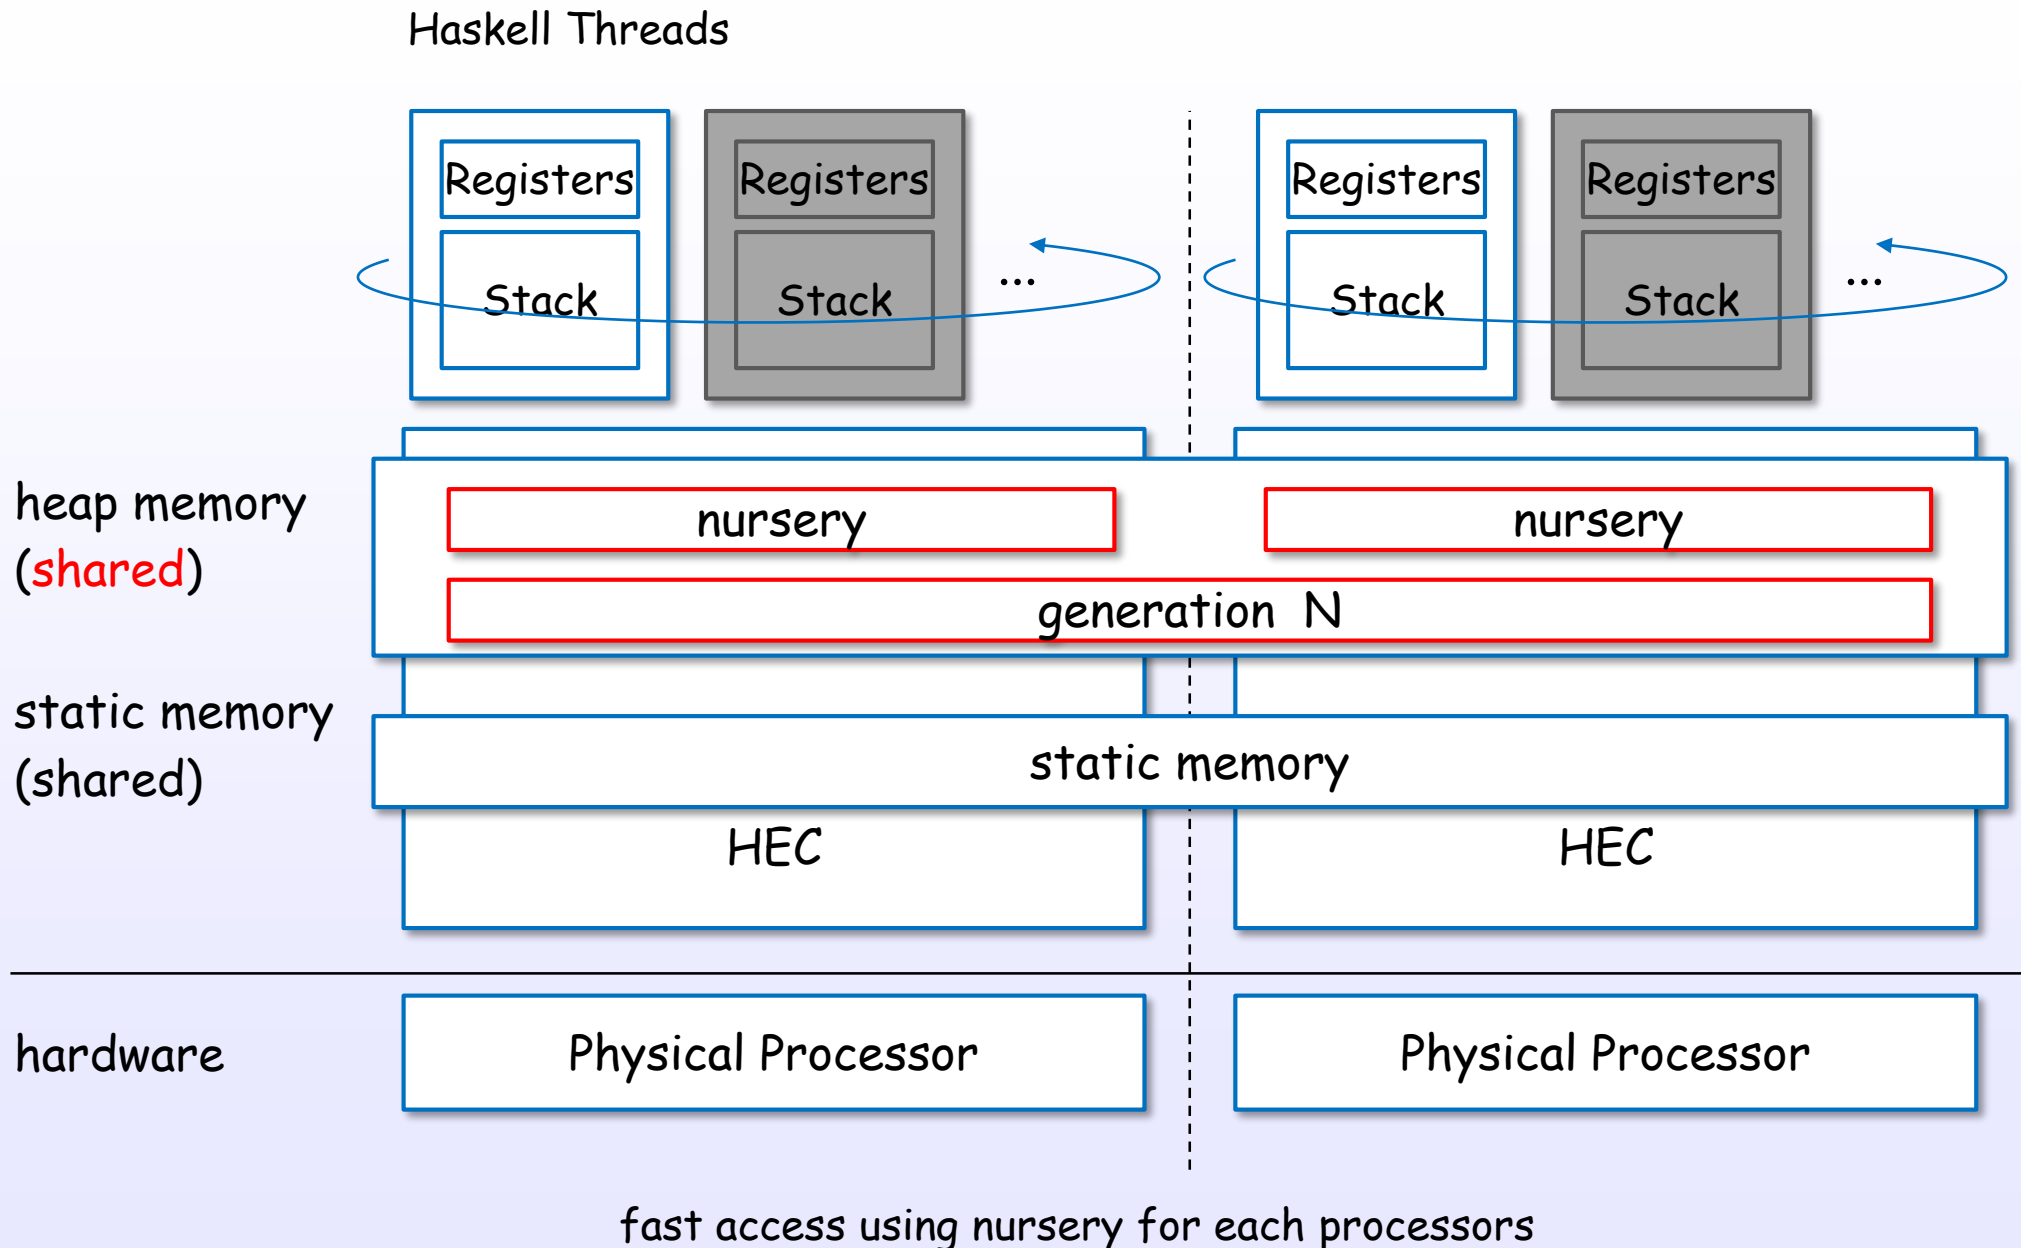
# Heap and Threads

# Threads share a heap

Haskell Threads

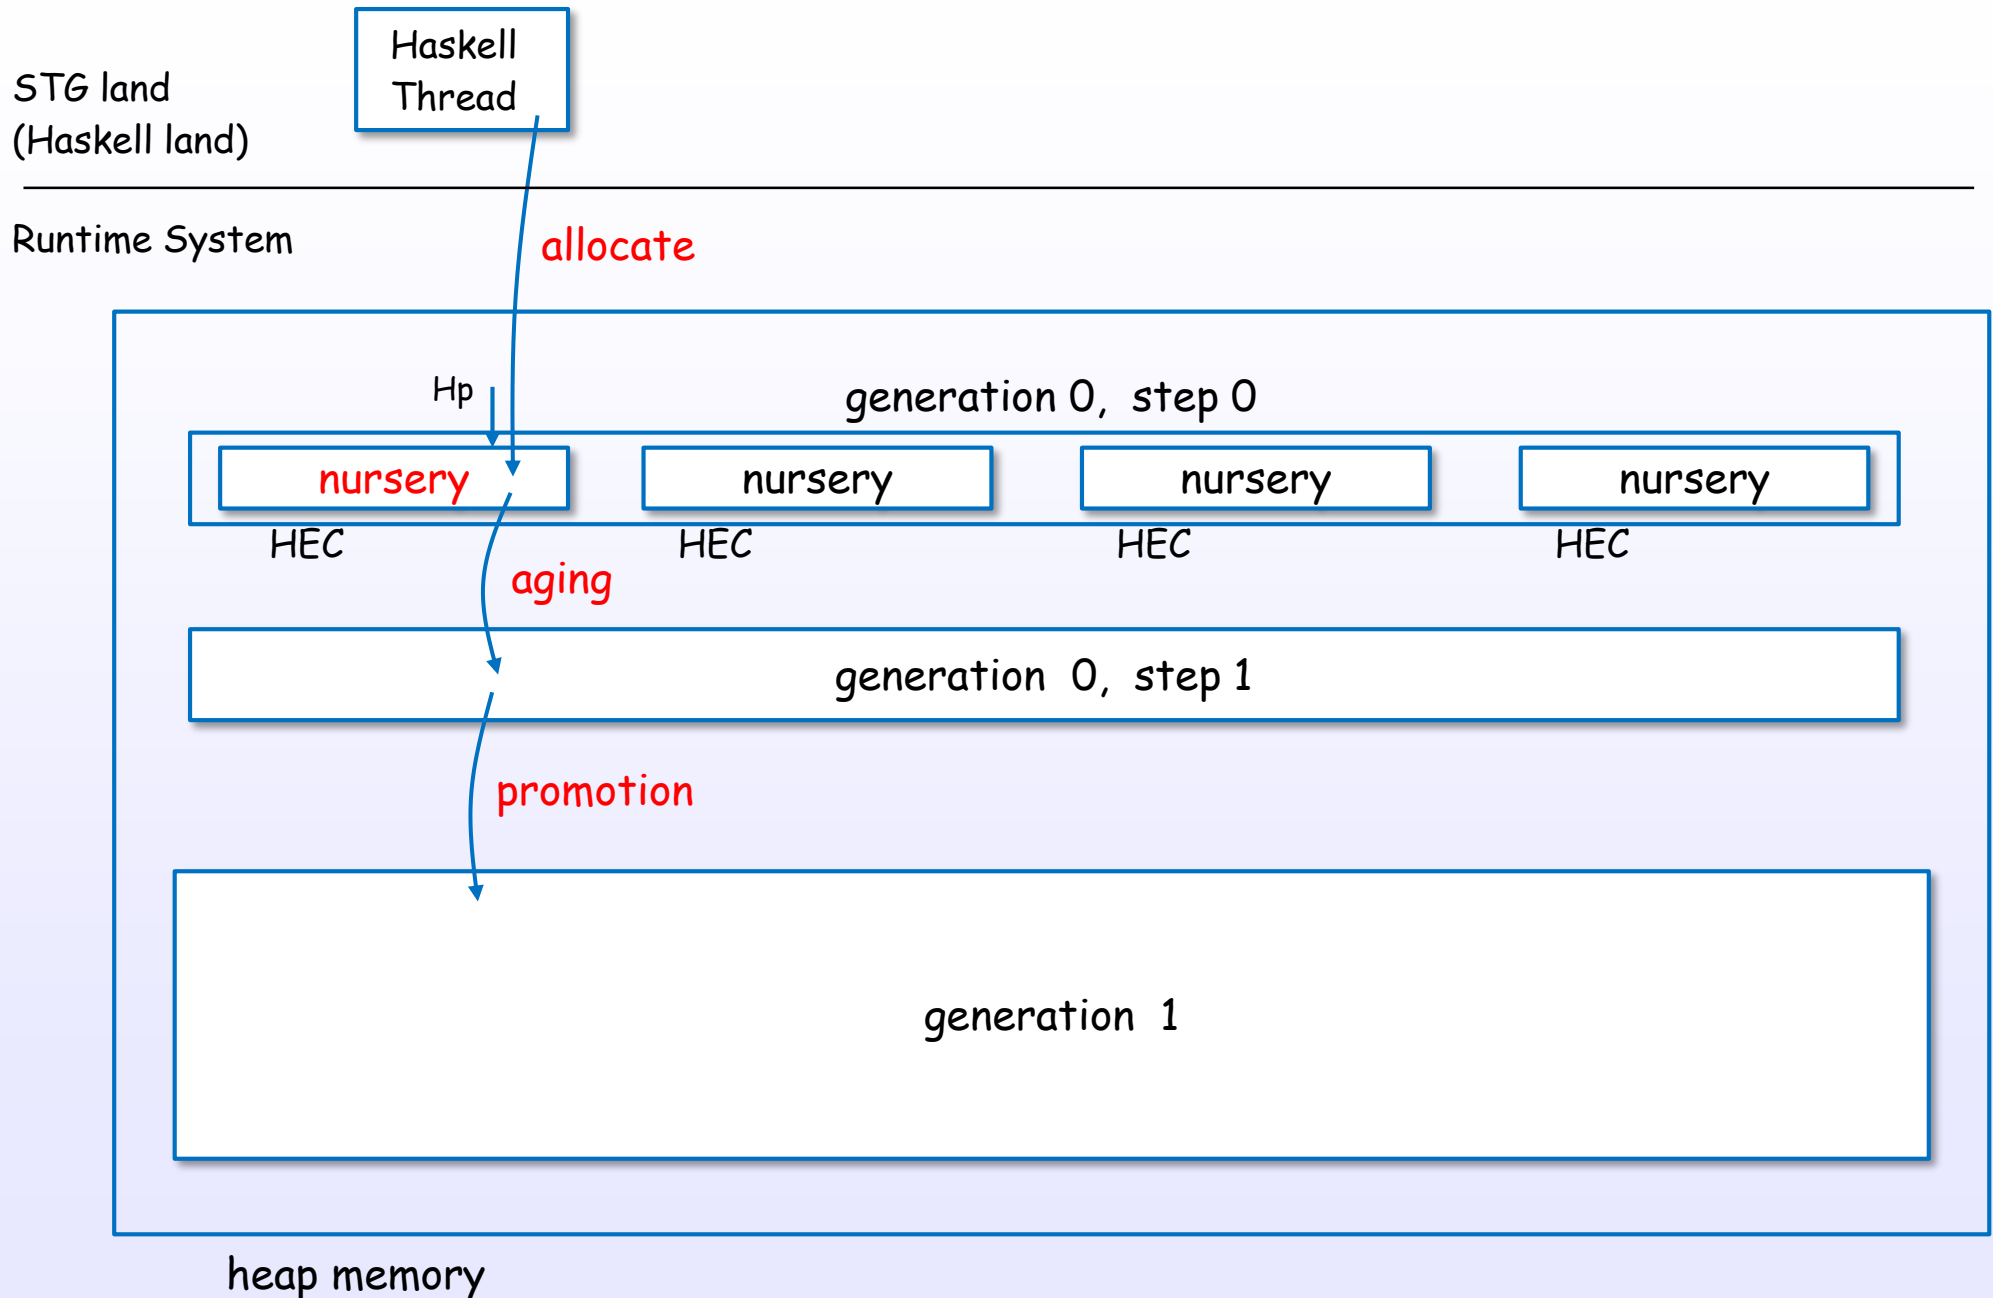


# Local allocation area (nursery)



# Threads and GC

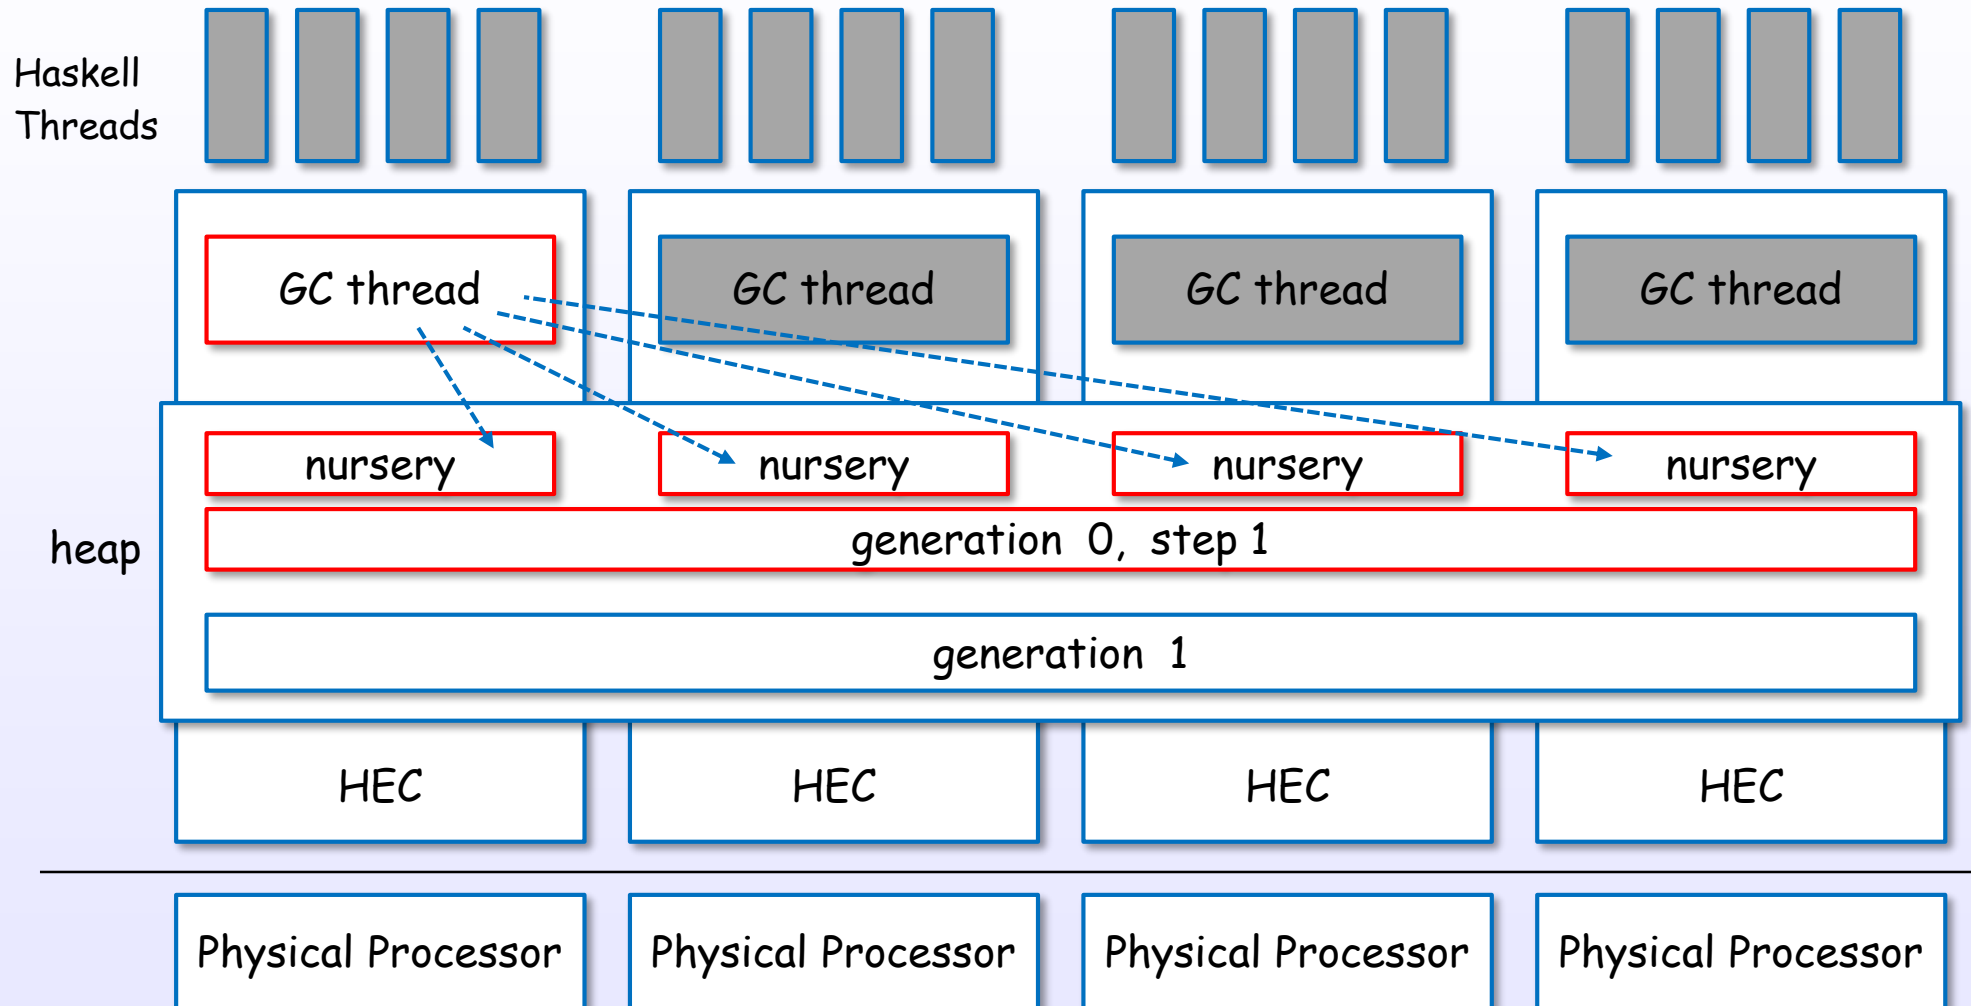
# GC, nursery, generation, aging, promotion



# Threads and minor GC

**sequential** GC for young generation (minor GC)

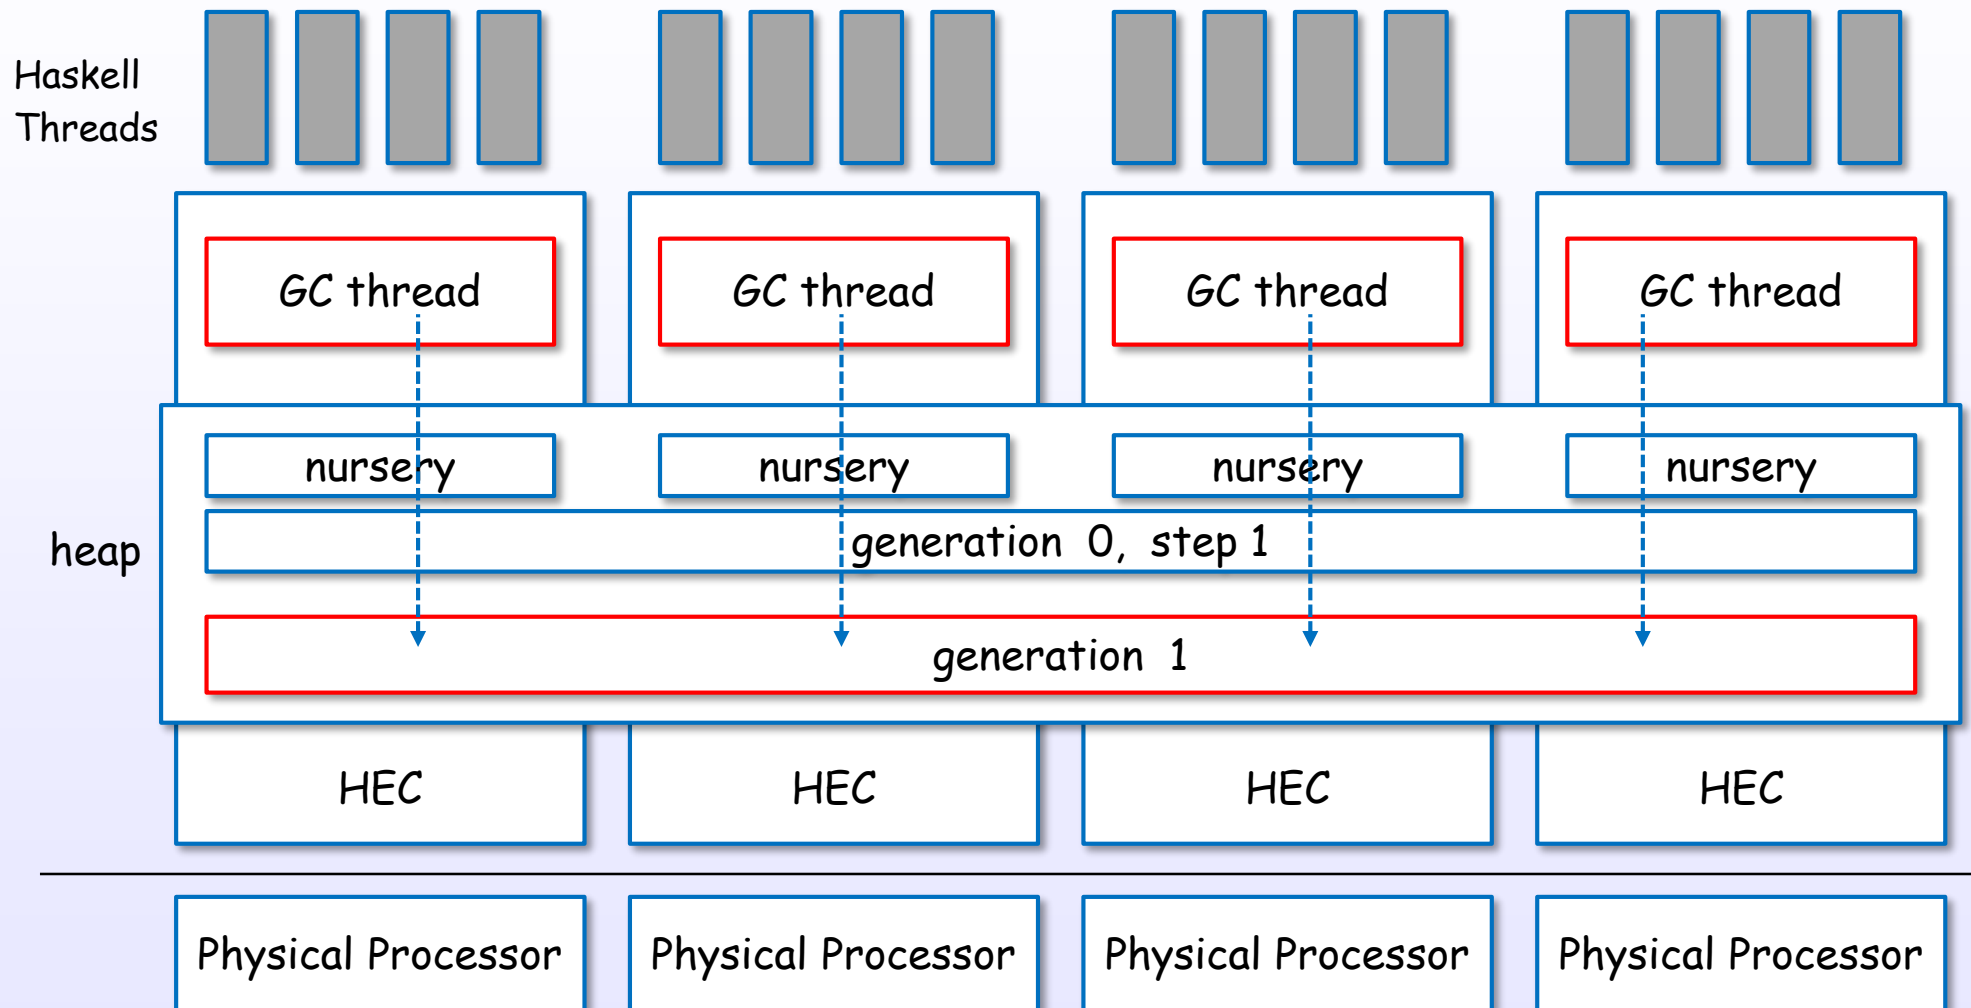
"stop-the-world" GC





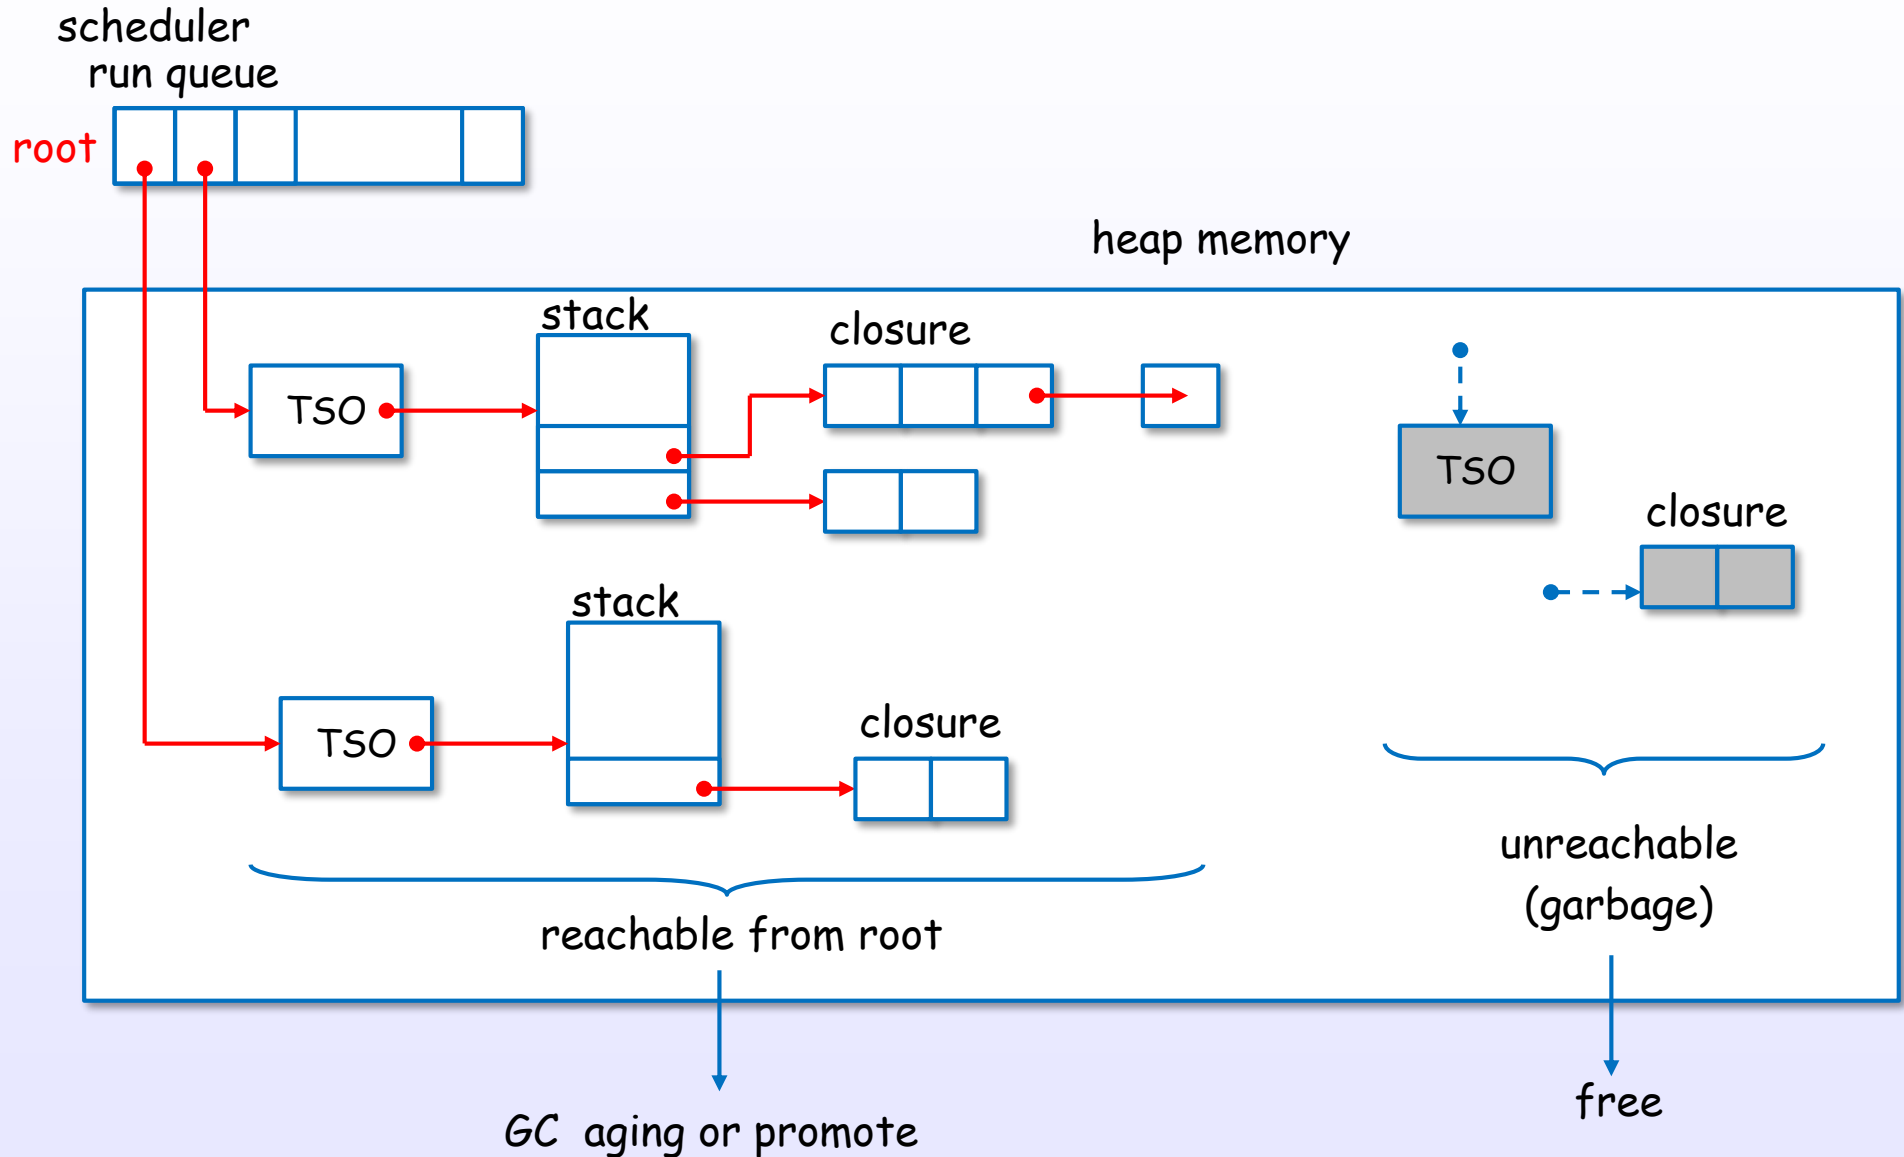
# Threads and major GC

**parallel** GC for oldest generation (major GC)  
"stop-the-world" GC



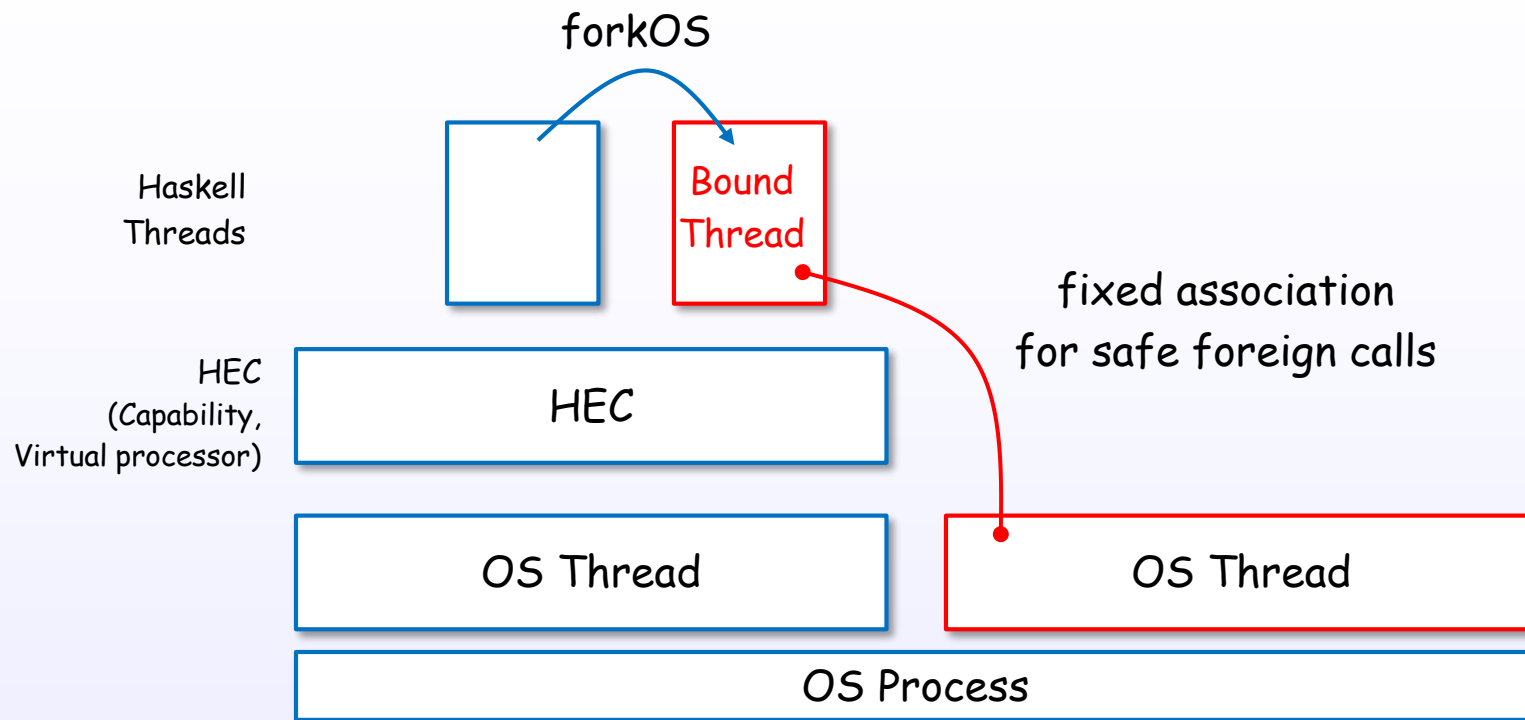
# GC discover live objects from the root

## Runtime System



Bound thread


# A bound thread has a fixed associated OS Thread




Foreign calls from a bound thread are all made by the same OS thread.  
A bound thread is created using `forkOS`.

The main thread is bound thread.


# forkIO, forkOn, forkOS



create a haskell unbound thread



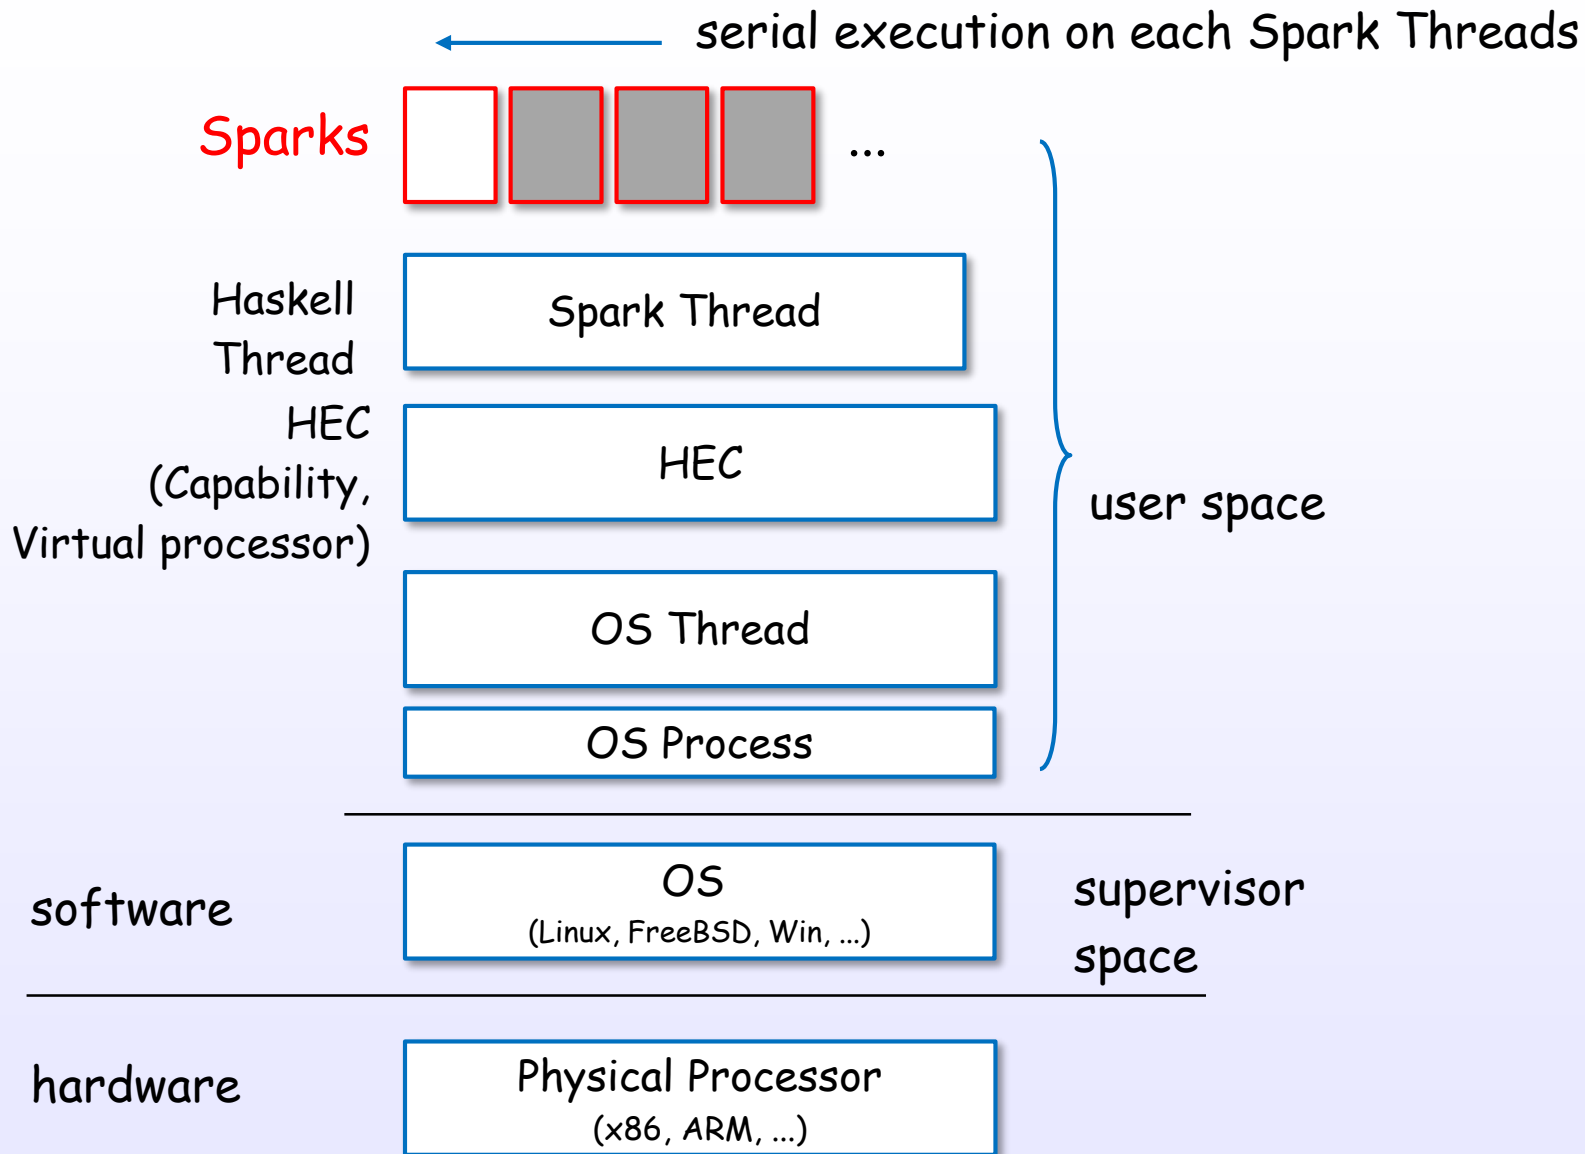
create a haskell unbound thread on the specified HEC



create a haskell **bound** thread and an OS thread

Spark

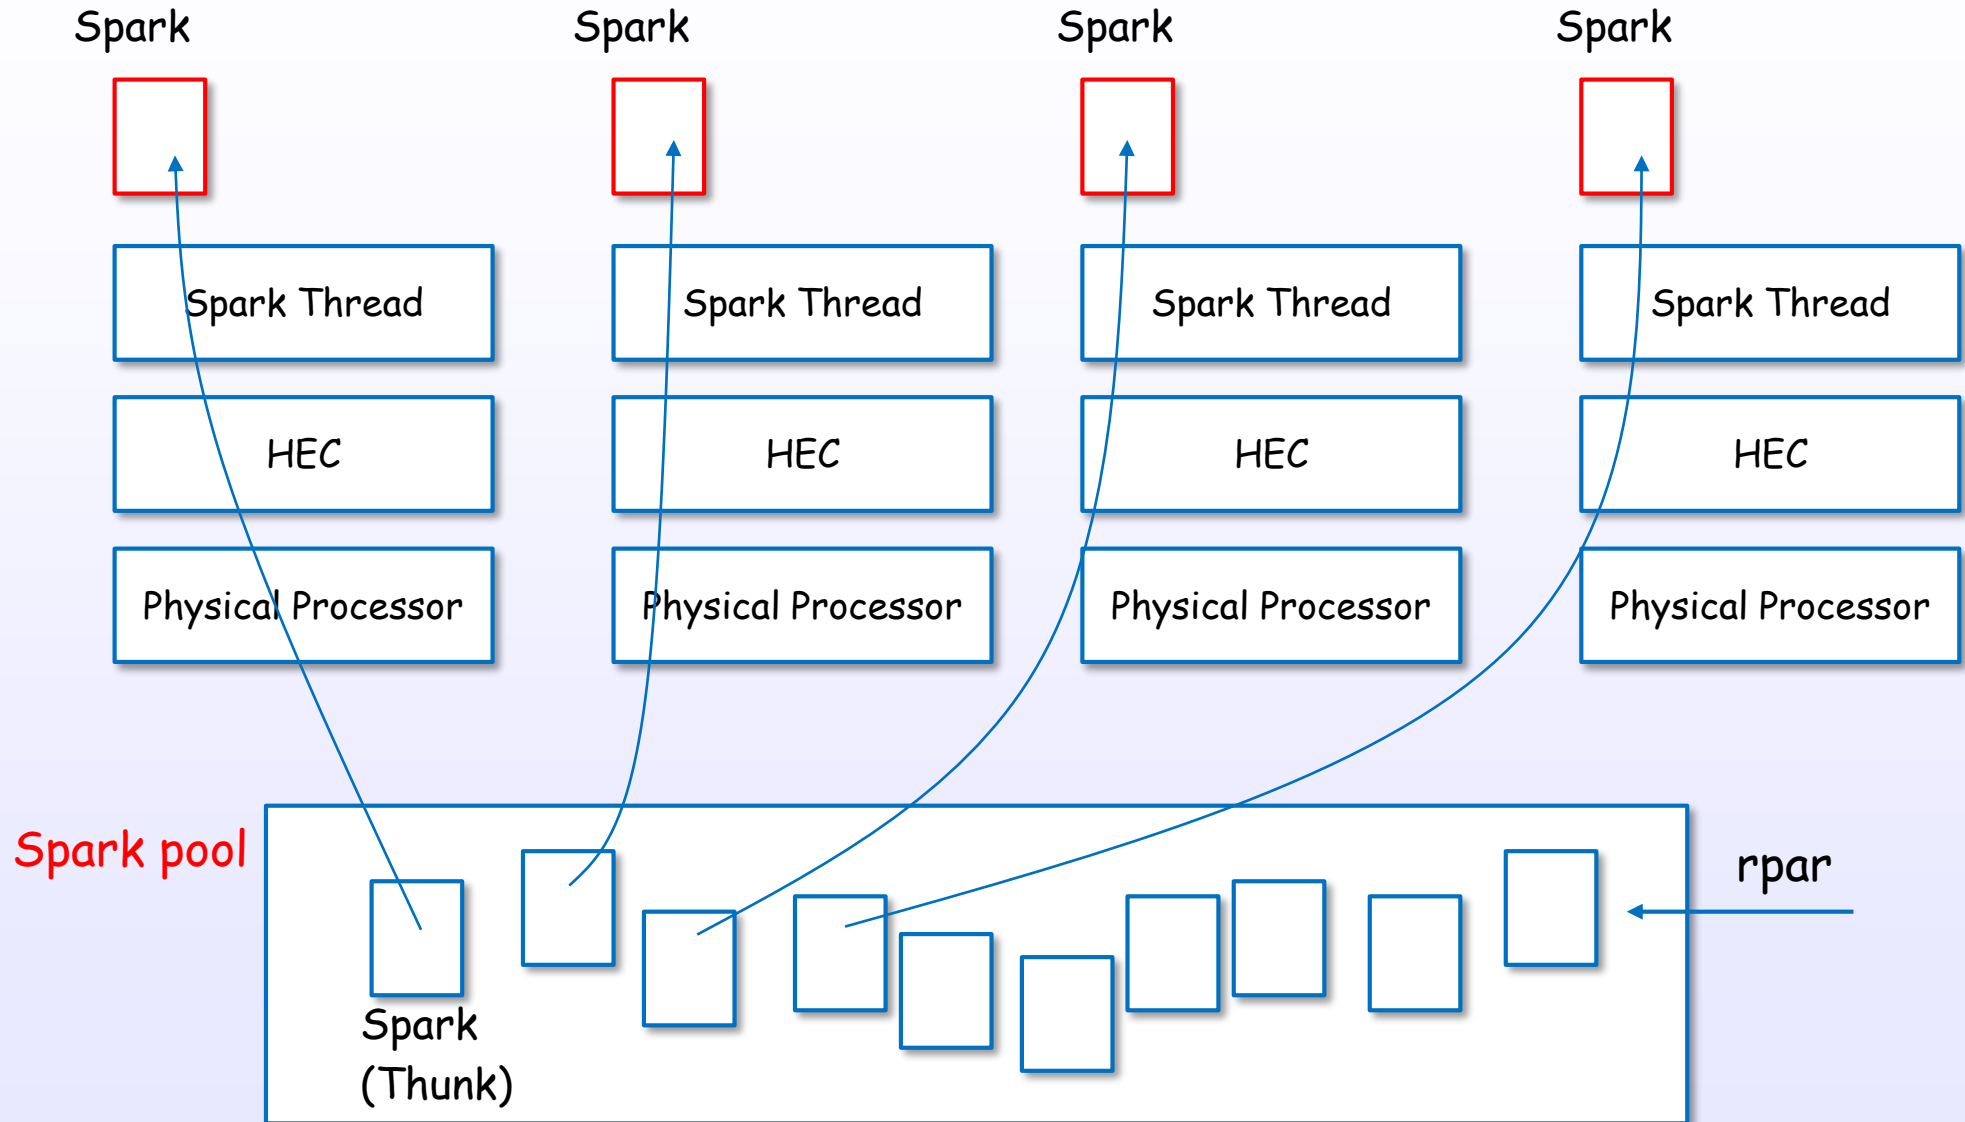
# Spark layer



Spark Threads are generated on idle HECs.

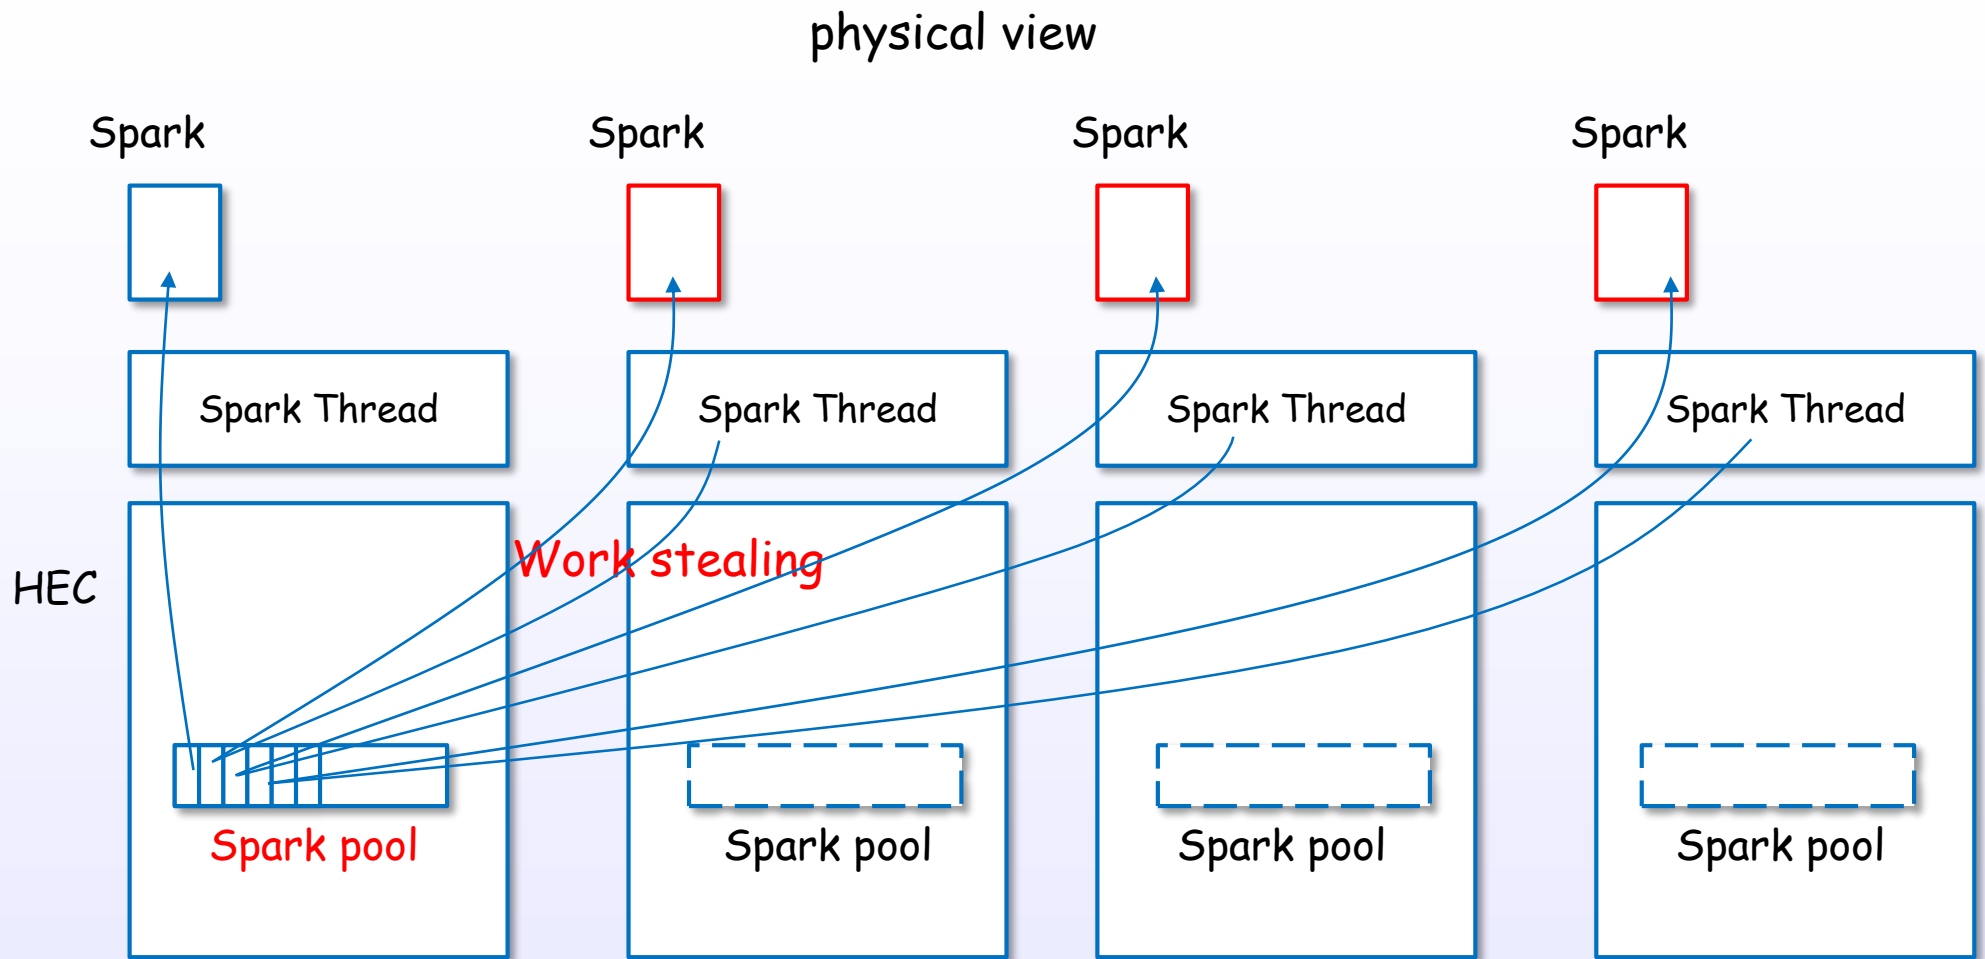
# Sparks and Spark pool

logical view






# Spark pool and work stealing



# Sparks and closures

Spark Thread



STG land  
(Haskell land)

getSpark

C land

HEC

Spark pool  
(WSDeque)



← rpar  
push

...

heap



(not TSO objects, but closures. therefore very lightweight)

MVar


# MVar

Haskell Thread #0




putMVar

Haskell Thread #1



takeMVar



MVar

# MVar and blocking


Haskell Thread



putMVar



BLOCKED  
if full



MVar

Haskell Thread



takeMVar

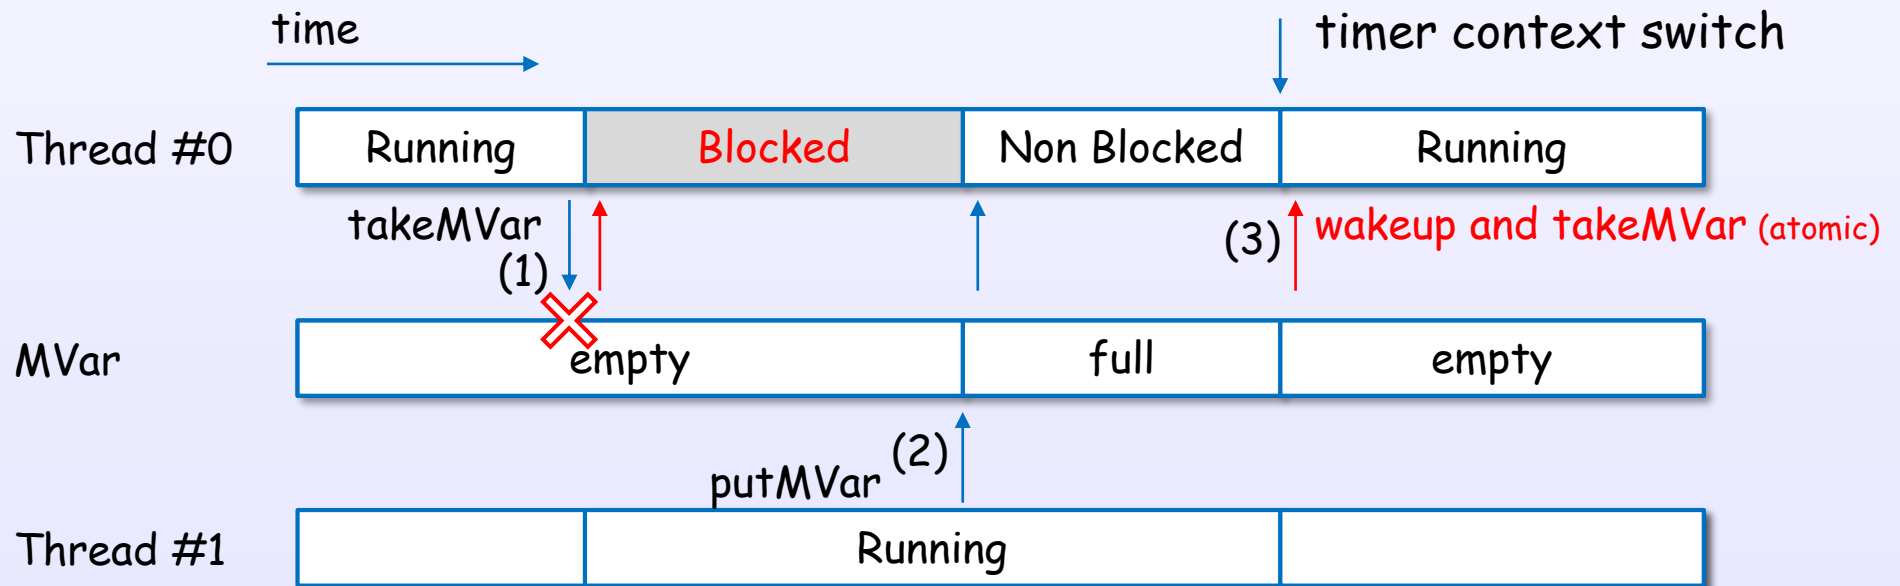



BLOCKED  
if empty



MVar

# MVar example




\* single core case

References : [16], [18], [19], [S31], [S12]

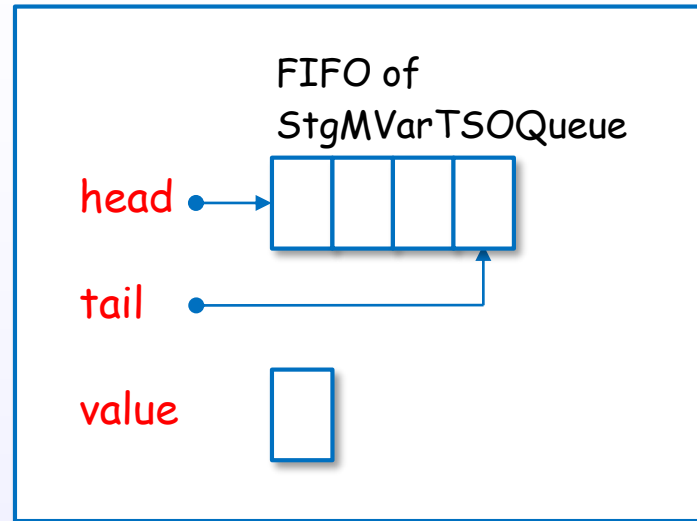
# MVar object view

User view

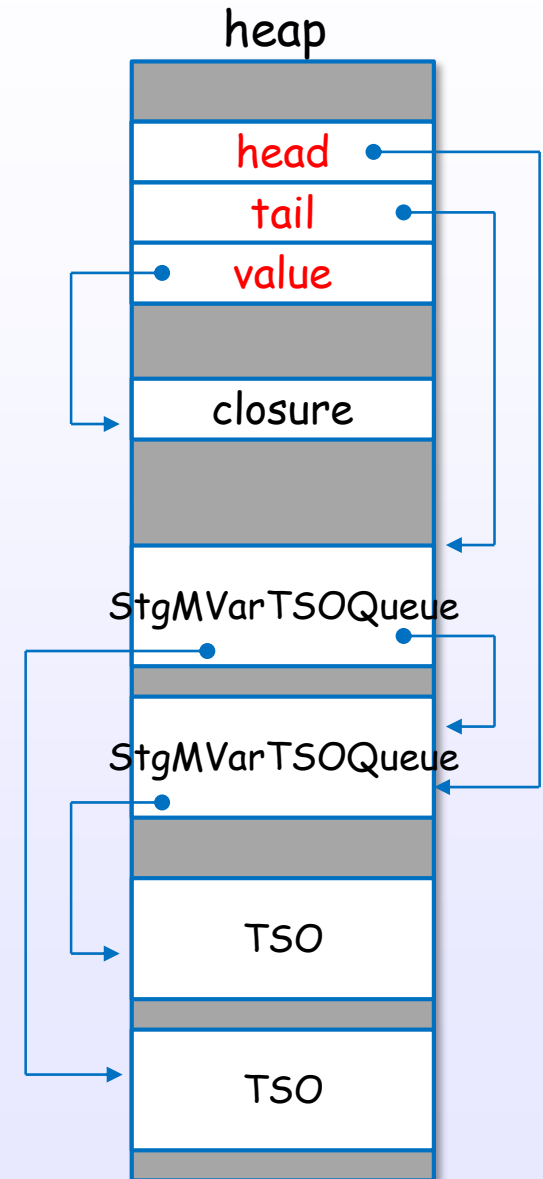
MVar



logical MVar object



physical MVar object



# newEmptyMVar

Haskell Threads


```
newEmptyMVar  
newMVar#
```

(1) **call** the Runtime primitive

Runtime System

```
stg_newMVarzh  
  ALLOC_PRIM_  
  SET_HDR  
  StgMVar_head  
  StgMVar_tail  
  StgMVar_value
```

(2) **create a MVar** object in the heap



(3) **link** each fields

References : [16], [18], [19], [S31], [S12]



# takeMVar (empty case)

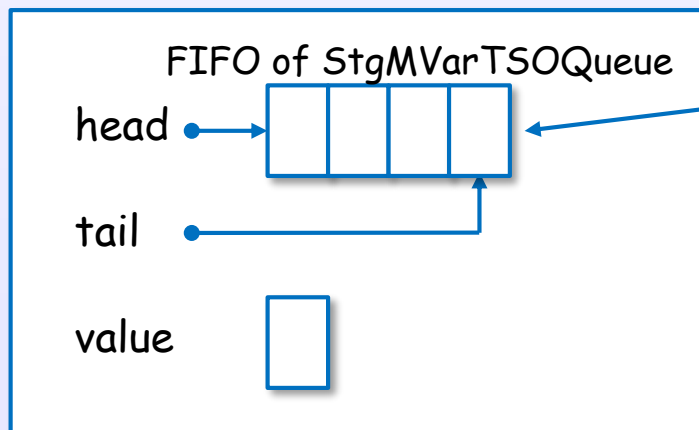
Haskell Threads

```
takeMVar  
takeMVar#
```

Runtime System

```
stg_takeMVarzh  
  create StgMVarTSOQueue ... (1)  
  append ... (2)  
  StgReturn ... (3)
```

(3) return to the scheduler



MVar object

(1) create

StgMVarTSOQueue

(2) append

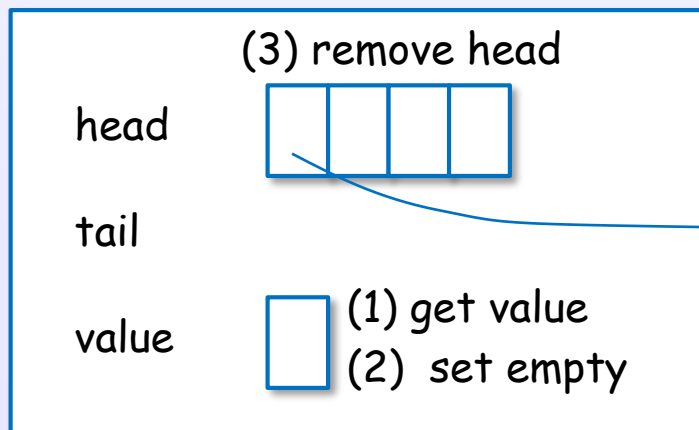
# takeMVar (full case)

## Haskell Threads

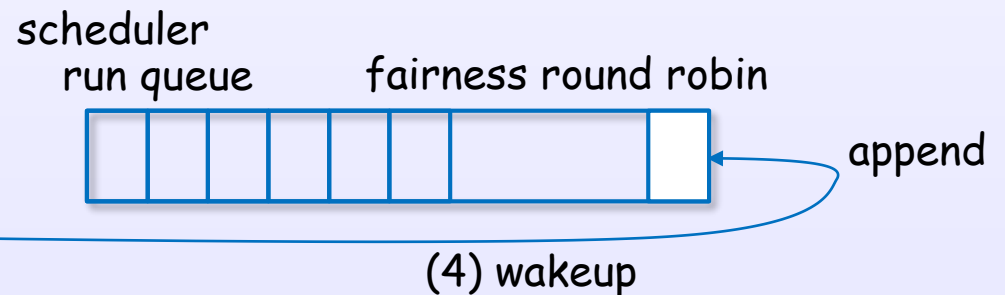
```
takeMVar  
takeMVar#
```

## Runtime System

```
stg_takeMVarzh  
(1) get value  
(2) set empty  
(3) remove head  
(4) tryWakeupThread
```



MVar object



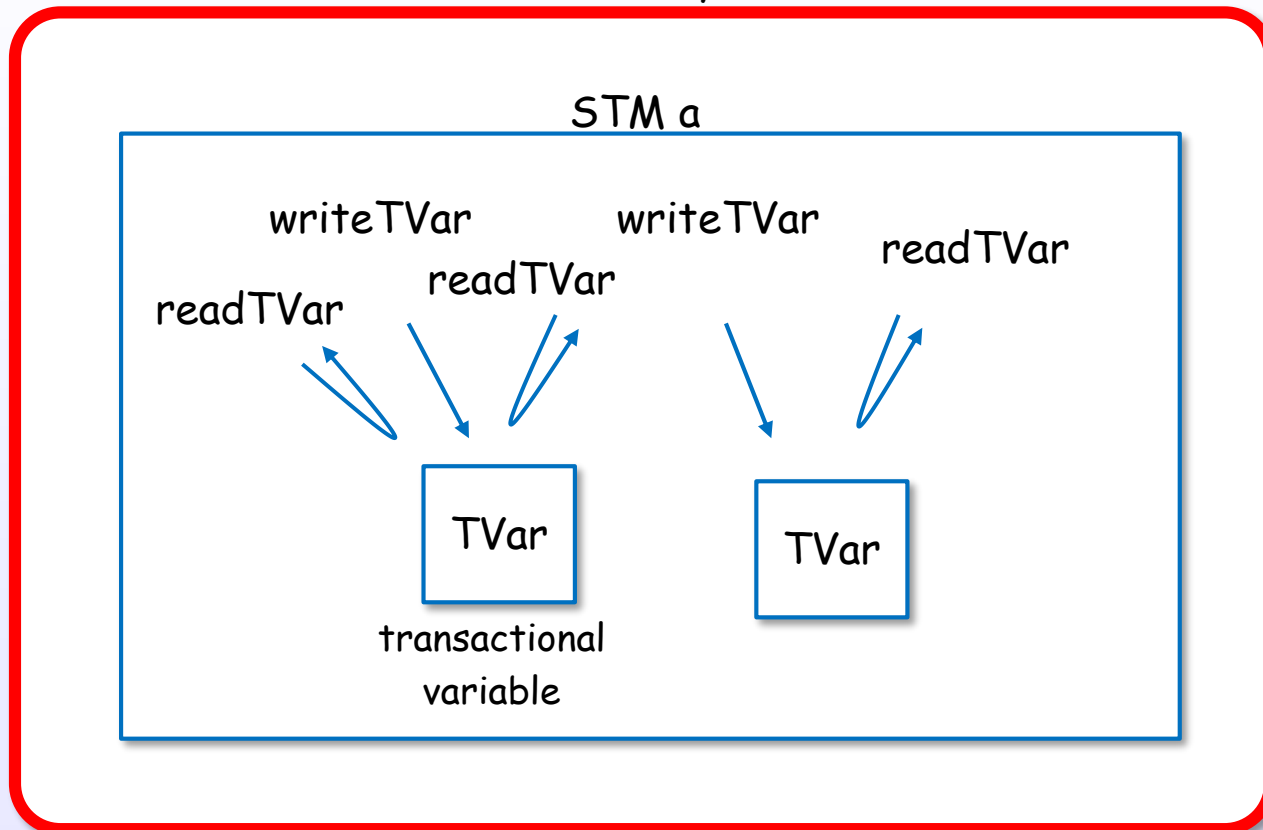
Only one of the blocked threads becomes unblocked.

# Software transactional memory

# Create a atomic block using atomically

**atomically** :: STM a -> IO a

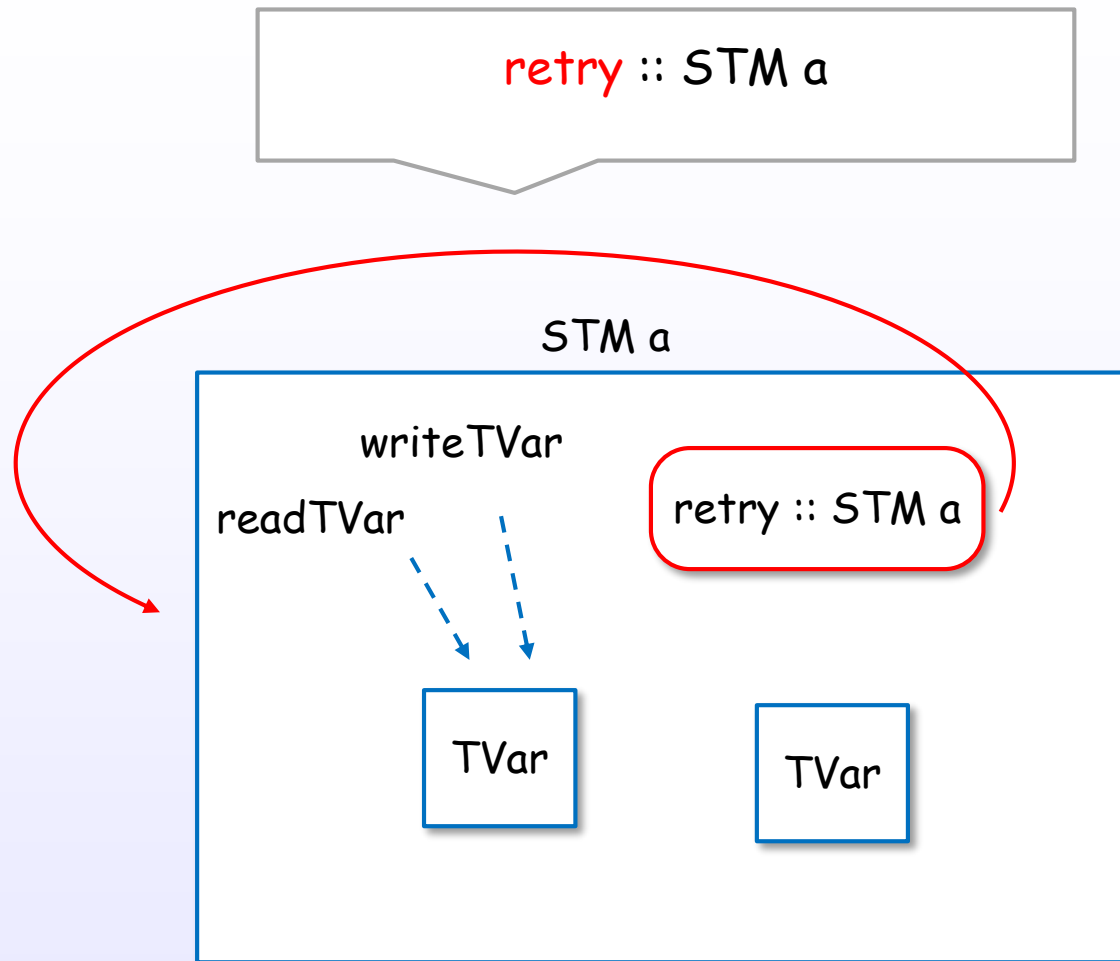
atomically



Create and evaluate a **composable** "atomic block"

Atomic block = All or Nothing

# Rollback and blocking control using retry

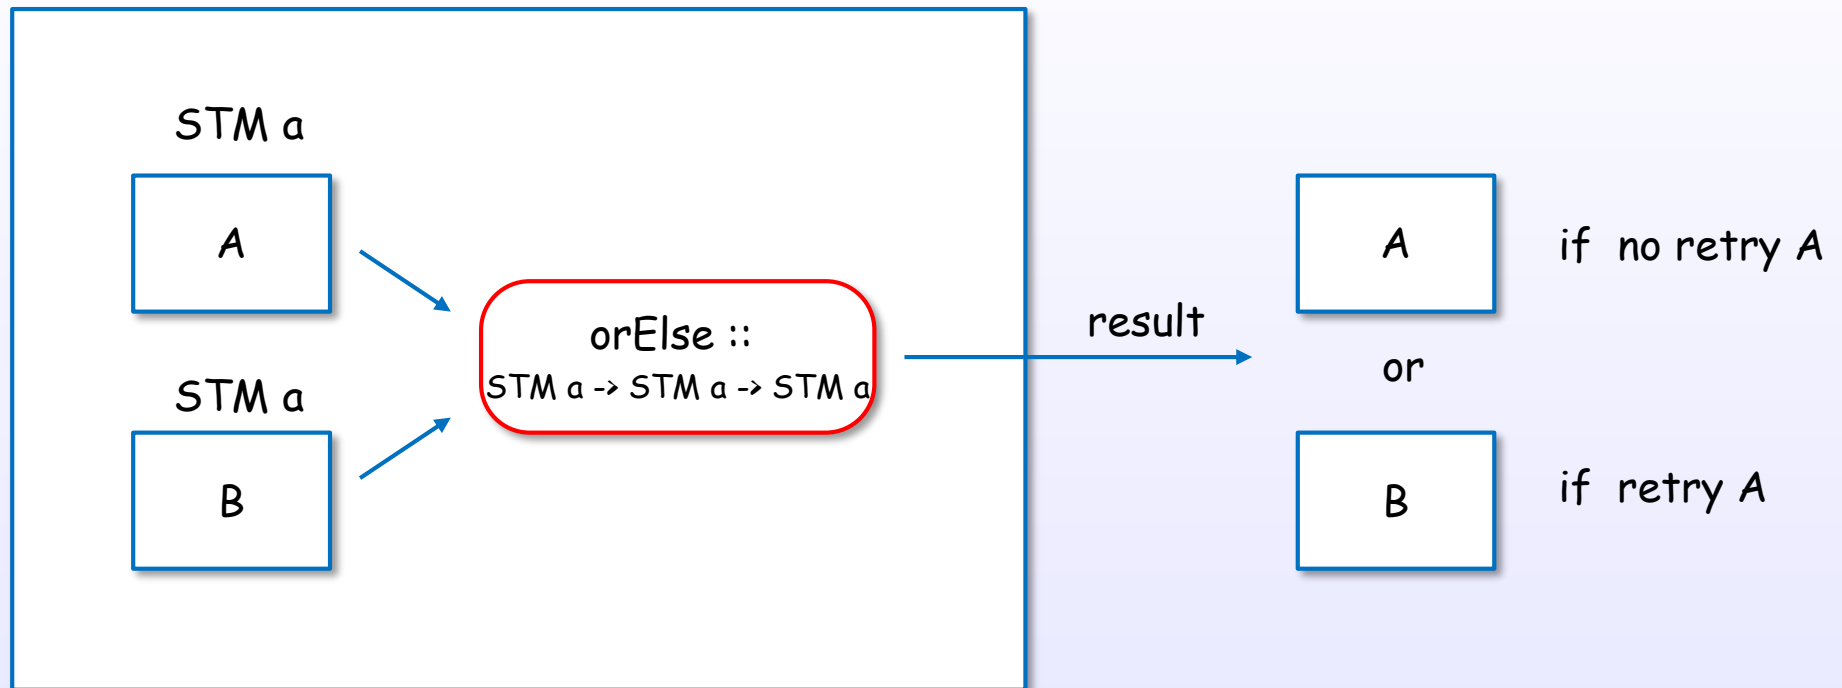


Discard, blocking and try again

# Compose OR case using orElse

`orElse :: STM a -> STM a -> STM a`

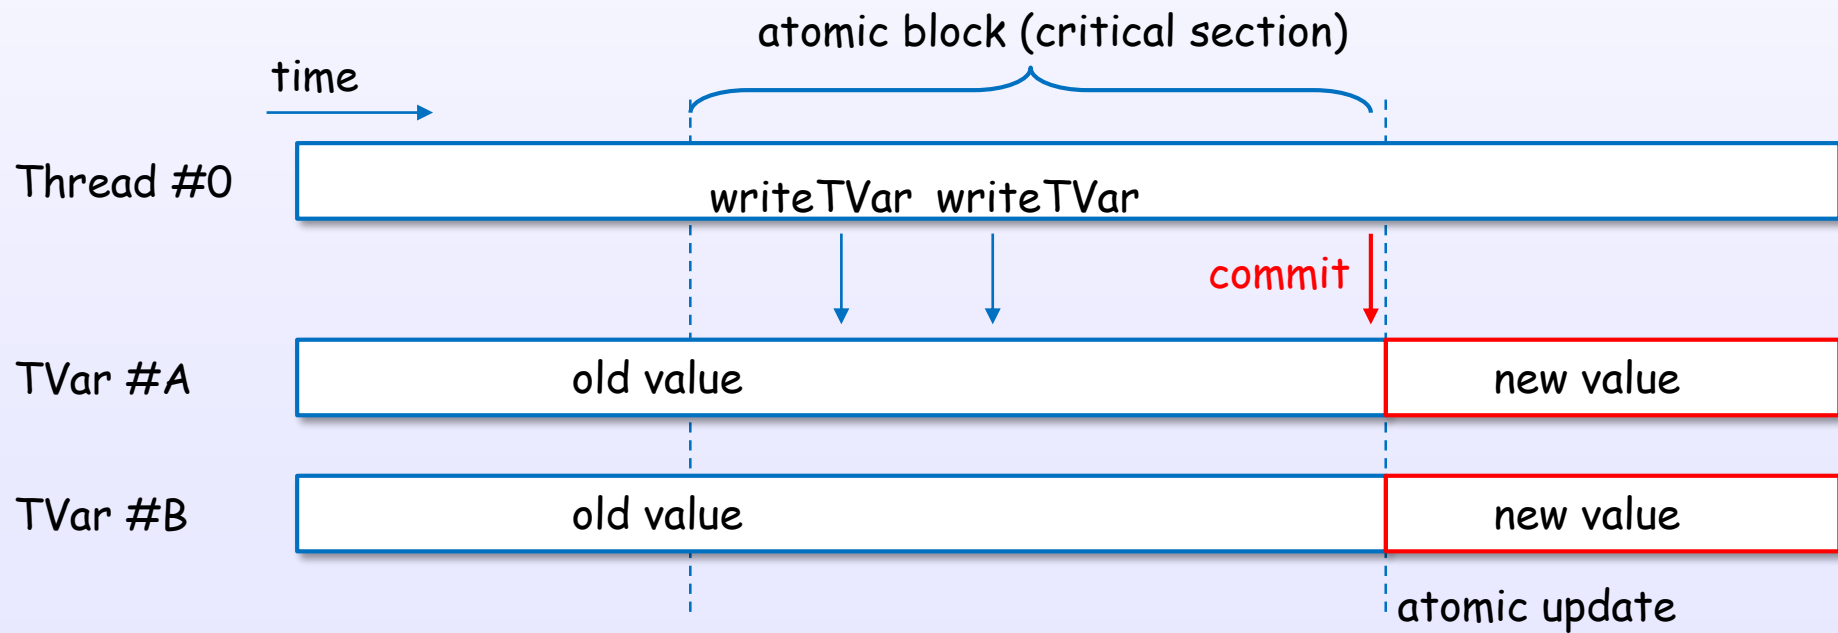
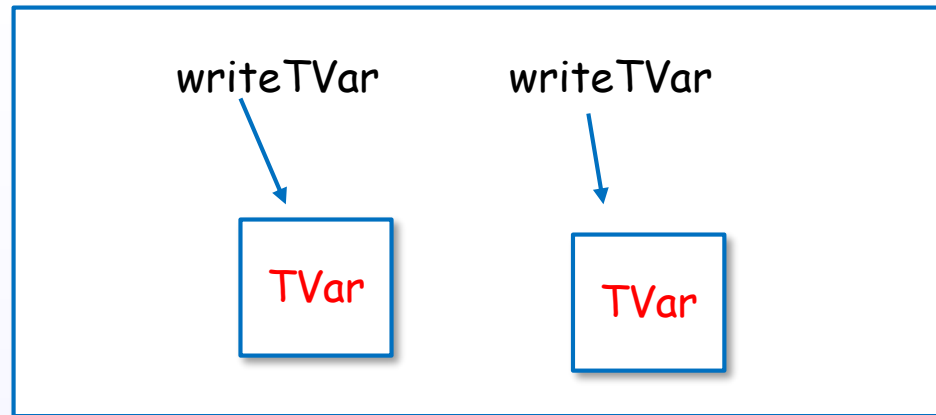
STM a



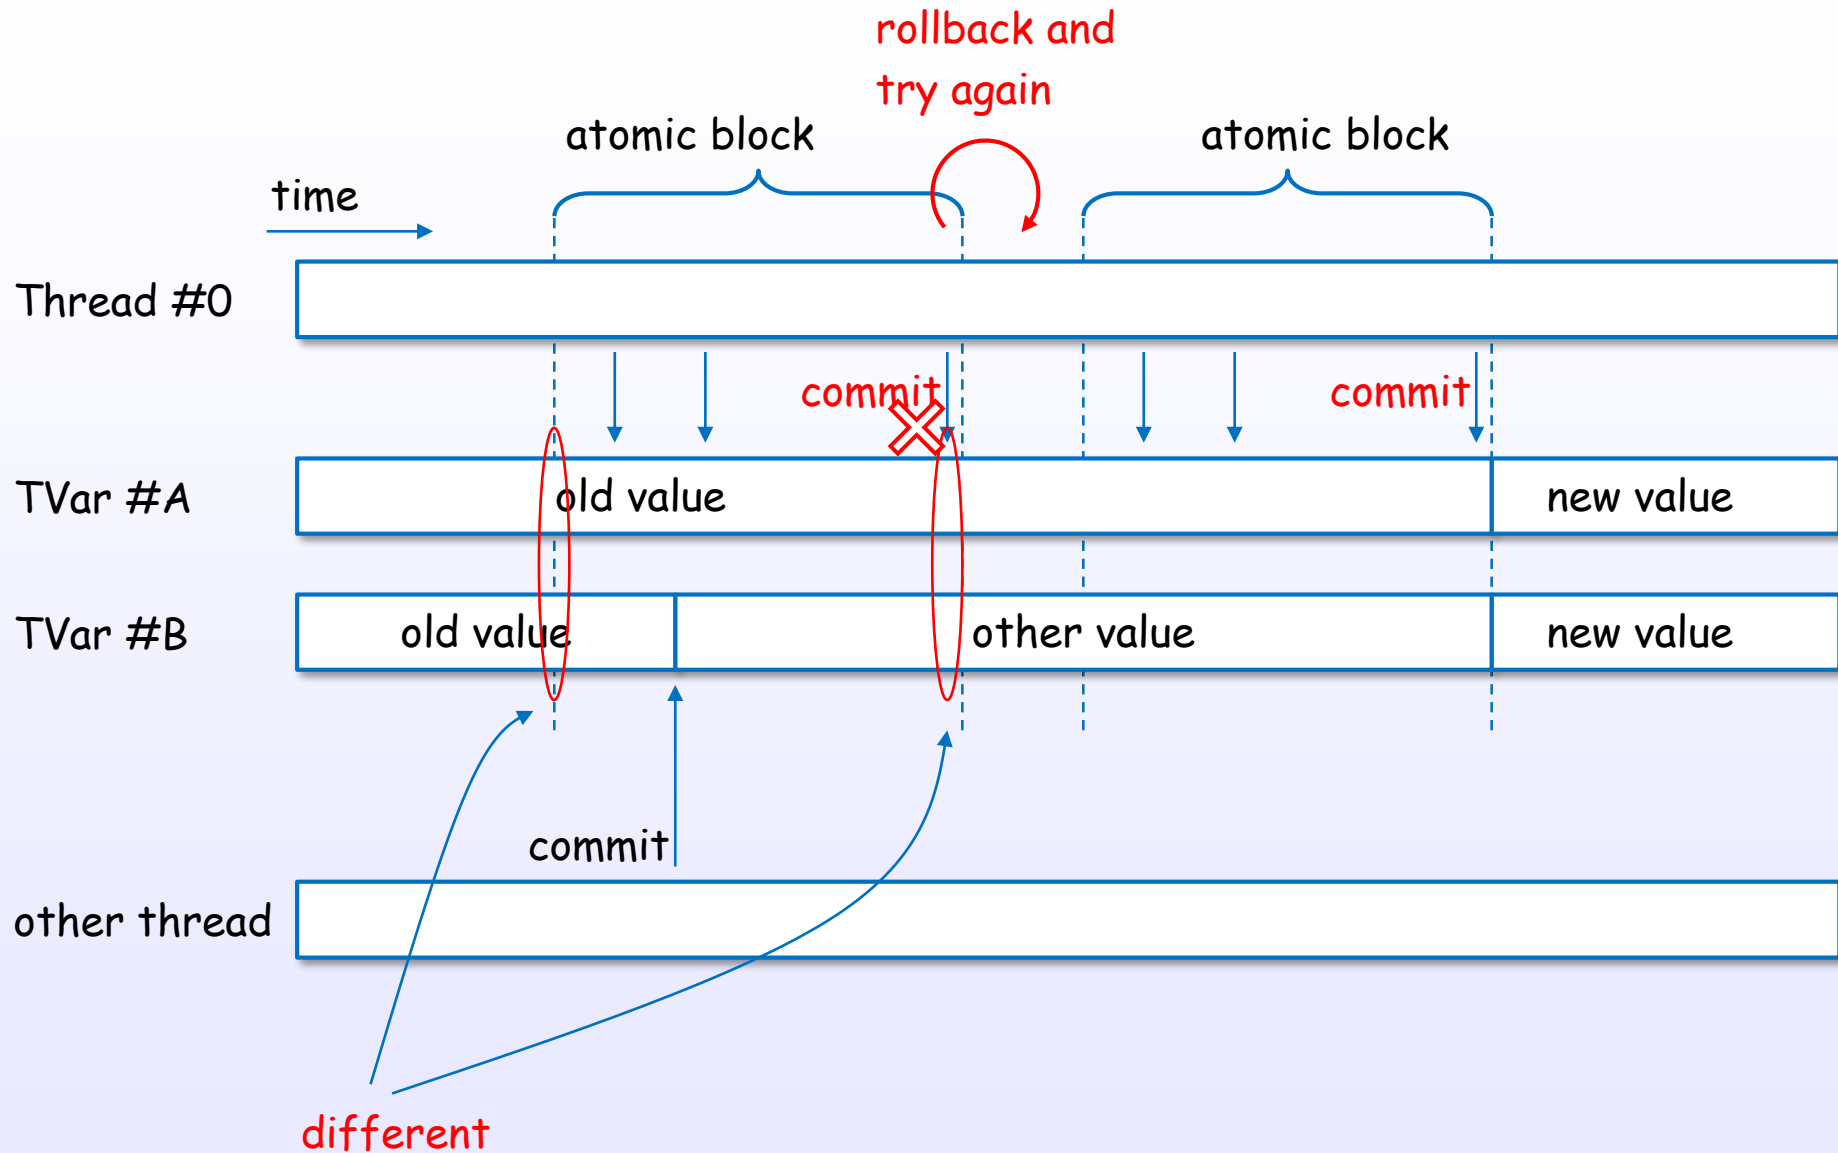
A **or** B or Nothing

# STM, TVar example (normal case)

STM a

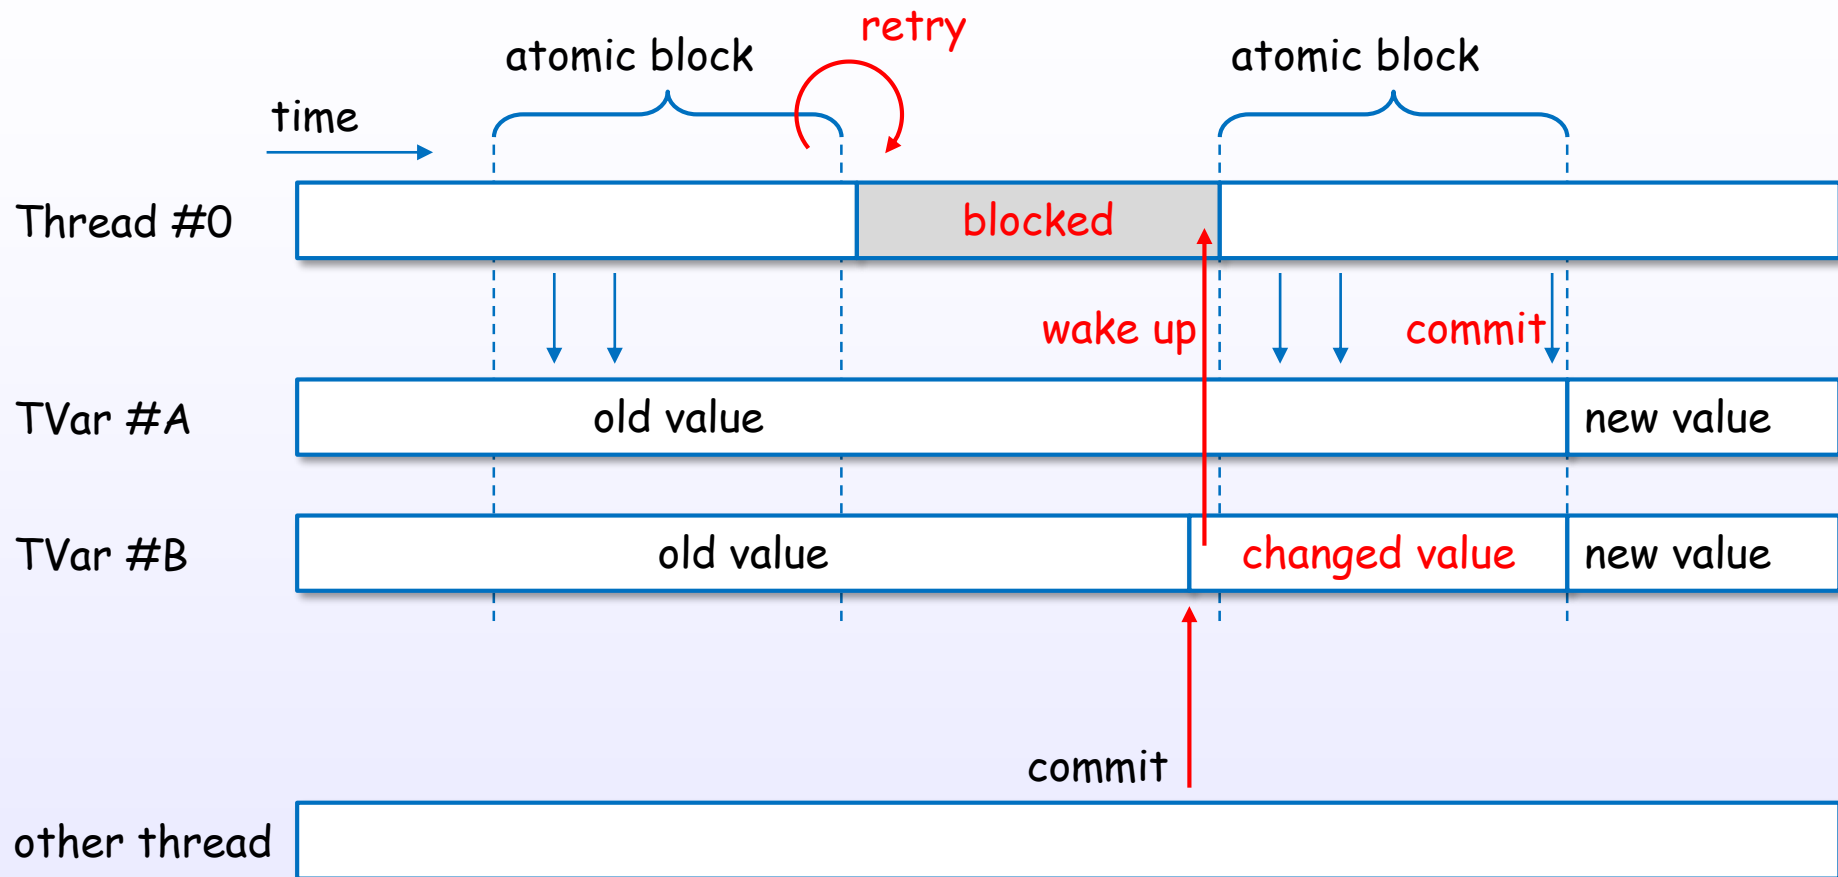


# STM, TVar example (conflict case)

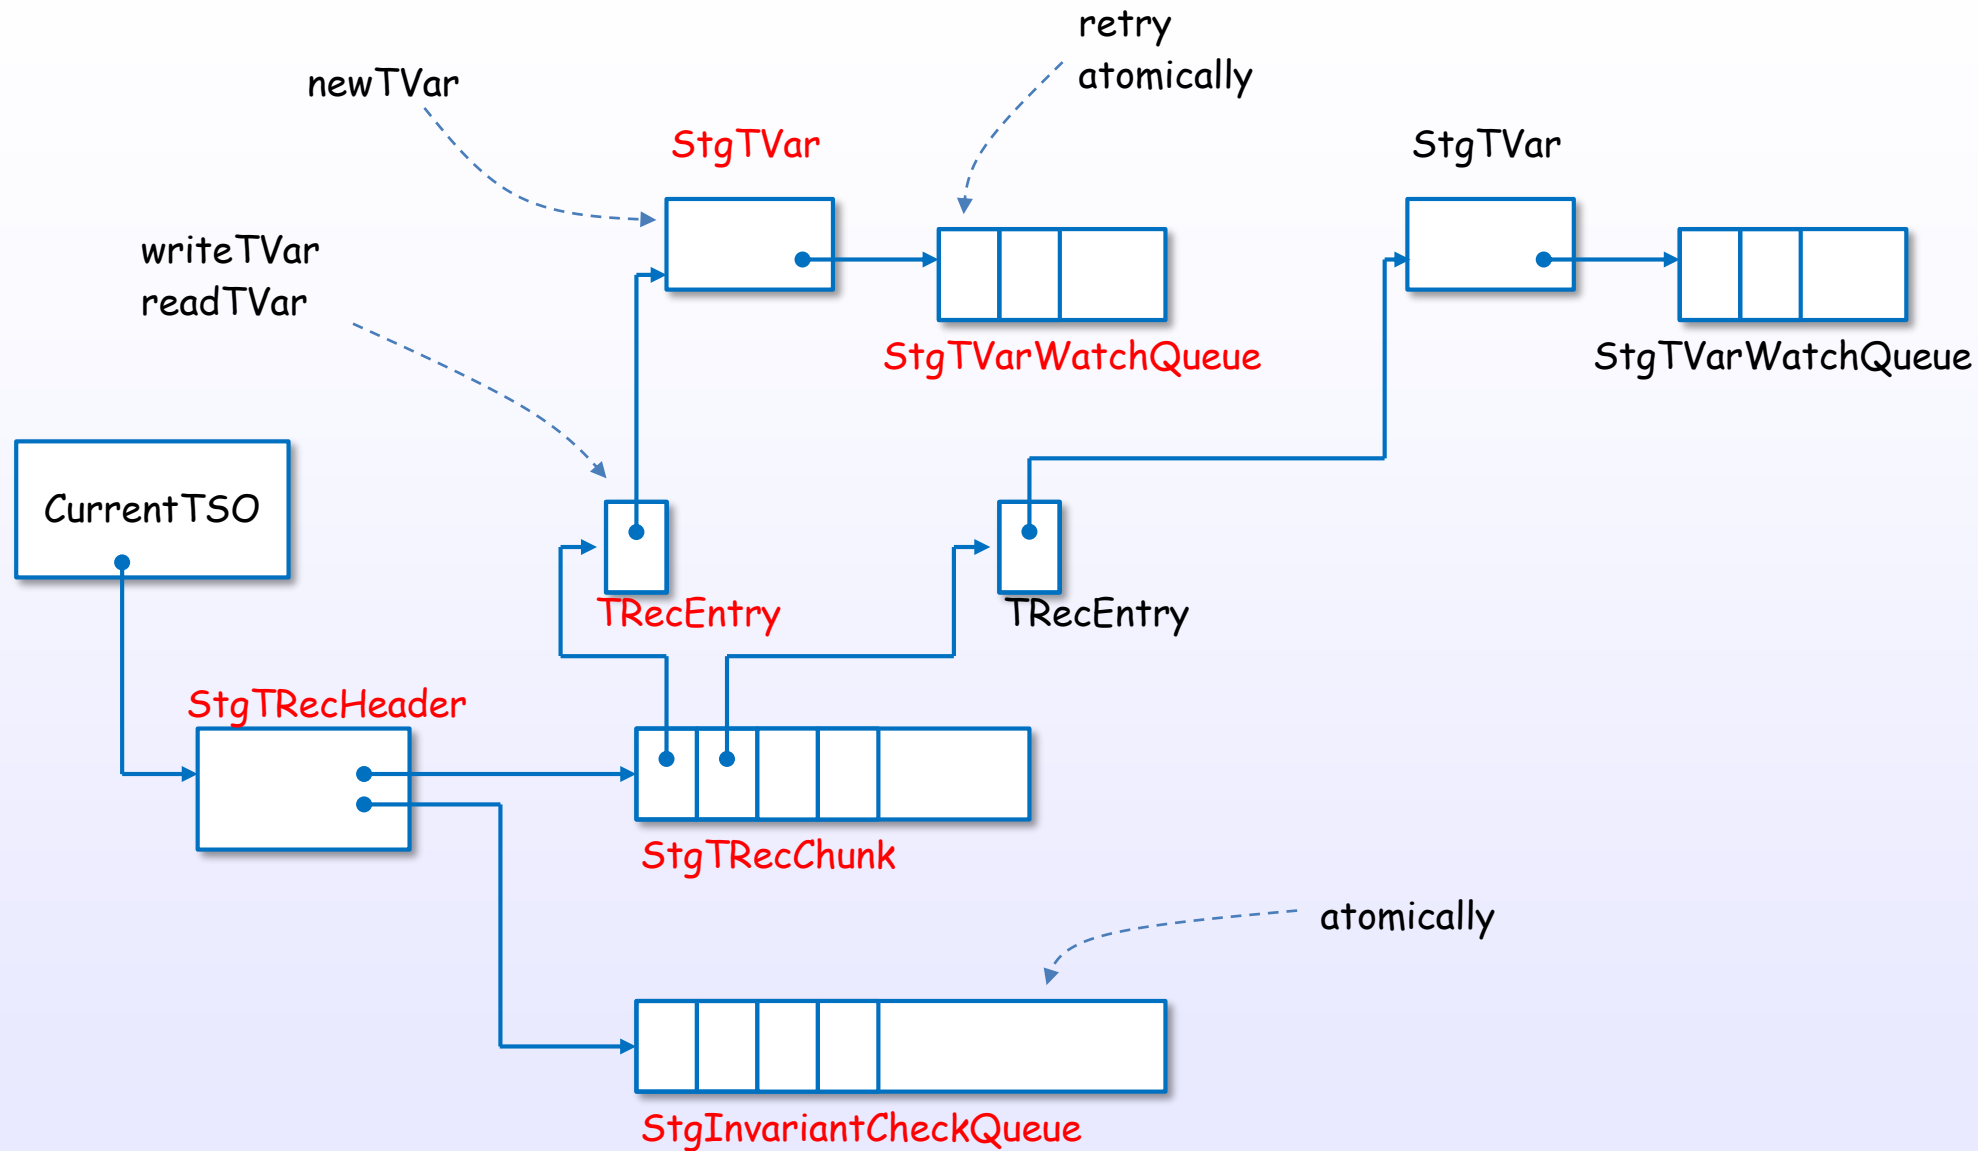




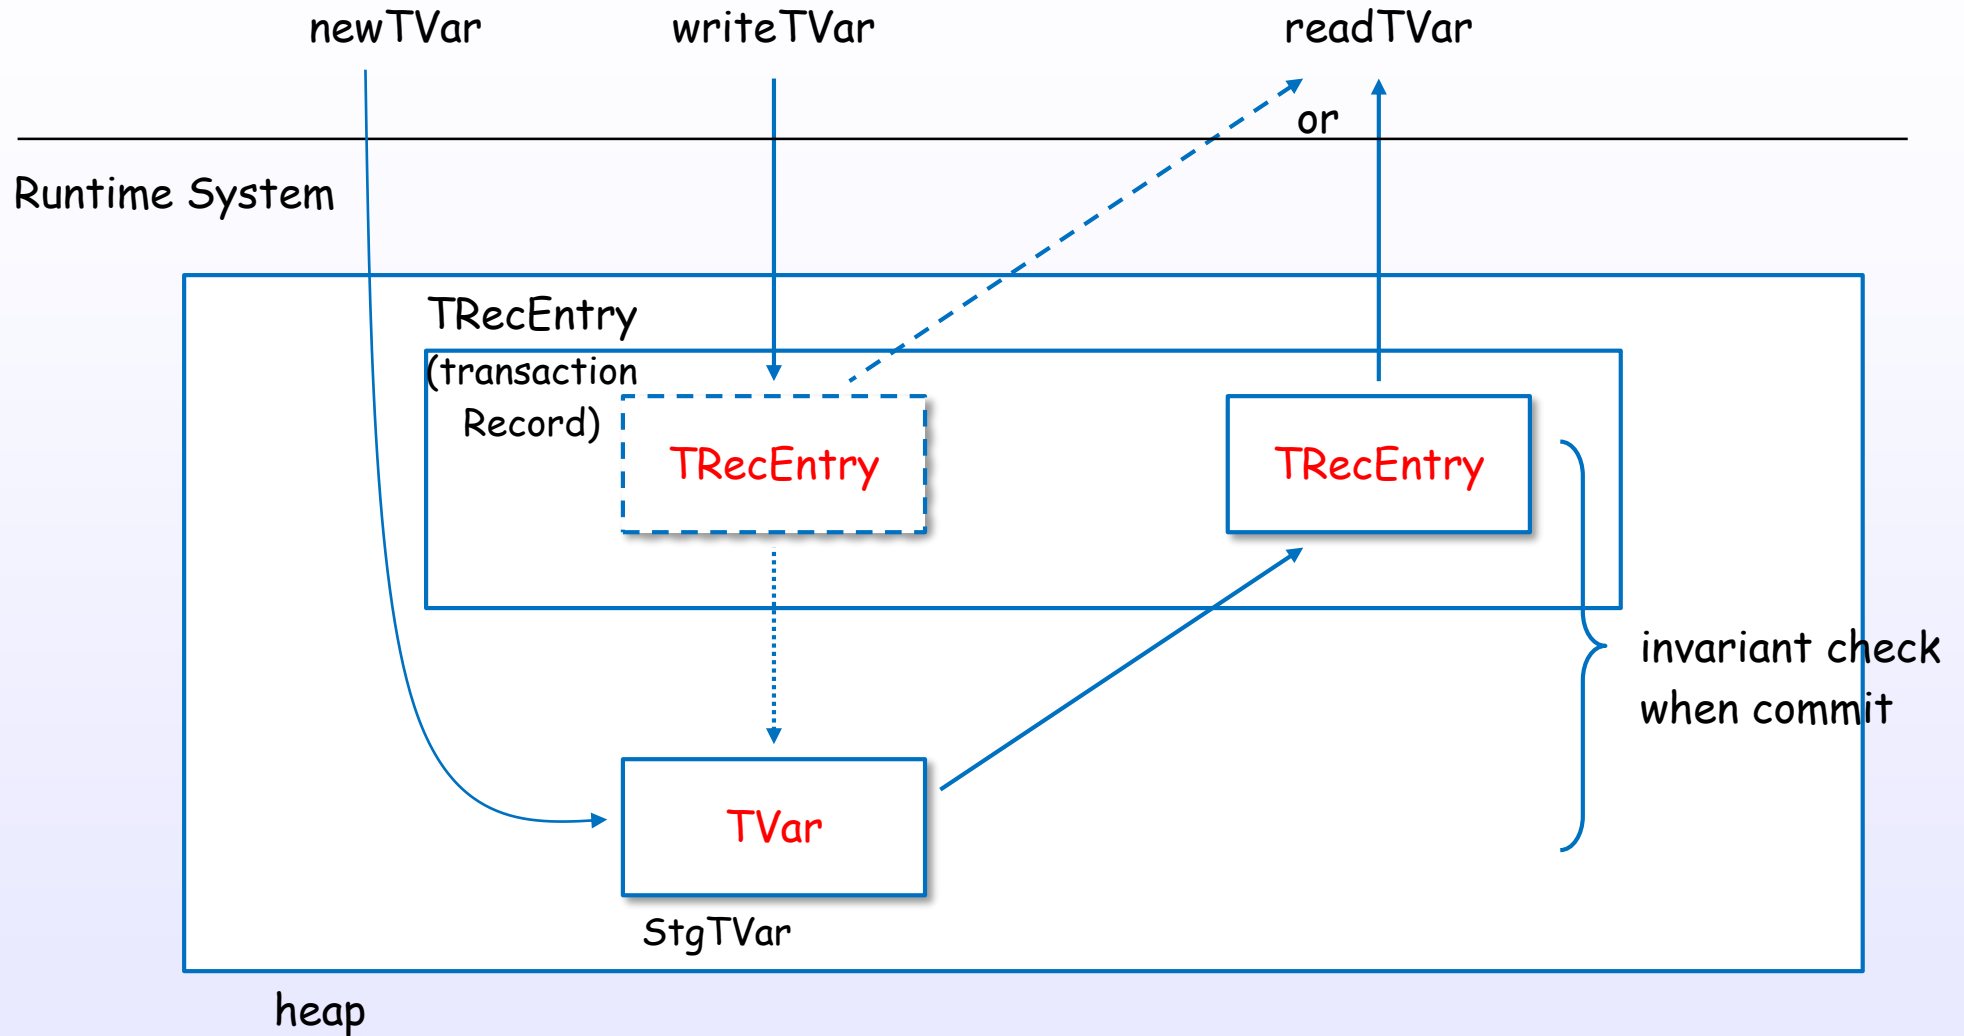
# retry example



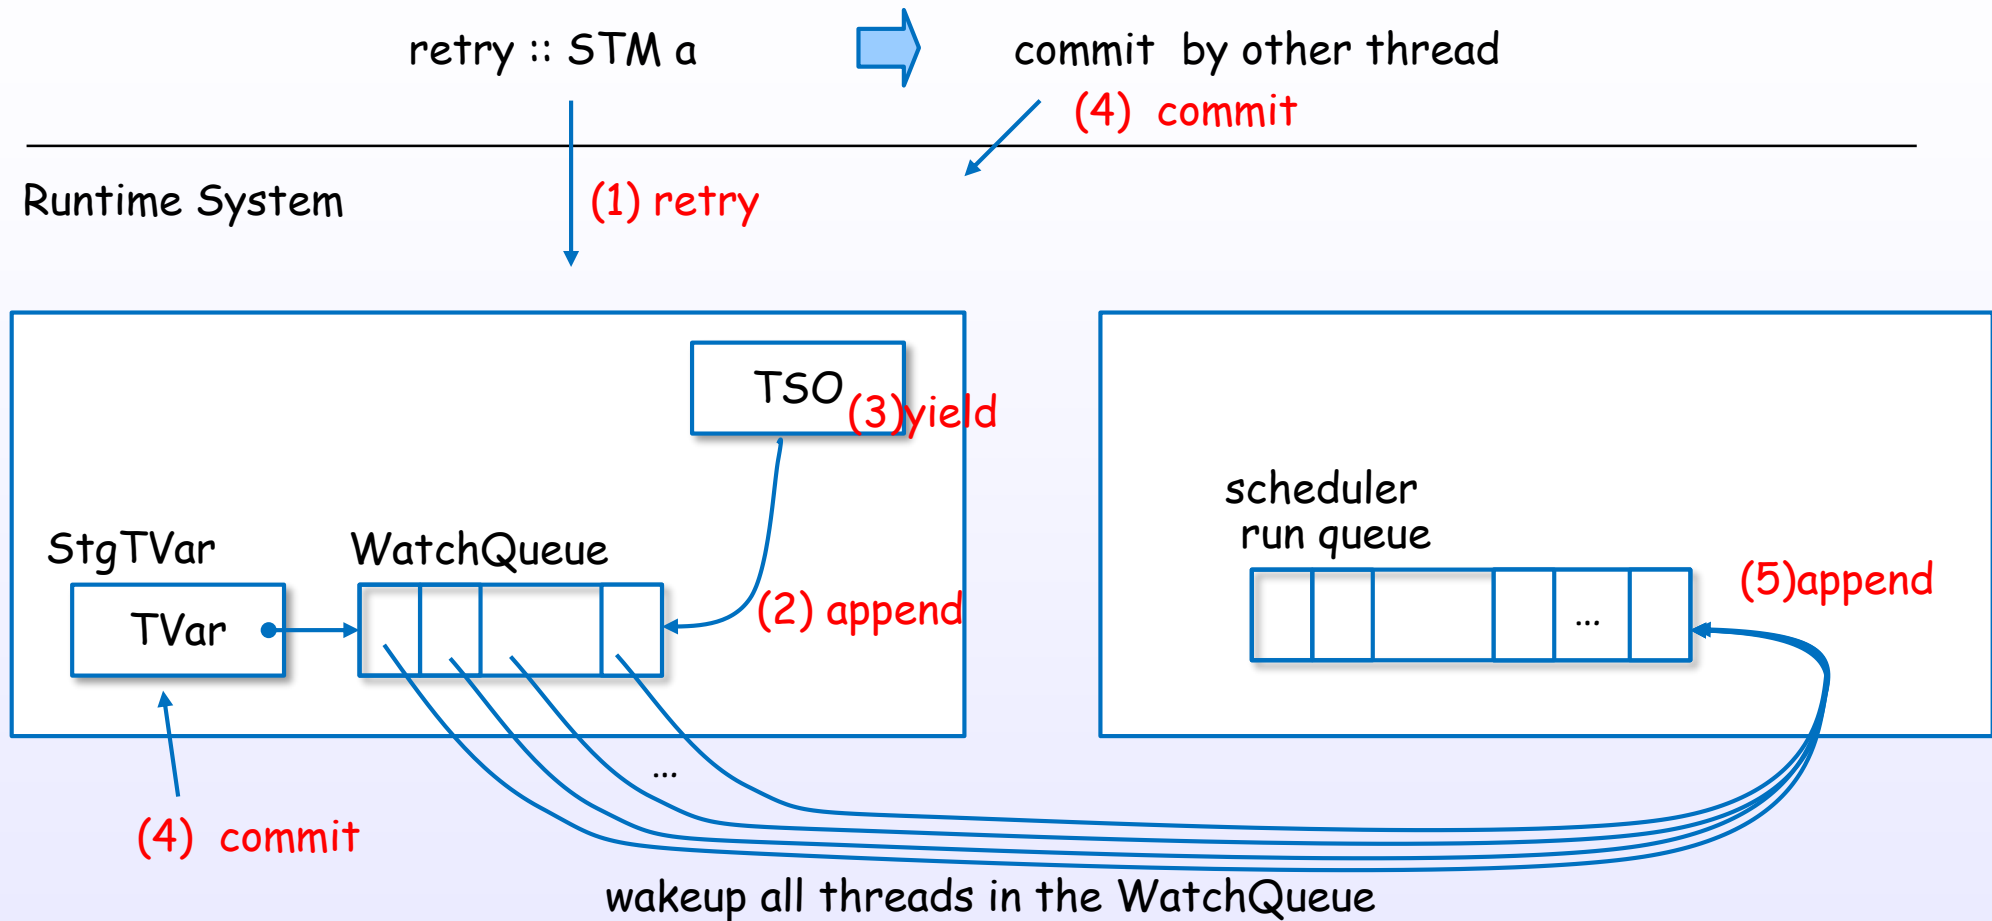
# STM, TVar data structure



# newTVar, writeTVar, readTVar



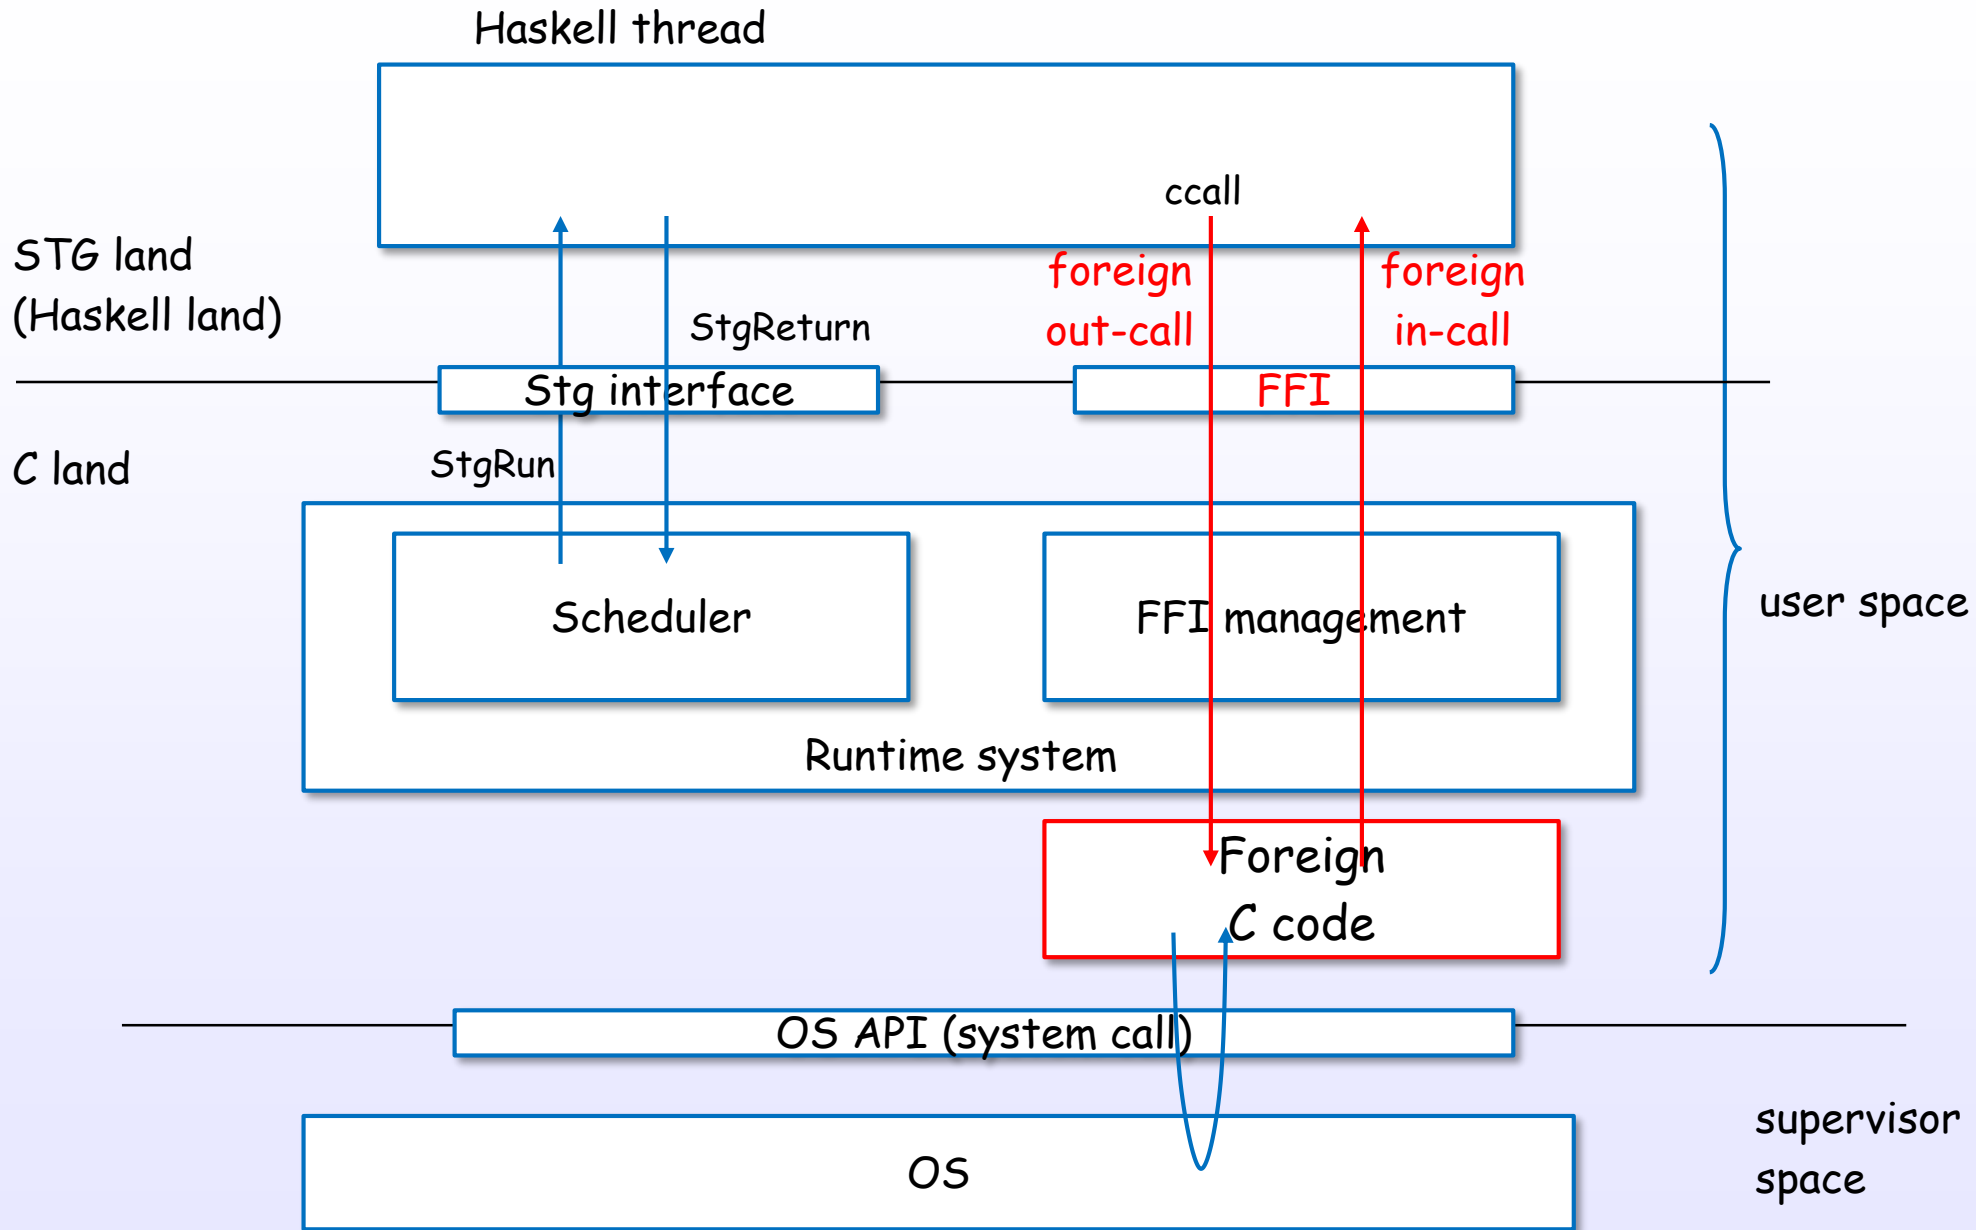
# block by retry, wake up by commit



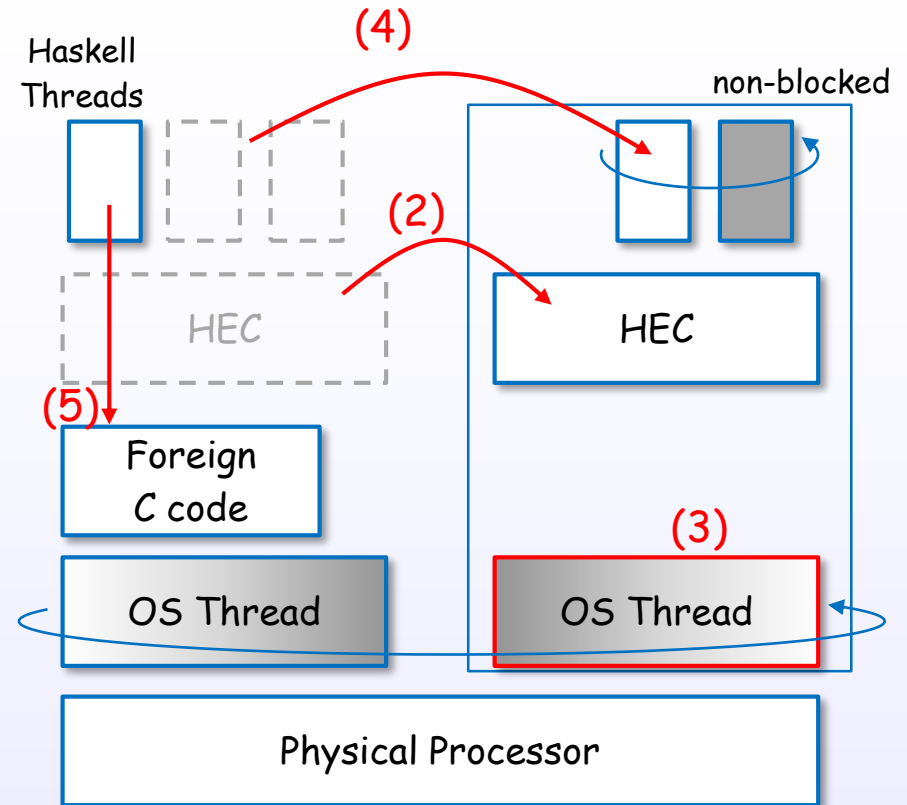
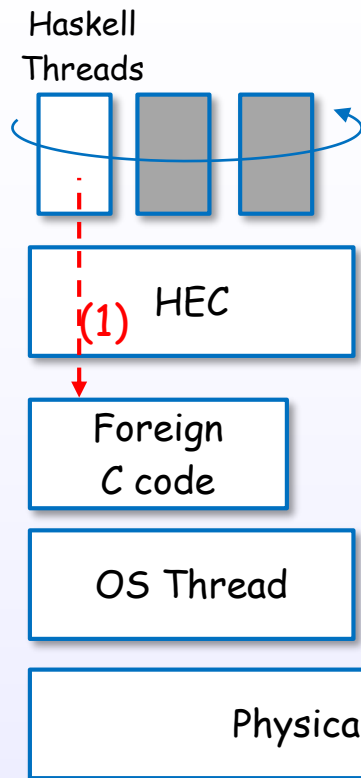
no guarantee of fairness,  
because the RTS has to run all the blocked transaction.

FFI

# FFI (Foreign Function Interface)



# FFI and OS Threads



(1) a safe foreign call (FFI)

(2) move the HEC to other OS thread

(3) spawn or draw an OS thread

(4) move Haskell threads

(5) call the foreign C code

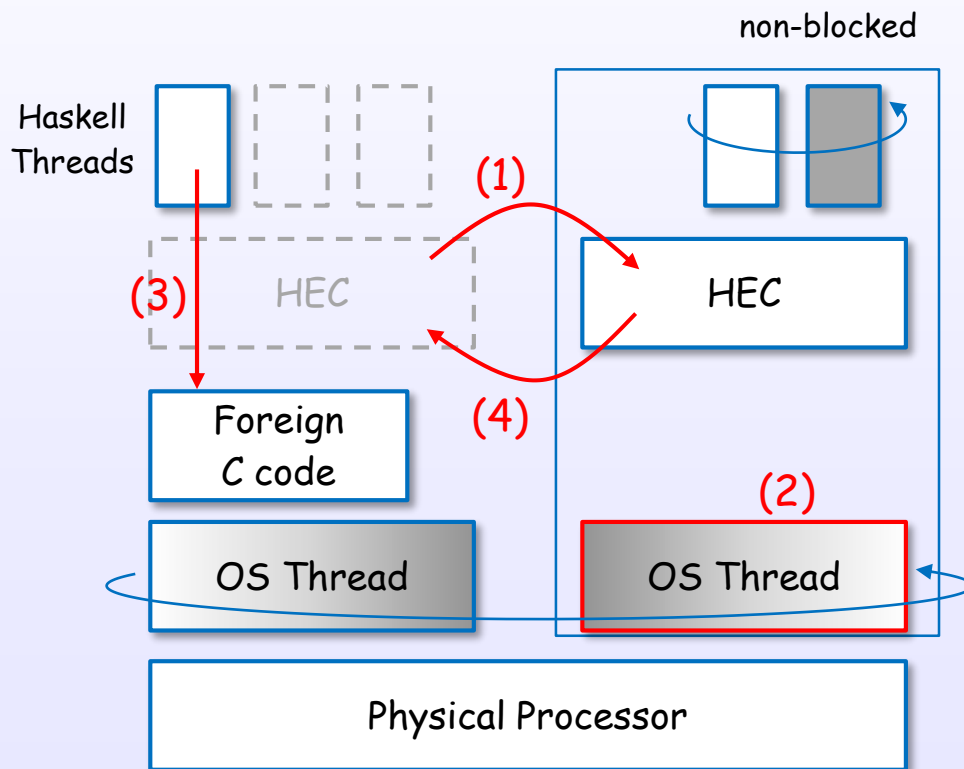
# A safe foreign call (code)

## Haskell Threads

```
ccall suspendThread  
ccall FOREIGN_C_CODE ... (3)  
ccall resumeThread
```

```
releaseCapability_  
giveCapabilityToTask ... (1)  
startWorkerTask  
createOSThread ... (2)
```

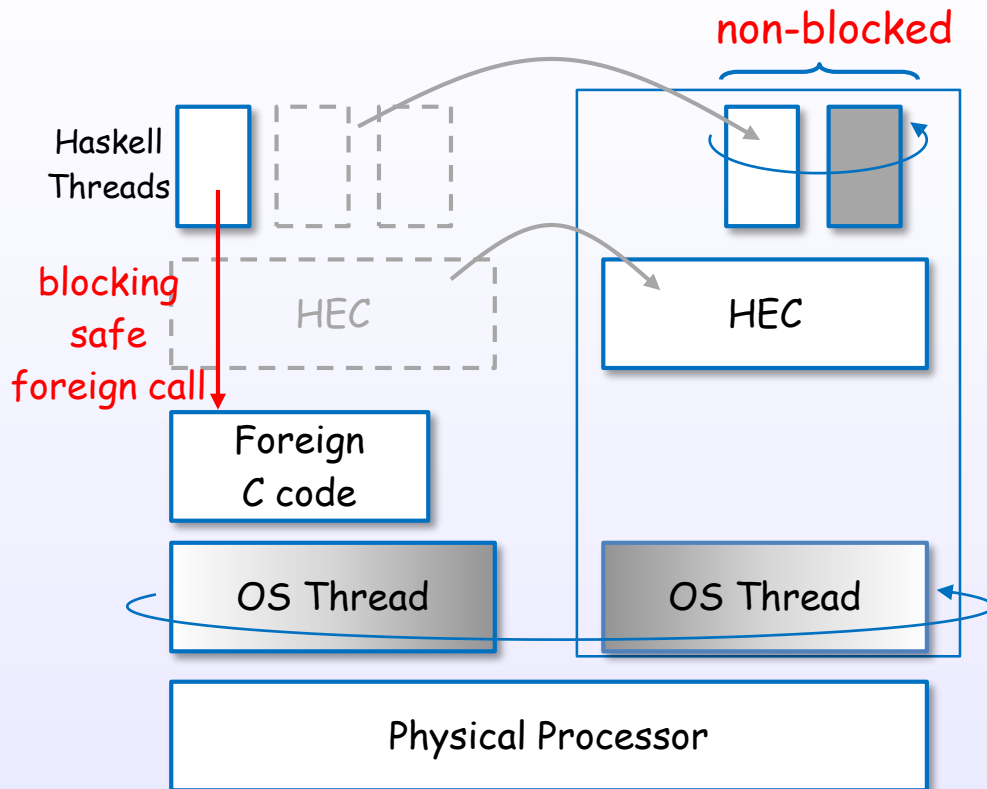
```
waitForReturnCapability ... (4)
```



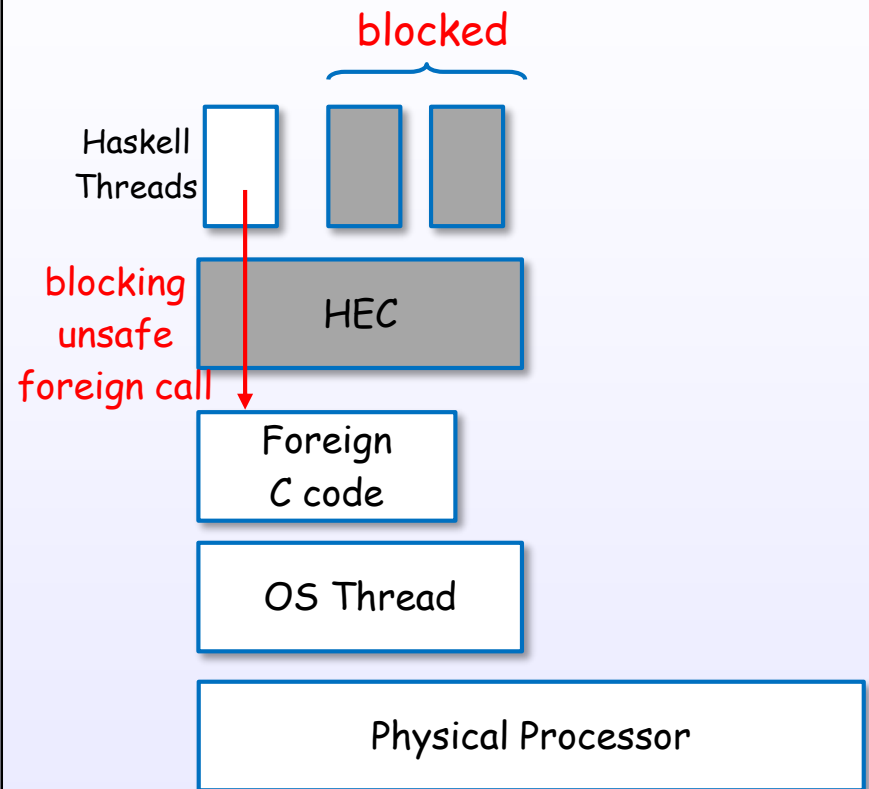


# a safe and an unsafe foreign call

a **safe** foreign call

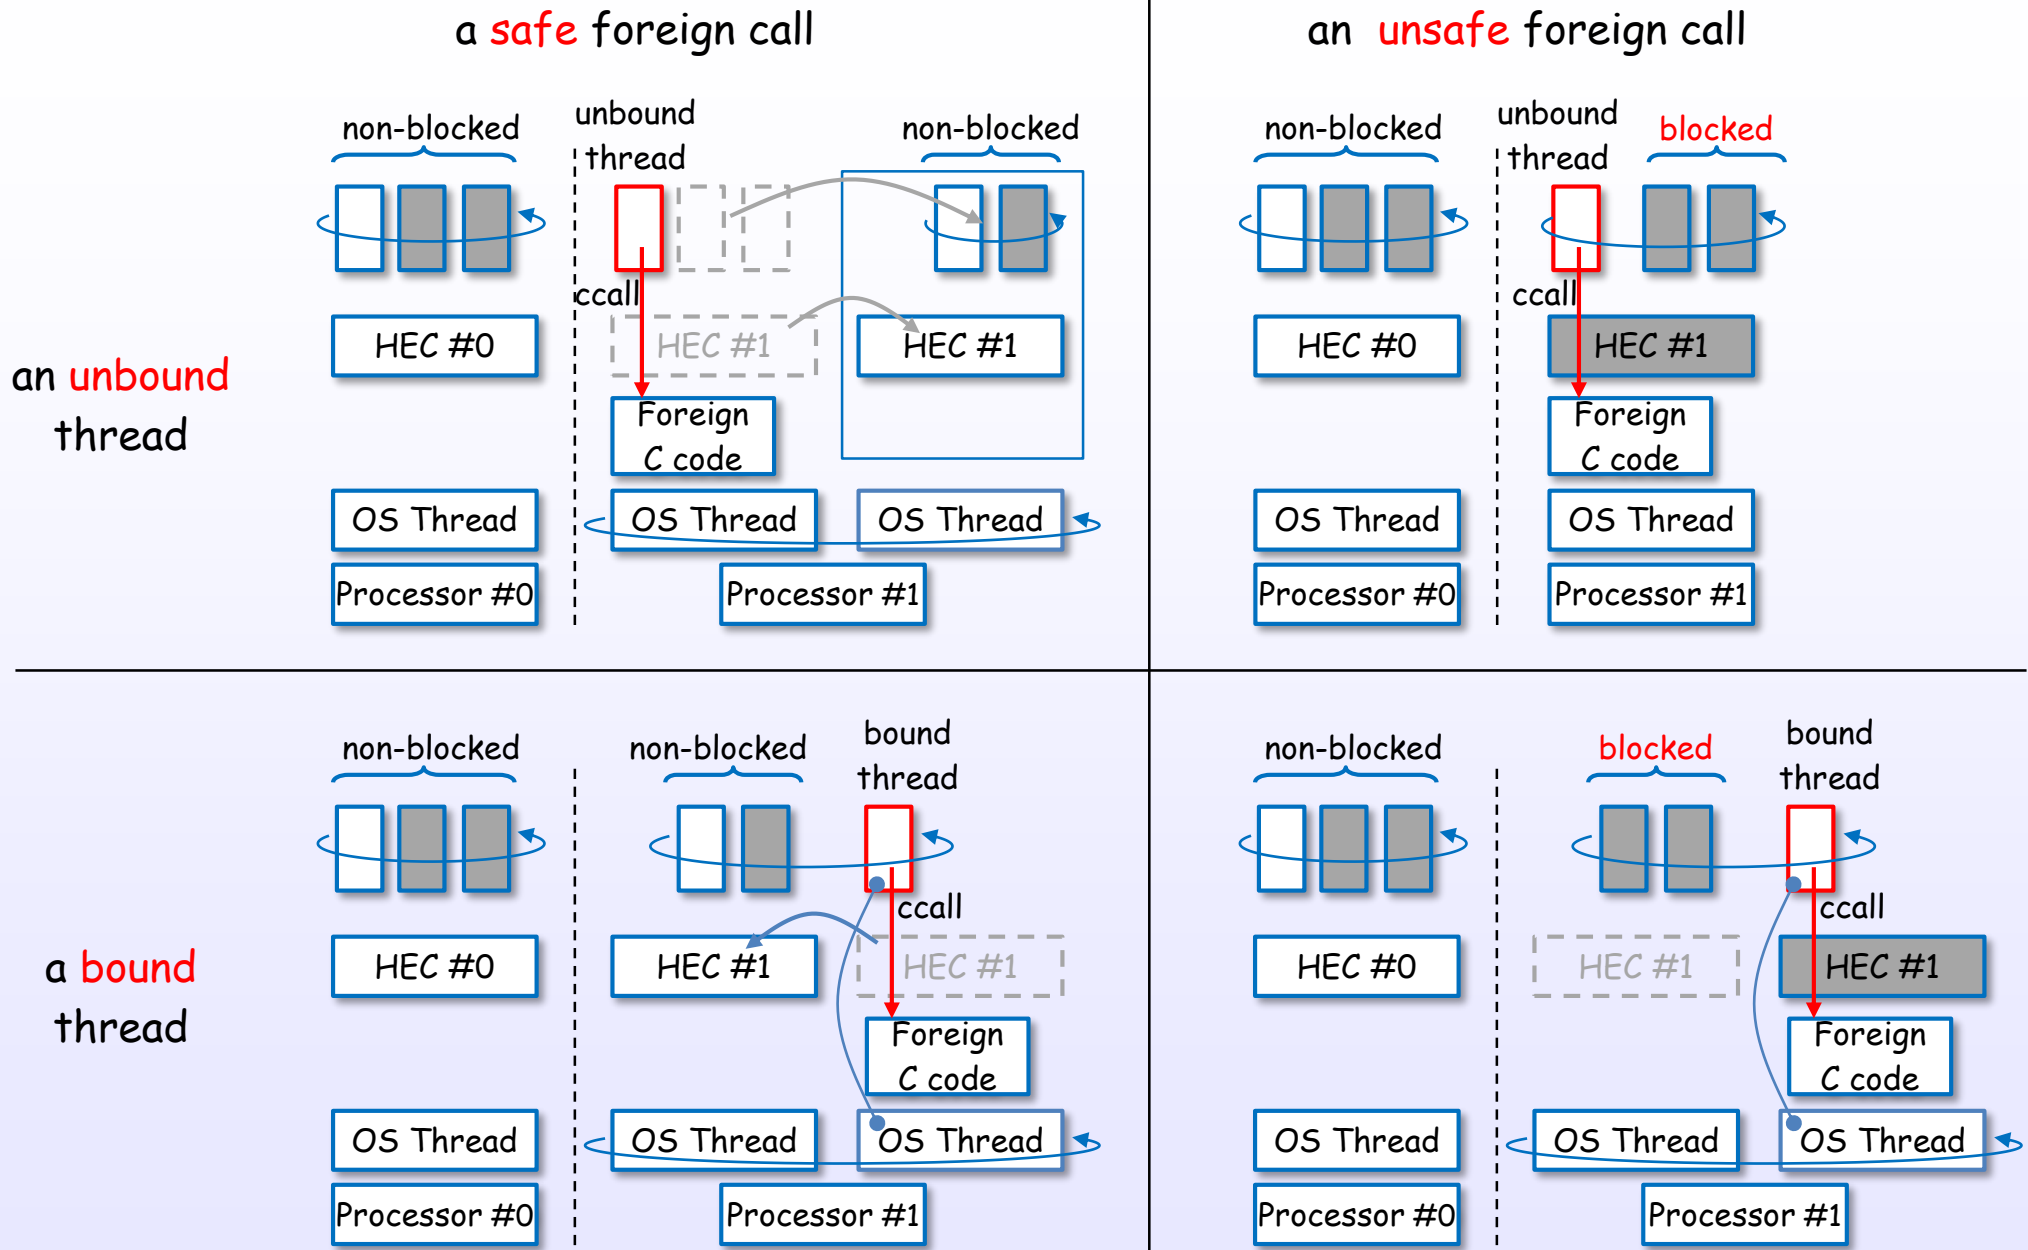


an **unsafe** foreign call



faster,  
but blocking to the other Haskell threads


# Safe/unsafe foreign call and bound/unbound thread



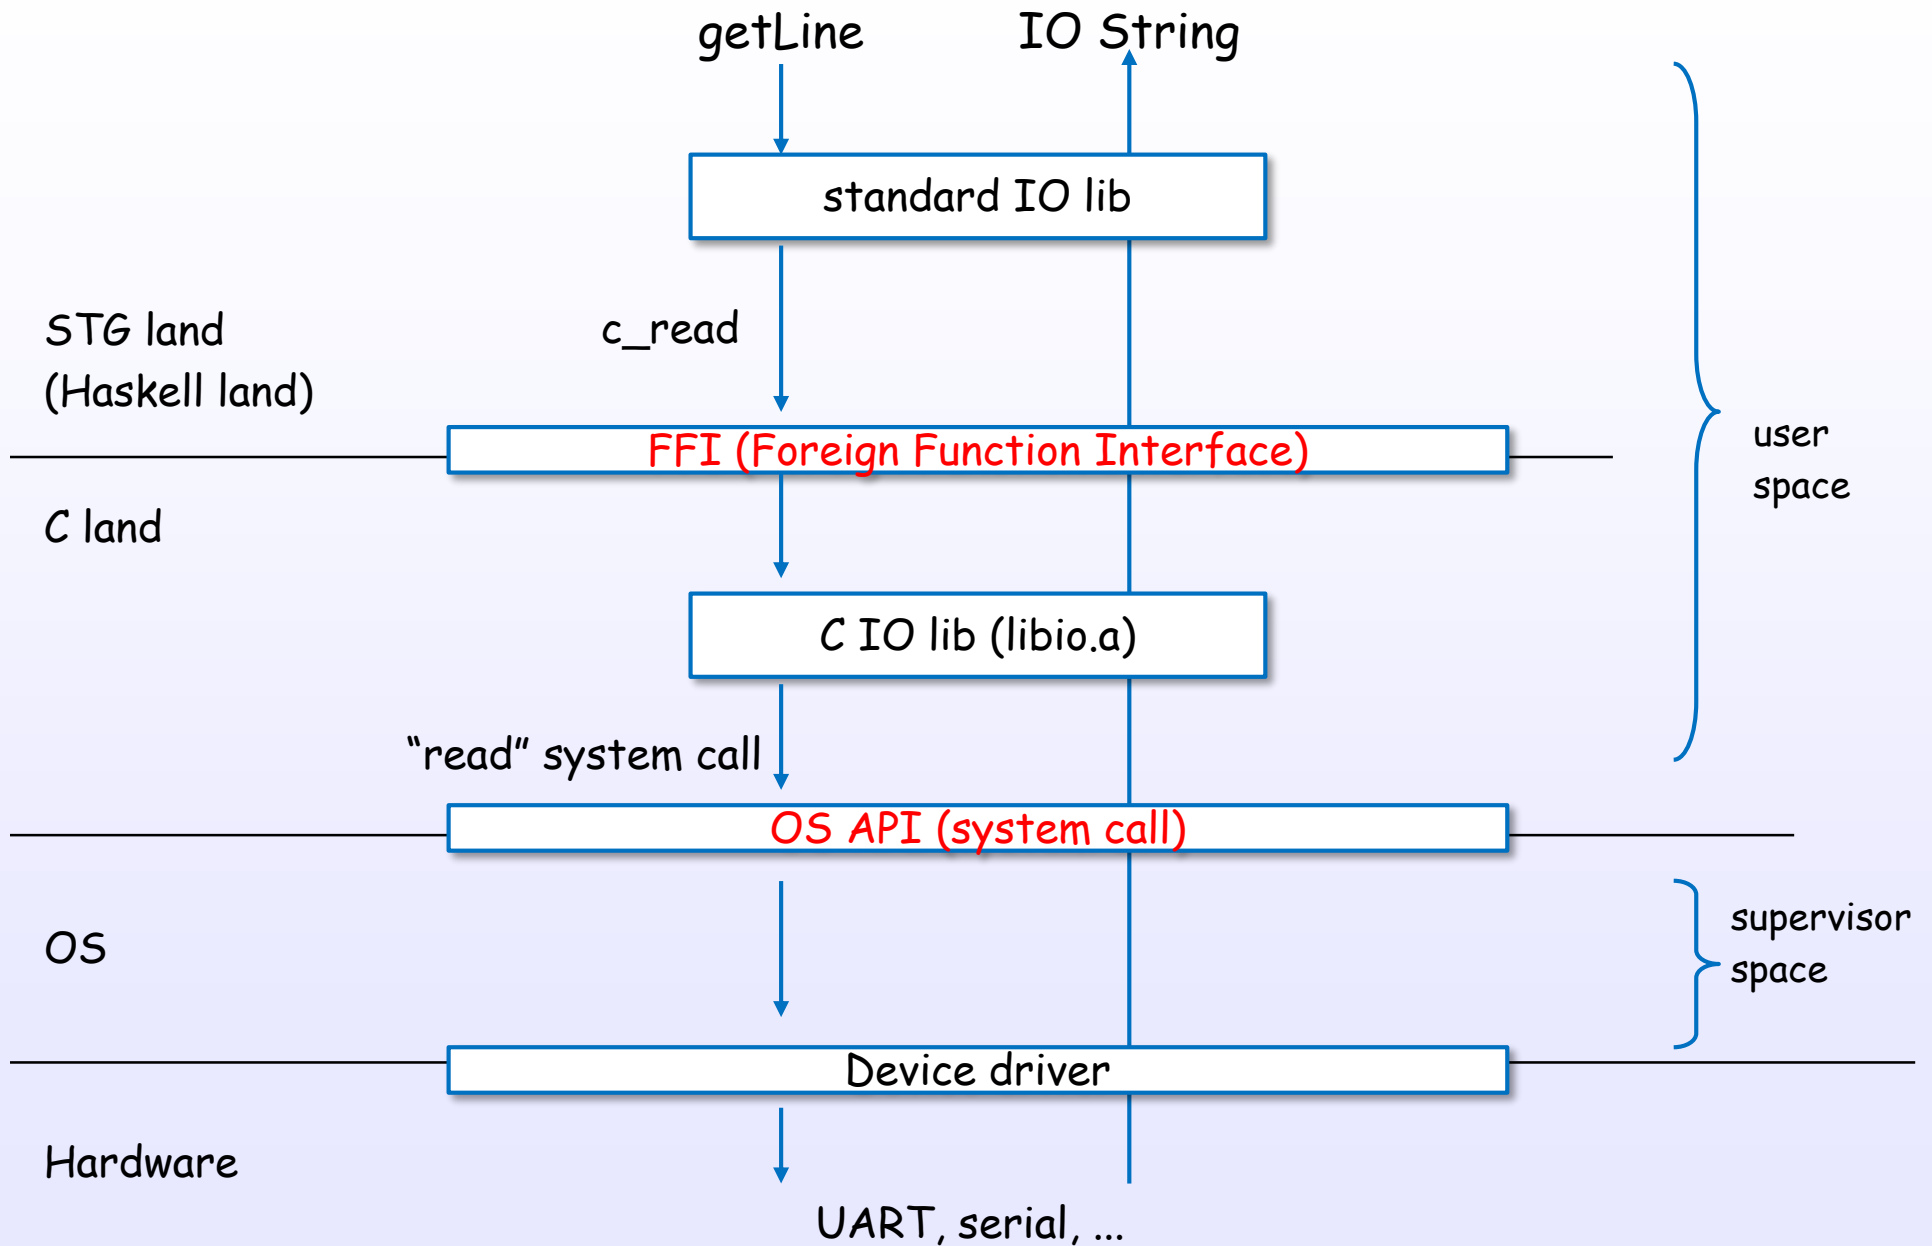
IO and FFI

# IO

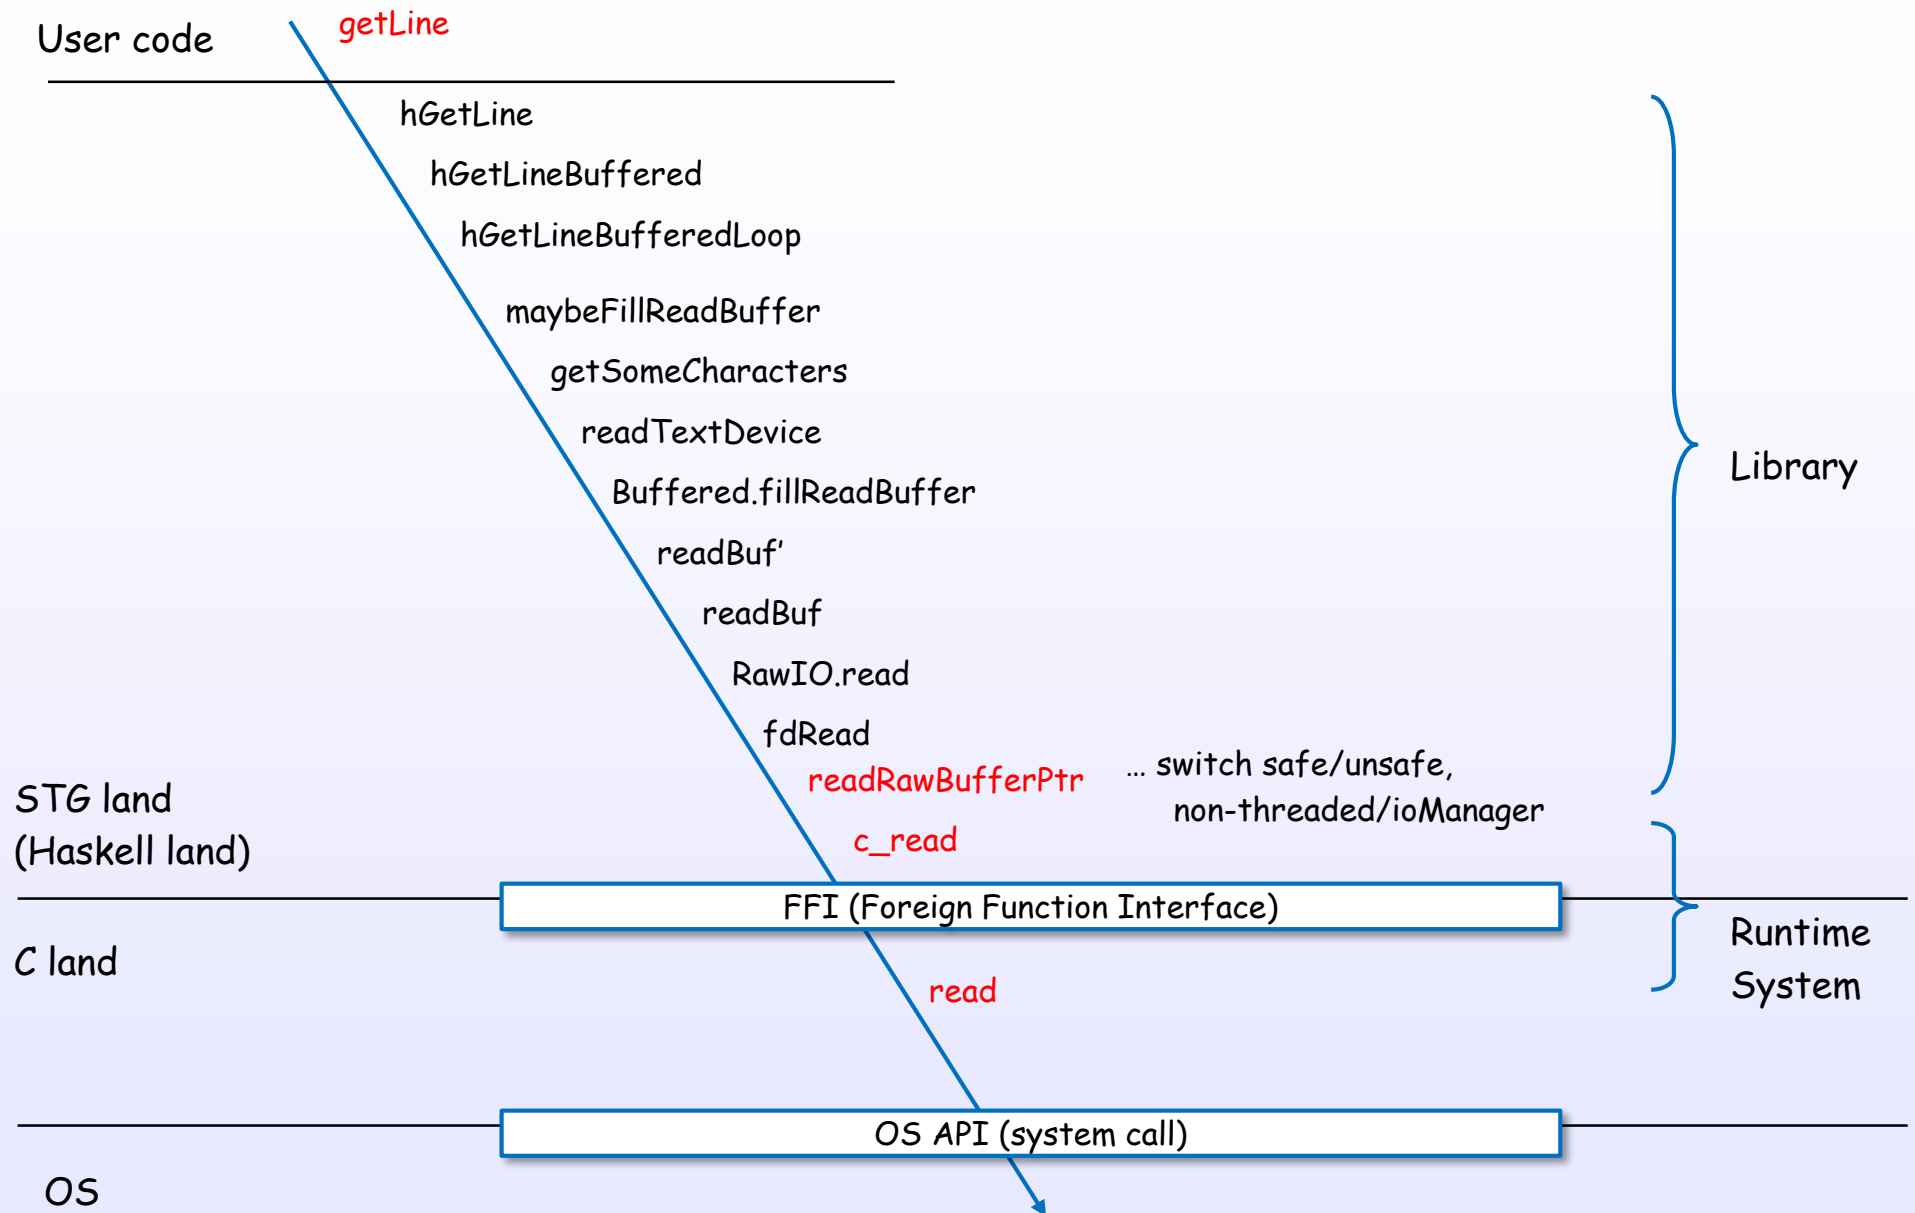
## Haskell Thread



# IO example: getLine

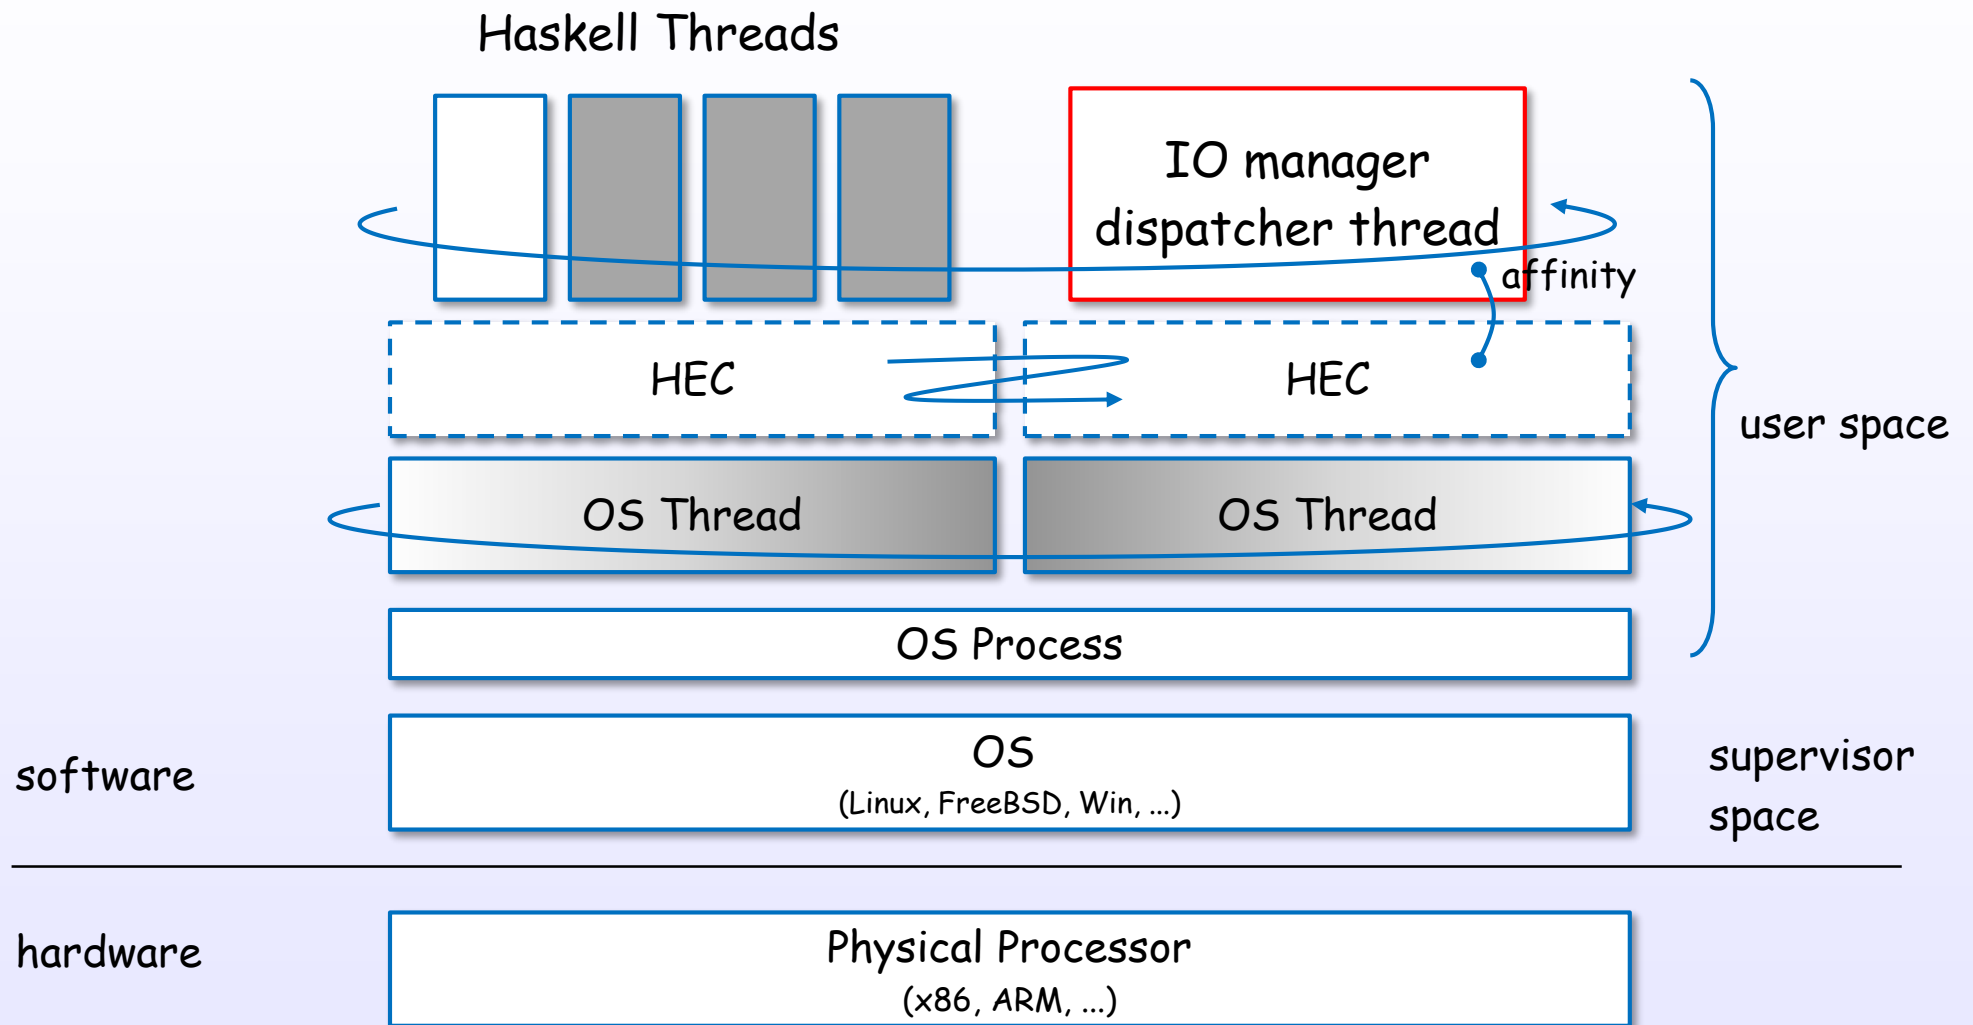


# IO example: getLine (code)



IO manager

# IO manager (single core)

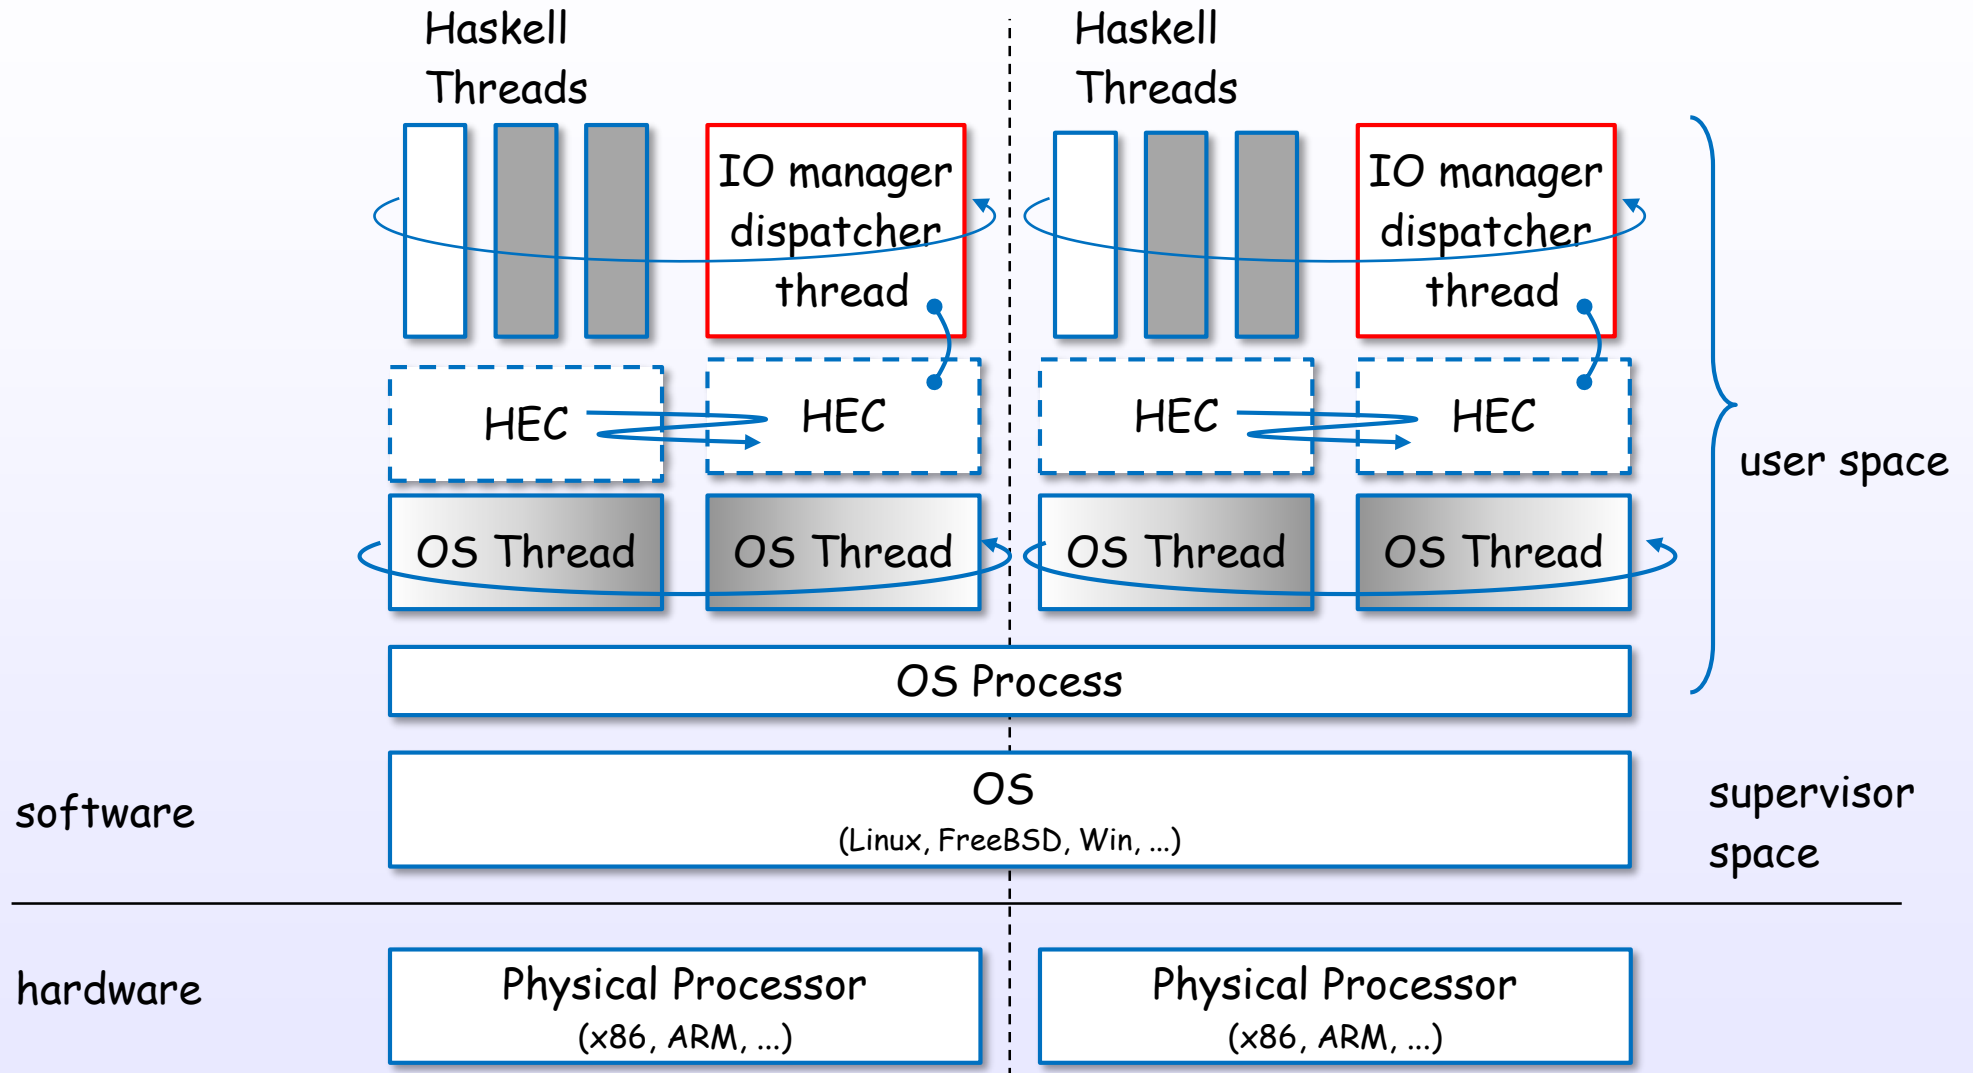


\*Threaded option case (ghc -threaded)

References : [7], [5], [8]

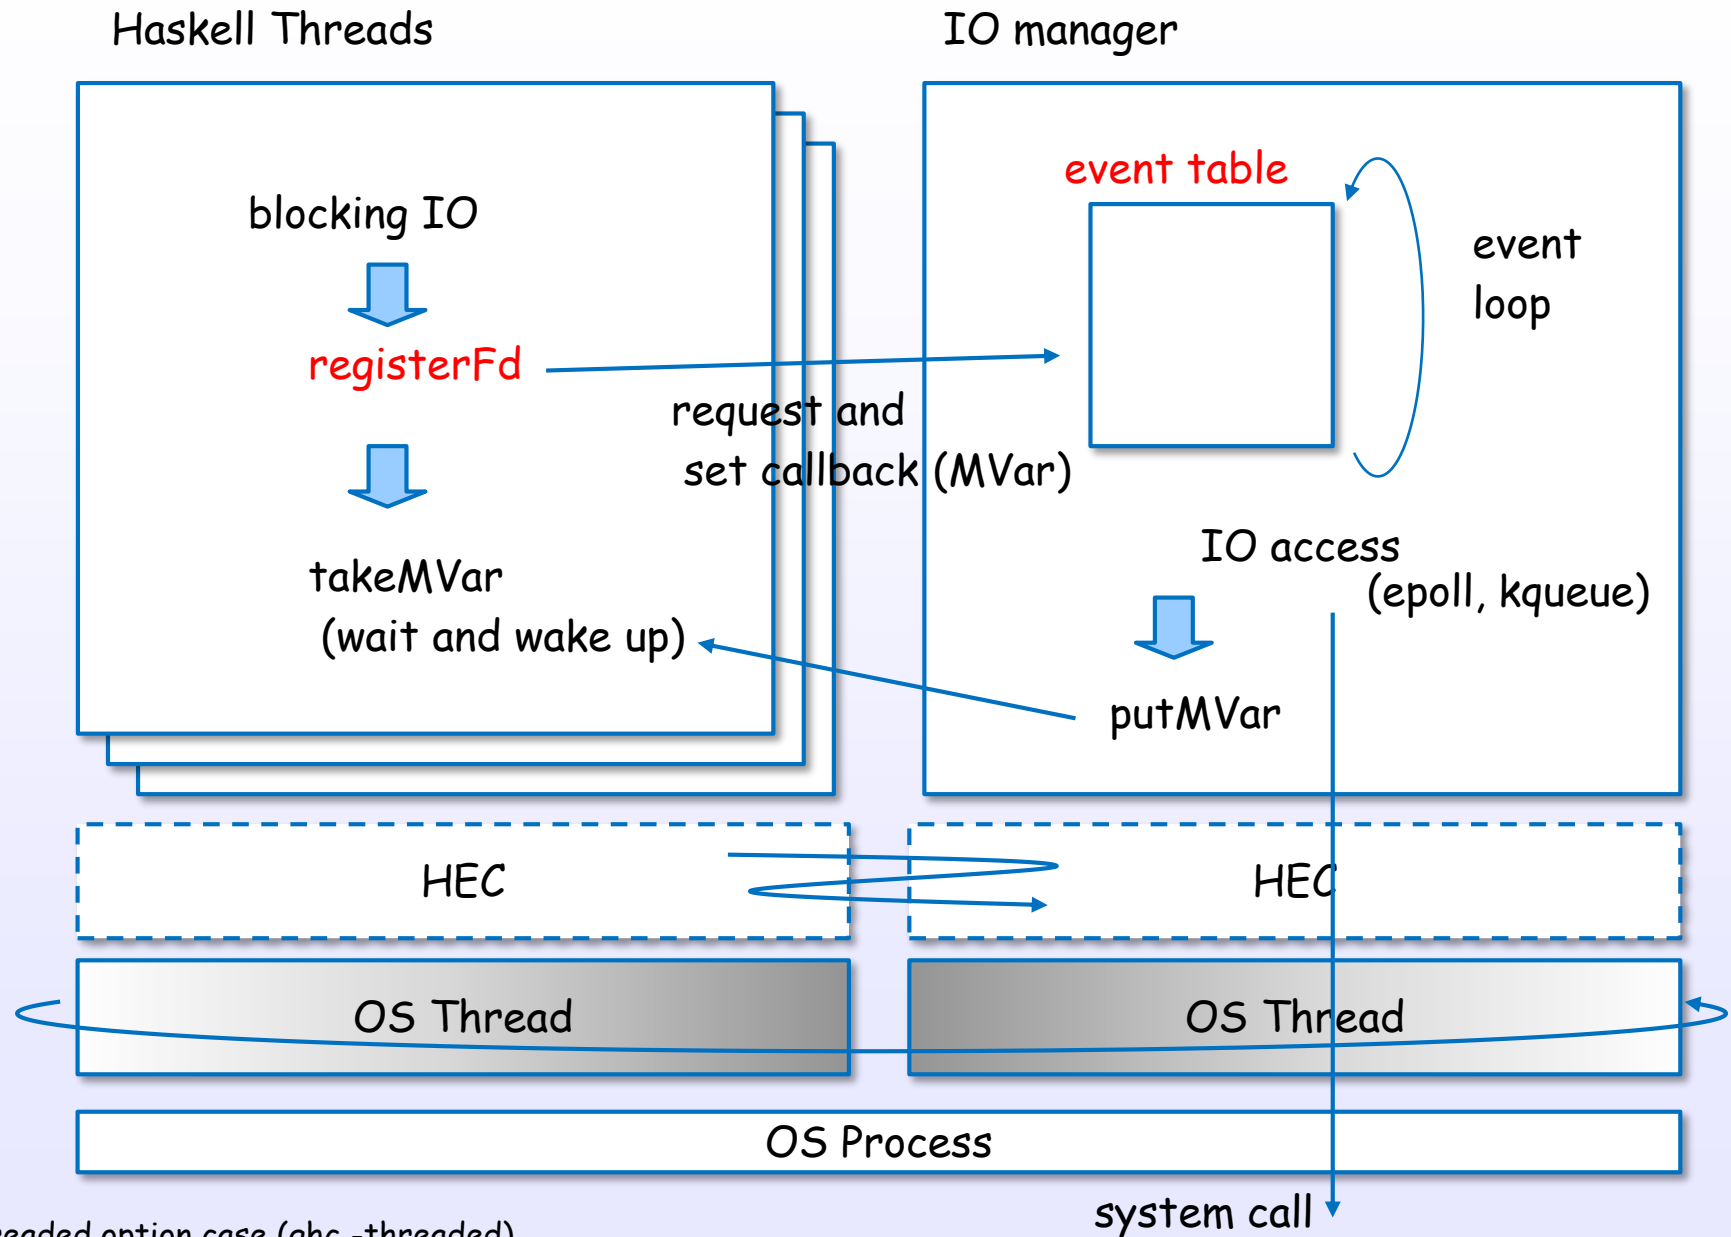


# IO manager (multi core)



\*Threaded option case (ghc -threaded)

# IO manager

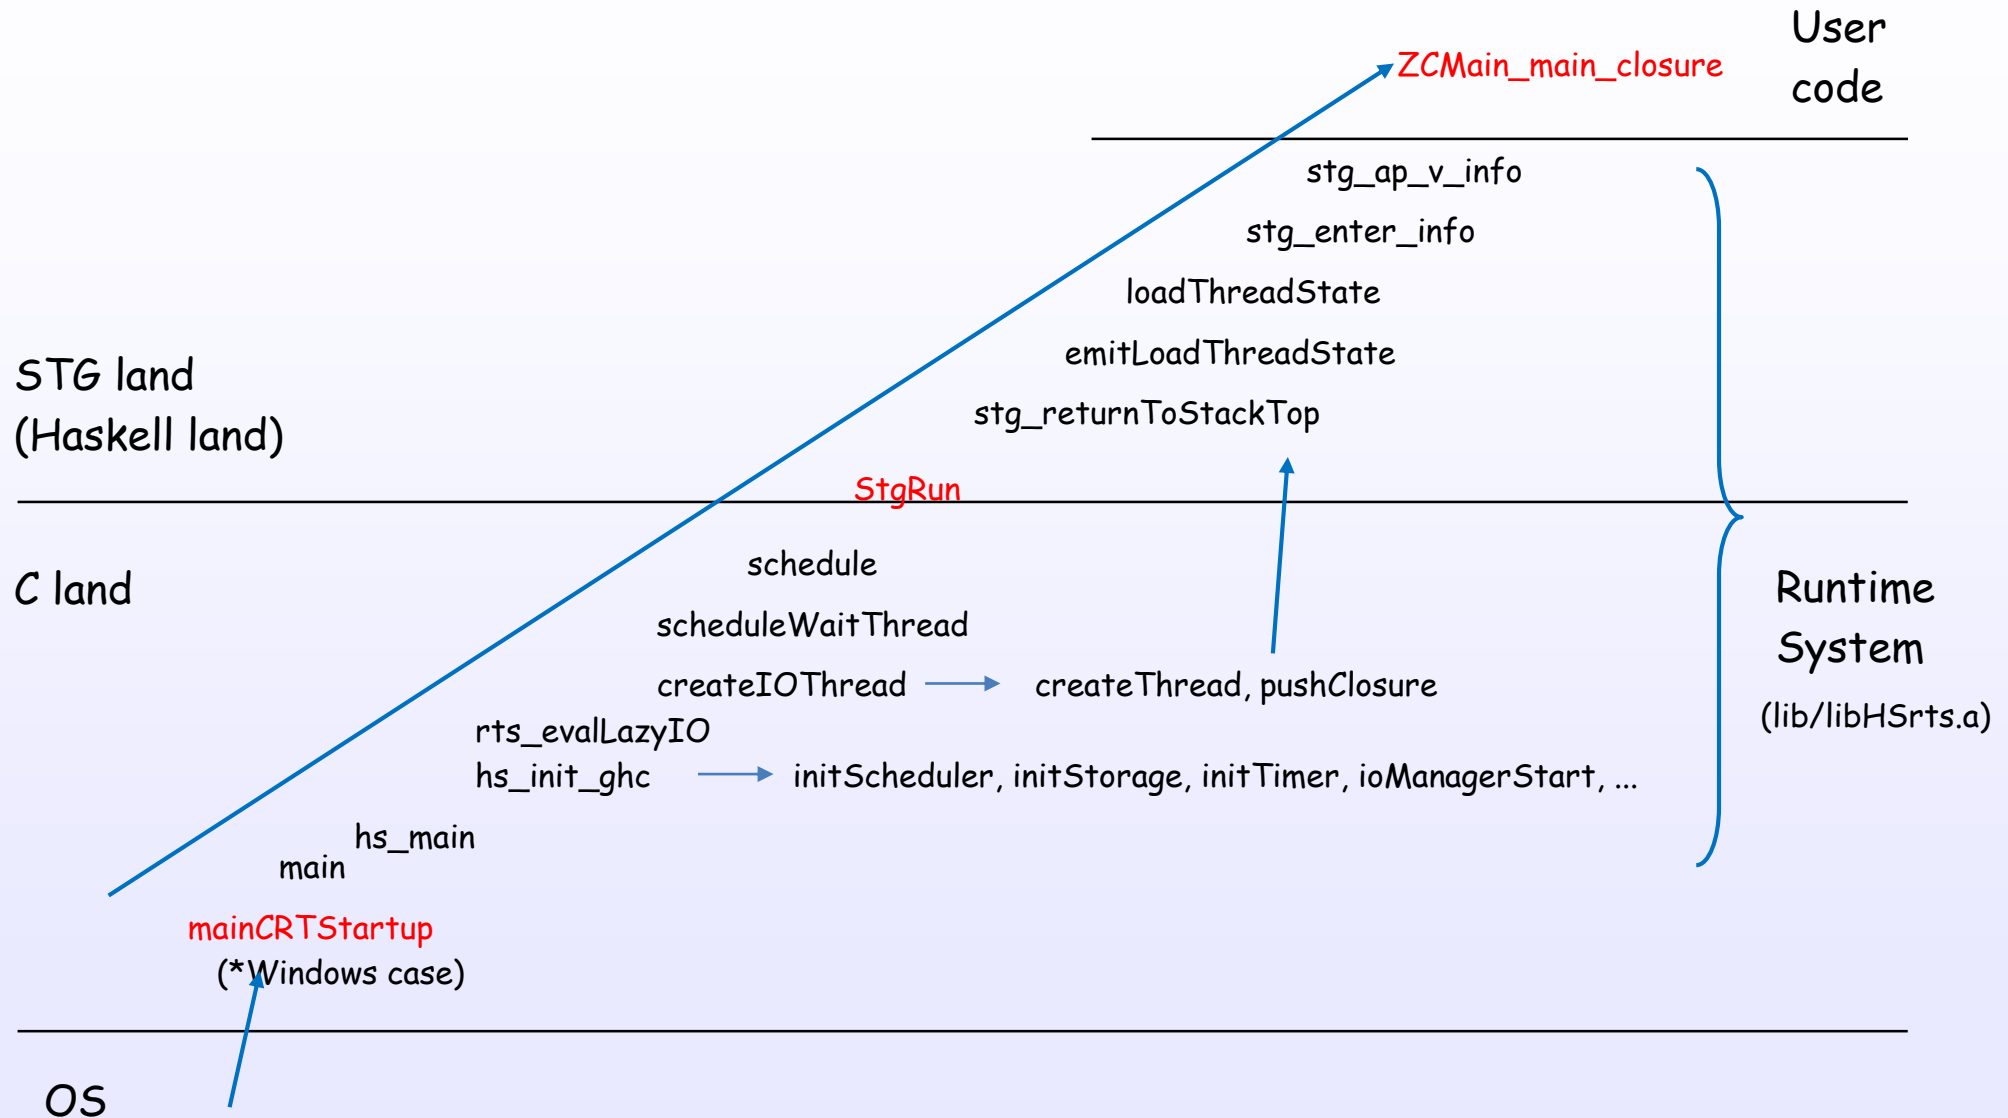


\*Threaded option case (ghc -threaded)

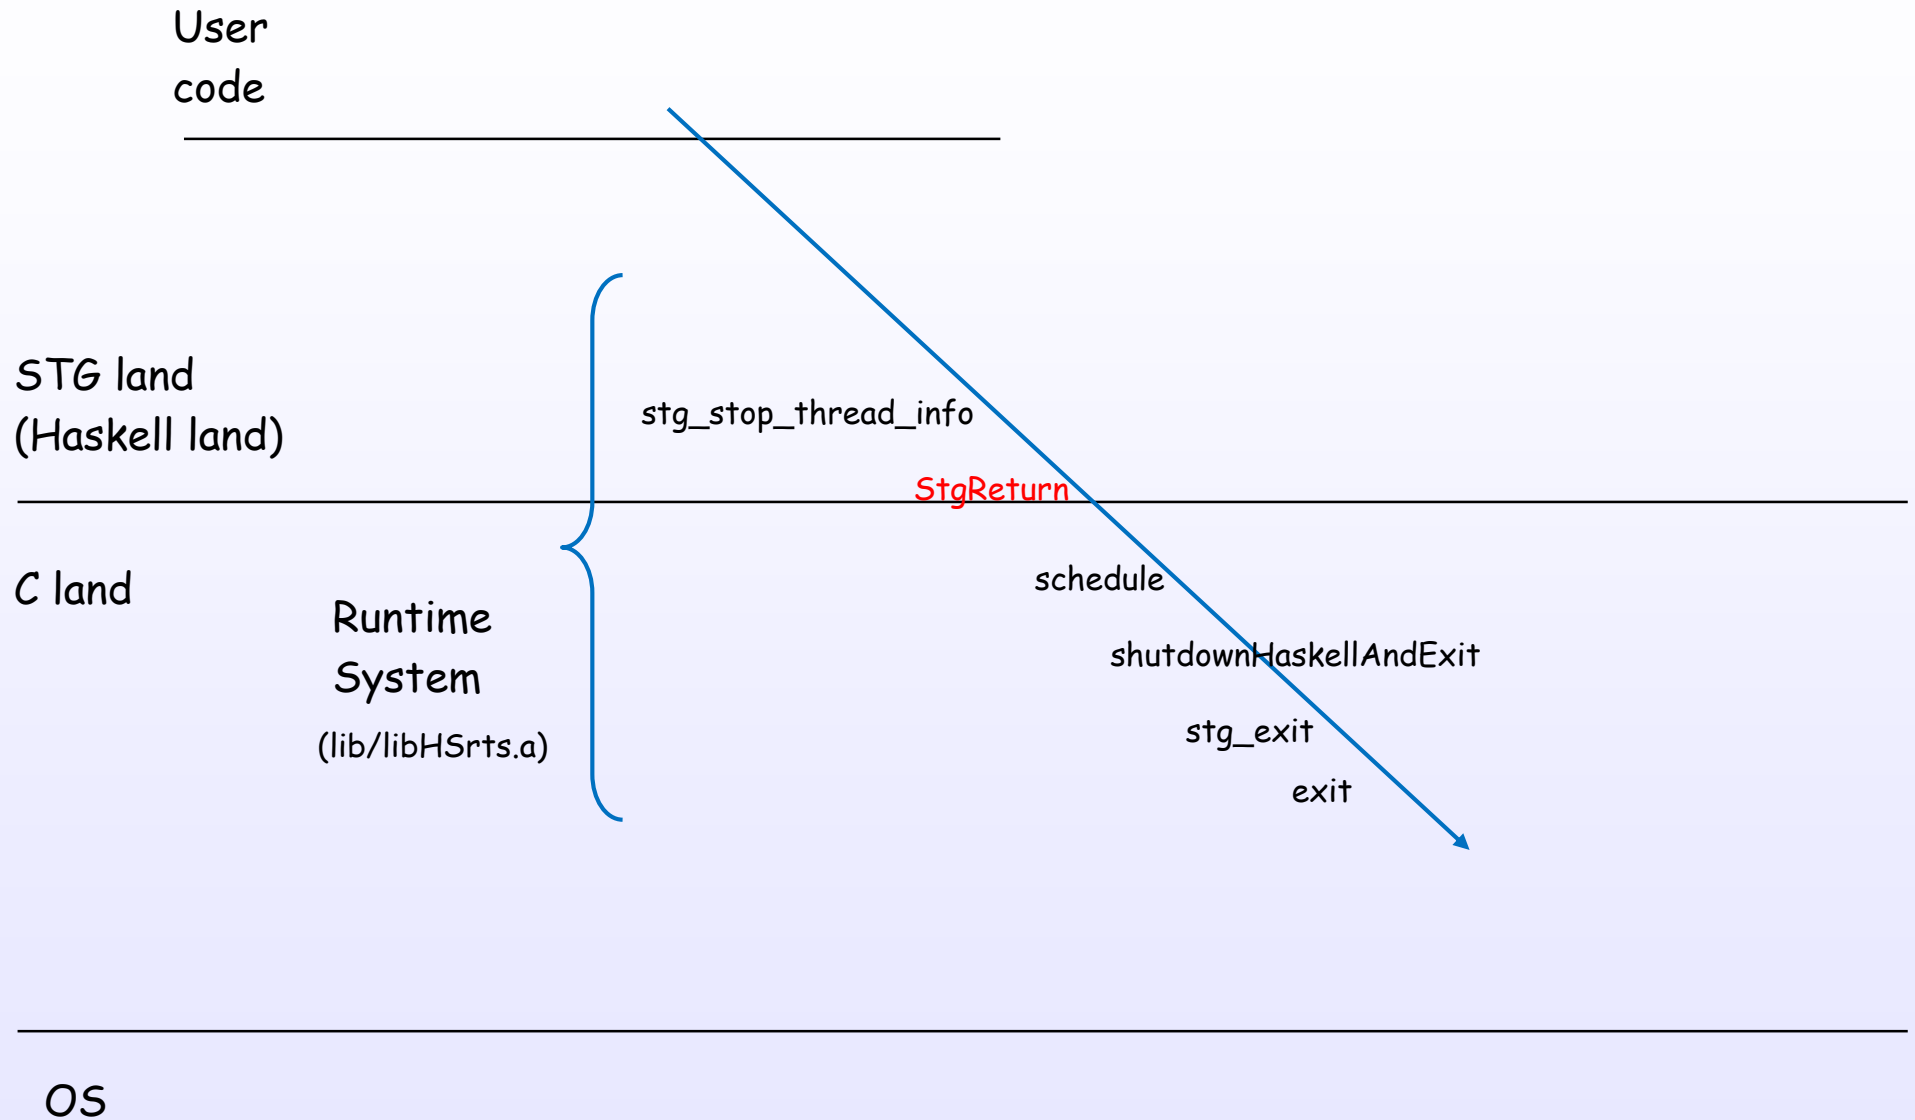
References : [7], [5], [8], [S29], [S30], [S32], [S37], [S35], [S3]

Bootstrap

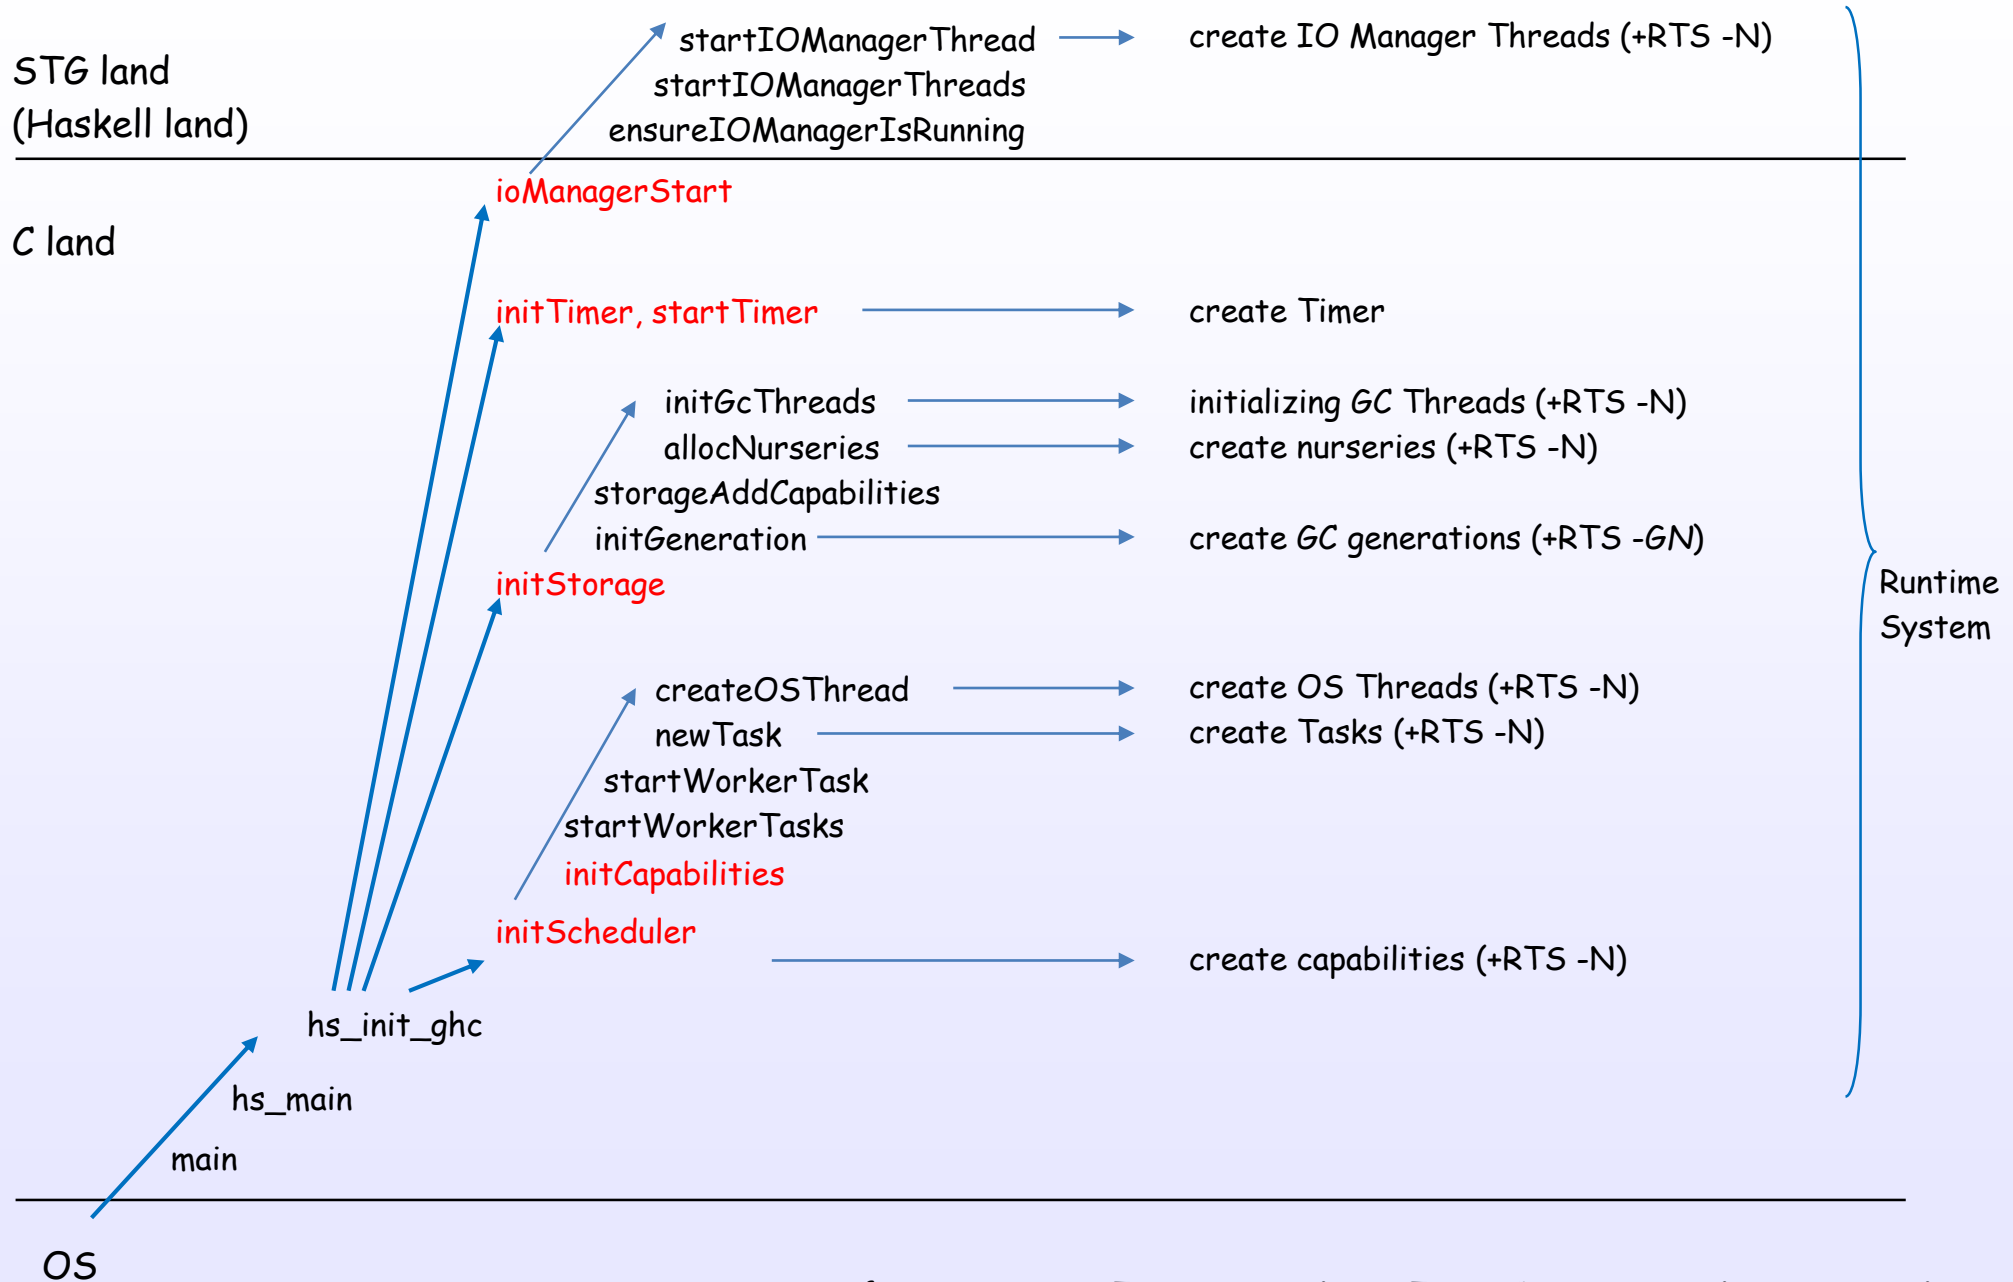
# Bootstrap sequence



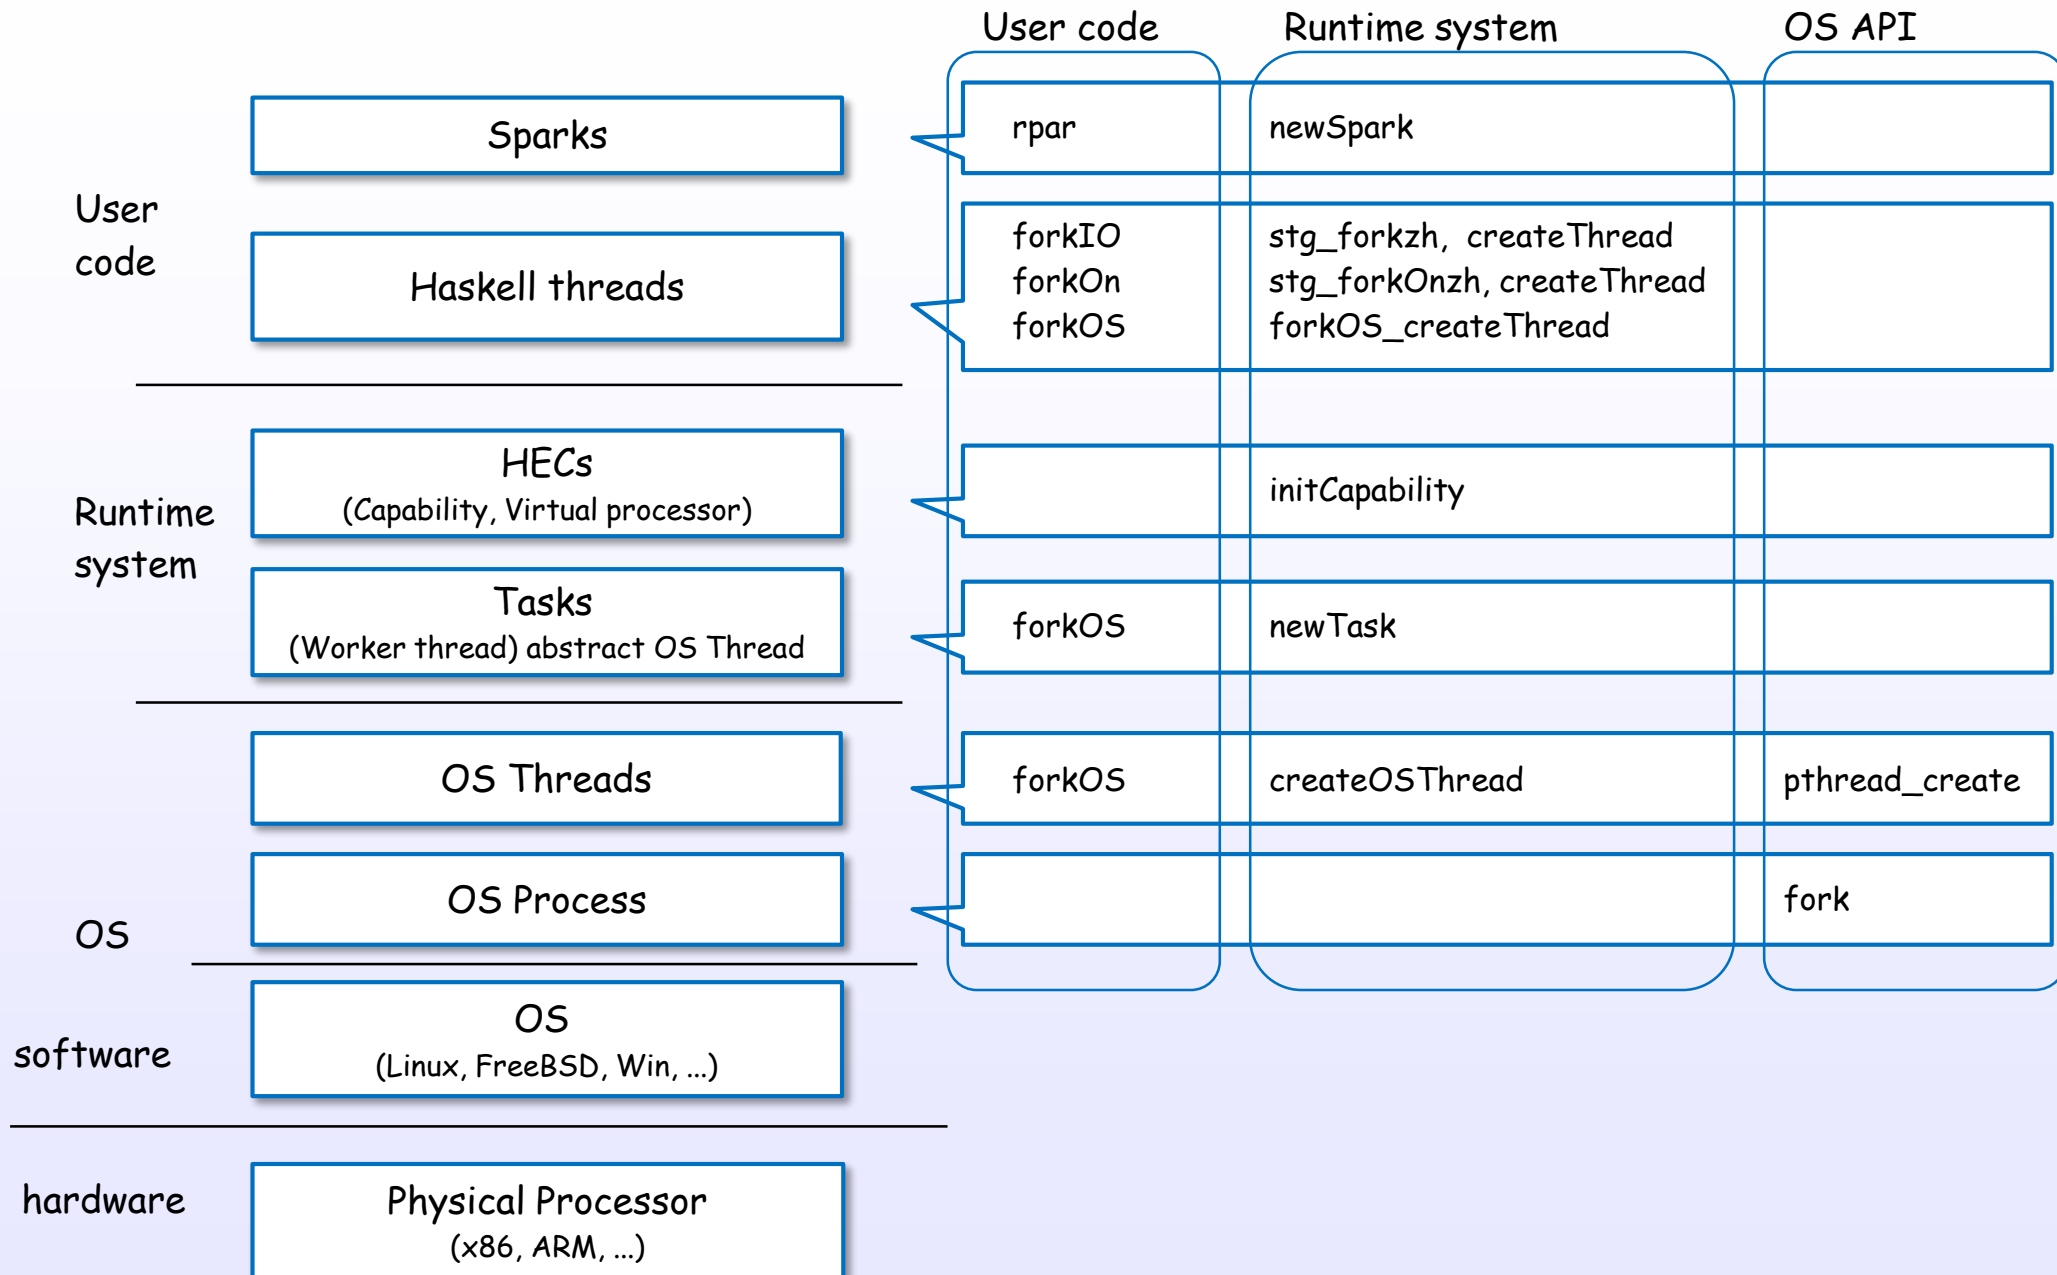
# Exit sequence



# Initializing



# Create each layers



## References



# References

- [1] The Glorious Glasgow Haskell Compilation System User's Guide  
[https://downloads.haskell.org/~ghc/latest/docs/html/users\\_guide/index.html](https://downloads.haskell.org/~ghc/latest/docs/html/users_guide/index.html)
- [2] Implementing lazy functional languages on stock hardware: the Spineless Tagless G-machine Version 2.5  
<http://research.microsoft.com/en-us/um/people/simonpj/Papers/spineless-tagless-gmachine.ps.gz>
- [3] Making a Fast Curry Push/Enter vs Eval/Apply for Higher-order Languages  
<http://research.microsoft.com/en-us/um/people/simonpj/papers/eval-apply/>
- [4] Faster Laziness Using Dynamic Pointer Tagging  
<http://research.microsoft.com/en-us/um/people/simonpj/papers/ptr-tag/ptr-tagging.pdf>
- [5] Runtime Support for Multicore Haskell  
<http://research.microsoft.com/en-us/um/people/simonpj/papers/parallel/multicore-ghc.pdf>
- [6] Extending the Haskell Foreign Function Interface with Concurrency  
<http://community.haskell.org/~simonmar/papers/conc-ffi.pdf>
- [7] Mio: A High-Performance Multicore IO Manager for GHC  
<http://haskell.cs.yale.edu/wp-content/uploads/2013/08/hask035-voellmy.pdf>
- [8] The GHC Runtime System  
<http://web.mit.edu/~ezyang/Public/jfp-ghc-rts.pdf>
- [9] The GHC Runtime System  
<http://www.scs.stanford.edu/14sp-cs240h/slides/ghc-rts.pdf>
- [10] Evaluation on the Haskell Heap  
<http://blog.ezyang.com/2011/04/evaluation-on-the-haskell-heap/>

# References

- [11] IO evaluates the Haskell Heap  
<http://blog.ezyang.com/2011/04/io-evaluates-the-haskell-heap/>
- [12] Understanding the Stack  
<http://www.well-typed.com/blog/94/>
- [13] Understanding the RealWorld  
<http://www.well-typed.com/blog/95/>
- [14] The GHC scheduler  
<http://blog.ezyang.com/2013/01/the-ghc-scheduler/>
- [15] GHC's Garbage Collector  
[http://www.mm-net.org.uk/workshop190404/GHC's\\_Garbage\\_Collector.ppt](http://www.mm-net.org.uk/workshop190404/GHC's_Garbage_Collector.ppt)
- [16] Concurrent Haskell  
<http://www.haskell.org/ghc/docs/papers/concurrent-haskell.ps.gz>
- [17] Beautiful Concurrency  
<https://www.fpcomplete.com/school/advanced-haskell/beautiful-concurrency>
- [18] Anatomy of an MVar operation  
<http://blog.ezyang.com/2013/05/anatomy-of-an-mvar-operation/>
- [19] Parallel and Concurrent Programming in Haskell  
<http://community.haskell.org/~simonmar/pcph/>
- [20] Real World Haskell  
<http://book.realworldhaskell.org/>

# References

- [21] A Haskell Compiler  
<http://www.scs.stanford.edu/16wi-cs240h/slides/ghc-compiler-slides.html>
- [22] Dive into GHC  
[http://www.stephendiehl.com/posts/ghc\\_01.html](http://www.stephendiehl.com/posts/ghc_01.html)

# References

## The GHC Commentary

- [C1] <https://ghc.haskell.org/trac/ghc/wiki/Commentary>
- [C2] <https://ghc.haskell.org/trac/ghc/wiki/Commentary/SourceTree>
- [C3] <https://ghc.haskell.org/trac/ghc/wiki/Commentary/Compiler>
- [C4] <https://ghc.haskell.org/trac/ghc/wiki/Commentary/Compiler/HscMain>
- [C5] <https://ghc.haskell.org/trac/ghc/wiki/Commentary/Compiler/CoreSynType>
- [C6] <https://ghc.haskell.org/trac/ghc/wiki/Commentary/Compiler/StgSynType>
- [C7] <https://ghc.haskell.org/trac/ghc/wiki/Commentary/Compiler/CmmType>
- [C8] <https://ghc.haskell.org/trac/ghc/wiki/Commentary/Compiler/GeneratedCode>
- [C9] <https://ghc.haskell.org/trac/ghc/wiki/Commentary/Compiler/SymbolNames>
- [C10] <https://ghc.haskell.org/trac/ghc/wiki/Commentary/Rts>
- [C11] <https://ghc.haskell.org/trac/ghc/wiki/Commentary/Rts/Storage/HeapObjects>
- [C12] <https://ghc.haskell.org/trac/ghc/wiki/Commentary/Rts/Storage/Stack>
- [C13] <https://ghc.haskell.org/trac/ghc/wiki/Commentary/Rts/Storage/GC>
- [C14] <https://ghc.haskell.org/trac/ghc/wiki/Commentary/Rts/HaskellExecution>
- [C15] <https://ghc.haskell.org/trac/ghc/wiki/Commentary/Rts/HaskellExecution/Registers>
- [C16] <https://ghc.haskell.org/trac/ghc/wiki/Commentary/Rts/HaskellExecution/PointerTagging>
- [C17] <https://ghc.haskell.org/trac/ghc/wiki/Commentary/Rts/Scheduler>
- [C18] <https://ghc.haskell.org/trac/ghc/wiki/Commentary/Rts/STM>
- [C19] <https://ghc.haskell.org/trac/ghc/wiki/Commentary/Libraries>

# References

## Source code

- [S1] `includes/stg/Regs.h`
- [S2] `includes/stg/MachRegs.h`
- [S3] `includes/rts/storage/ClosureTypes.h`
- [S4] `includes/rts/storage/Closures.h`
- [S5] `includes/rts/storage/TSO.h`
- [S6] `includes/rts/storage/InfoTables.h`
- [S7] `compiler/main/DriverPipeline.hs`
- [S8] `compiler/main/HscMain.hs`
- [S9] `compiler/cmm/CmmParse.y.source`
- [S10] `compiler/codeGen/StgCmmForeign.hs`
- [S11] `compiler/codeGen/Stg*.hs`
- [S12] `rts/PrimOps.cmm`
- [S13] `rts/RtsMain.c`
- [S14] `rts/RtsAPI.c`
- [S15] `rts/Capability.h`
- [S16] `rts/Capability.c`
- [S17] `rts/Schedule.c`
- [S18] `rts/StgCRun.c`
- [S19] `rts/StgStartup.cmm`
- [S20] `rts/StgMiscClosures.cmm`
- [S21] `rts/HeapStackCheck.cmm`
- [S22] `rts/Threads.c`
- [S23] `rts/Task.c`
- [S24] `rts/Timer.c`
- [S25] `rts/sm/GC.c`
- [S26] `rts/Sparks.c`
- [S27] `rts/WSDeque.c`
- [S28] `rts/STM.h`
- [S29] `rts/posix/Signals.c`
- [S30] `rts/win32/ThrIOManager.c`
- [S31] `libraries/base/GHC/MVar.hs`
- [S32] `libraries/base/GHC/Conc/IO.hs`
- [S33] `libraries/base/GHC/Conc/Sync.hs`
- [S34] `libraries/base/GHC/Event/Manager.hs`
- [S35] `libraries/base/GHC/Event/Thread.hs`
- [S36] `libraries/base/GHC/IO/BufferedIO.hs`
- [S37] `libraries/base/GHC/IO/FD.hs`
- [S38] `libraries/base/GHC/IO/Handle/Text.hs`
- [S39] `libraries/base/System/IO.hs`
- [S40] `libraries/base/System/Posix/Internals.hs`
- [S41] `AutoApply.o (utils/genapply/GenApply.hs)`

*Connect the algorithm and transistor*